Multi-Task Particle Swarm Optimization with Knowledge-Guided Elite Refinement and Two-Phase Weight Scheduling for High-Dimensional Feature Selection

Anonymous Author[a]

[a]*Anonymous Institution (masked for double-blind review),*

Abstract

Feature selection in high-dimensional data remains a central challenge for evolutionary computation. Particle swarm optimization suffers from three structural deficiencies: random initialization provides insufficient guidance in sparse high-dimensional spaces, uniform mutation discards cross-iteration feature-importance signals, and static objective preferences cannot adapt to different evolutionary phases. To address these limitations, we propose KDS-EL (Knowledge-Guided Dual-Phase Swarm Evolutionary Learning), a closed-loop knowledge-guided evolutionary framework. First, to overcome cold-start inefficiency, we propose mutual-information-guided initialization with adaptive re-initialization, injecting prior knowledge and recovering population diversity when stagnation occurs. Second, to exploit discarded cross-iteration feature-importance signals, we propose a knowledge-guided elite refinement strategy that accumulates feature-importance knowledge and steers elite solutions toward historically effective features. Third, to align objective weights with search phases, we propose two-phase weight scheduling that emphasizes early dimensionality reduction and late-stage accuracy refinement. On 20 high-dimensional benchmarks, KDS-EL ranks first among all competitors, achieving a mean accuracy of 0.9574 with only 87.8 selected features, outperforming the strongest single- and multi-objective competitors by +3.97 and +4.42 percentage points; Wilcoxon and Holm–Bonferroni tests confirm statistically significant improvements over all 13 competitors.

Keywords: feature selection, particle swarm optimization, mutual information, knowledge-guided elite refinement, two-phase weight scheduling

1. Introduction

High-dimensional data have become a common challenge in classification tasks across genomics, text mining, and computer vision. Datasets in these domains often contain tens of thousands of features while providing only tens or hundreds of labeled samples, a regime denoted $p \gg n$. Under such conditions, standard classifiers are prone to overfitting, computational overhead grows rapidly, and model interpretability degrades, which are typical consequences of the curse of dimensionality [1, 2]. Feature selection (FS) mitigates these effects by identifying a compact subset of informative features before model training. However, the space of candidate feature subsets expands exponentially with the number of dimensions D, making FS an NP-hard combinatorial optimization problem [3]; in addition, marginal feature scores cannot fully capture the feature interactions that determine classifier performance.

FS methods are commonly categorized into filter, wrapper, and embedded approaches [4, 1]. Filter methods evaluate features with statistical measures and are computationally efficient, but they often ignore feature interactions and classifier-specific relevance. Embedded methods integrate selection into model training and offer a balance between efficiency and accuracy, while remaining tied to the assumptions of the underlying model. Wrapper methods train a classifier on each candidate subset and therefore tend to produce higher accuracy, yet their search cost grows rapidly with dimensionality. To search the large, non-convex solution space of FS, evolutionary computation has been widely adopted [5, 6, 7]. Population-based stochastic search can explore candidate feature subsets without gradient information or restrictive model assumptions, and evolutionary algorithms such as particle swarm optimization (PSO), genetic algorithms, and differential evolution have shown strong performance on high-dimensional FS benchmarks.

PSO is among the most widely used evolutionary algorithms for FS because of its compact representation, simple update rules, and fast convergence. Nevertheless, standard PSO-based FS methods face three structural limitations when applied to sparse, high-dimensional search spaces. First, random initialization scatters the initial swarm across the entire binary hypercube, while an effective feature subset usually activates only a small fraction of the available dimensions; a substantial portion of the limited evaluation budget is therefore spent on dense, low-quality solutions before search can concentrate on sparse regions. Second, conventional uniform bit-flip mutation treats every feature identically and discards the per-feature improvement information accumulated during previous evaluations, so the search cannot progressively exploit historically effective dimensions. Third, most PSO-based FS methods combine classification error and feature ratio using a fixed weight vector throughout the entire search, although feature reduction and accuracy refinement have different temporal requirements: redundant features should be removed aggressively at the beginning, while accuracy should be refined on a compact subset at the later stage.

Recent evolutionary FS studies have attempted to mitigate these issues through improved initialization, adaptive mutation, or modified fitness formulations [8, 9, 10, 11]. These improvements are valuable, but they generally target individual components and do not jointly address cold-start inefficiency, knowledge accumulation, and rigid objective weighting. Multi-task PSO has recently shown particular promise for fixed-length, high-dimensional FS: for example, MEL [12] uses fixed-length binary encoding and cooperative subpopulations to scale to tens of thousands of features and demonstrates strong efficiency. However, existing multi-task PSO frameworks still face unresolved challenges: population construction usually provides little prior guidance in sparse spaces, feature-importance evidence obtained during search is not systematically accumulated, and objective preferences remain fixed across evolutionary phases.

To address these limitations, this paper proposes KDS-EL (Knowledge-Guided Dual-Phase Swarm Evolutionary Learning), a closed-loop knowledge-guided evolutionary framework that integrates three complementary mechanisms into the multi-task PSO architecture. Knowledge is progressively introduced, accumulated, and reused throughout the evolutionary process. First, mutual-information-guided initialization with adaptive re-initialization injects prior feature-relevance knowledge into the initial population and restarts stagnated individuals, reducing cold-start waste. Second, the Knowledge-guided Elite Refinement (KER) strategy accumulates a cross-iteration feature-importance matrix from fitness comparisons already performed during search, applies exponential decay to discount obsolete evidence, and steers elite solutions toward historically effective features within a bounded refinement budget. Third, Two-Phase Weight Scheduling switches the dual-objective weights from aggressive feature reduction in the early search to fine-grained accuracy refinement in the later search, aligning the optimization preference with the temporal structure of FS. The three mechanisms operate at complementary levels of the search process and are independently verifiable through component contribution analysis.

The main contributions of this paper are summarized as follows:

1. We propose an MI-guided initialization scheme with adaptive re-initialization, described in Section 4.2, that injects feature-relevance knowledge into population construction and restores diversity when stagnation occurs.
2. We propose the KER strategy, described in Section 4.4, which accumulates cross-iteration feature-importance evidence from fitness comparisons already performed during search, discounts obsolete evidence by exponential decay, and performs knowledge-biased elite perturbation with 90% guided and 10% random flips within a bounded refinement budget.
3. We propose Two-Phase Weight Scheduling, described in Section 4.5, which replaces static dual-objective weights with an early feature-

Email address: corresponding.author@example.com (Anonymous Author)

reduction phase and a late accuracy-refinement phase, aligning optimization preferences with the temporal structure of FS.

4. We conduct systematic experiments on 20 high-dimensional benchmarks spanning 500 to 54,675 features and 14 fixed-length methods, including component contribution analysis and parameter sensitivity analysis.

Extensive experiments show that KDS-EL ranks first with a mean accuracy of 0.9574 and 87.8 selected features on average, outperforming the strongest single-objective and multi-objective competitors by 3.97 and 4.42 percentage points, respectively, while achieving a 3.4× speedup over MEL (51.4 s vs. 175.0 s); Wilcoxon and Holm–Bonferroni tests confirm statistically significant improvements over all 13 competitors.

The remainder of this paper is organized as follows. Section 2 reviews feature selection and evolutionary FS methods. Section 3 introduces the necessary background, including PSO, MI-based feature filtering, and the MEL framework. Section 4 presents the KDS-EL framework and its three mechanisms in detail. Section 5 describes the experimental configuration, reports the main results, quantifies the contribution of each mechanism, and examines parameter sensitivity. Section 6 concludes the paper and discusses future directions.

2. Related Work

2.1. Feature Selection Paradigms

Feature selection is a fundamental pre-processing step in high-dimensional data analysis, aiming to identify a minimal subset of relevant features that preserves or improves classification performance while reducing computational cost and enhancing interpretability [1, 2]. Methods fall broadly into three paradigms: filter, wrapper, and embedded.

Filter methods evaluate features independently of any learning algorithm, relying on intrinsic data properties such as variance, correlation, or eigenvector centrality [13, 14]. Representative techniques include the minimum Redundancy Maximum Relevance (mRMR) criterion [15] and the Fast Correlation-Based Filter (FCBF) [16]. These methods are computationally efficient and scale readily to high-dimensional settings, yet they disregard feature interactions and may retain redundant subsets under strong inter-feature correlation.

Wrapper methods treat feature selection as a search problem, evaluating candidate subsets by the predictive performance of a specific classifier [4]. Hybrid genetic algorithms [17] and other meta-heuristic wrappers can capture complex feature interdependencies, but their computational cost grows rapidly with dimensionality, as each candidate subset requires a full classifier training-and-evaluation cycle.

Embedded methods integrate feature selection into model training. Lasso and its bootstrapped variant Bolasso [18] enforce sparsity through L_1 regularization; gradient-boosting-based approaches such as FS-GBDT [19] leverage tree-split importance; and hybrid algorithms like ILRC [20] combine regularization with clustering. Embedded methods offer a favorable speed–accuracy trade-off, but their feature rankings are tied to specific model assumptions.

At high dimensionality, each paradigm faces distinct structural bottlenecks: filters are fast but oblivious to feature interactions, wrappers are accurate but face a search space that grows exponentially with dimension, and embedded methods depend on model-specific regularization biases. These complementary weaknesses motivate evolutionary computation, which decouples the search mechanism from the evaluation metric and navigates the exploration–exploitation trade-off through population-based global search.

2.2. Single-Objective Evolutionary Feature Selection

Evolutionary algorithms have been extensively applied to feature selection for their population-based global search and gradient-free nature. Early work adapted classical meta-heuristics—binary PSO (BPSO) [7], ant colony optimization [21], artificial bee colony [22], and genetic algorithms with support vector machines [23]—to the binary subset selection problem, establishing the template of a population of candidate subsets evolving under selection, crossover, and mutation guided by a scalar fitness function.

Because standard evolutionary operators degrade at high dimensionality, specialized single-objective methods have been developed: self-adaptive PSO adjusting swarm size and mutation rate [24], progressively refined PSO alternating between coarse- and fine-grained search [25], and problem-specific heuristics—FWPSO [8], which biases velocity updates with per-feature importance; SaWDE [9] with adaptive weighting; IGPSO [10], which estimates relevance via a multi-layer perceptron; and PSO-CSM [11] integrating multiple evaluation criteria. Learning-based control and space decomposition are also explored: reinforcement-learning-aided PSO [26] and divide-and-conquer search [27]. These methods refine search operators but do not systematically exploit accumulated feature-importance knowledge harvested from the search itself.

2.3. Mutual-Information-Guided and Importance-Aware Evolutionary Methods

Mutual information (MI) quantifies the reduction in label uncertainty provided by a feature [28] and is the most widely used relevance score in hybrid filter-wrapper feature selection. In filter form, mRMR [15] and FCBF [16] rank features before any classifier is trained, offering efficient coarse pruning at the price of ignoring feature interactions, as discussed in Section 2.1. Within evolutionary wrappers, MI and related importance scores have also been used to shape the search itself: Yu et al. [29] rank features by MI and initialize a multi-population differential evolution with the top-ranked subset; FWPSO biases velocity updates with per-feature importance scores [8]; and IGPSO learns a feature-importance estimate with a multilayer perceptron [10]. These methods demonstrate that prior relevance knowledge can improve the starting point or update direction of evolutionary search.

Such filter-driven schemes, however, consume the injected knowledge one-shot: MI rankings are computed before optimization and are never revised by the search itself. The per-feature evidence accumulated across iterations—which additions or removals improved fitness—is discarded, and subsequent operators remain uniform. KDS-EL instead treats feature-importance knowledge as an in-loop state: MI scores bootstrap the initial population, while the KER strategy maintains a cumulative feature-importance matrix that is updated from the fitness comparisons already performed during search and re-deployed to refine elite solutions within a bounded refinement budget.

2.4. Multi-Objective Evolutionary Feature Selection

Multi-objective feature selection (MOFS) treats classification accuracy and feature count as competing objectives, producing a Pareto front of non-dominated solutions. Xue et al. [30] pioneered multi-objective PSO for feature selection, showing that jointly optimizing accuracy and subset size yields more diverse and informative solutions than scalarized approaches. Hu et al. [31] extended this framework with fuzzy cost functions for imbalanced classification.

The multi-objective evolutionary algorithm family provides several architectures for this problem. NSGA-II-FS [32] uses non-dominated sorting and crowding distance to maintain Pareto-front diversity. MOEA/D-FS [33] decomposes the problem into scalar subproblems solved cooperatively. MOPSO-FS [34] extends single-objective PSO with an external archive and Pareto-dominance-based leader selection. Hancer et al. [35] combined evolutionary search with filter-based ranking under both single- and multi-objective settings.

Despite their conceptual advantages, MOFS methods incur substantial overhead: maintaining an external archive, non-dominated sorting, and Pareto-front management all cost significant time (in our experiments, over 110 seconds on average). Moreover, under fixed-length binary encoding at high dimensionality, the Pareto front expands combinatorially and convergence slows, limiting scalability at gene-expression scale. Surrogate-assisted multi-objective search [36] can mitigate evaluation cost, yet surrogate accuracy and training overhead under extreme dimensionality remain open problems. KDS-EL therefore adopts the single-objective weighted formulation, which delivers a single compact solution directly, and retains MOFS methods only as comparison methods.

2.5. Evolutionary Multi-Task Learning

Evolutionary multi-task learning solves multiple related tasks jointly, with knowledge transfer accelerating convergence. Zhang and Yang [37] surveyed the multi-task learning foundations on which such evolutionary approaches build, and several instantiations target feature selection: a minion-based multitasking framework [38], PSO-based multitasking for high-dimensional classification [39, 40], multitask bee colony optimization for hyperspectral band selection [41], and multitasking for large-scale multi-objective optimization [42].

Of particular relevance to fixed-length high-dimensional feature selection, MEL [12] divides a PSO population into two subpopulations that cooperate through cross-task knowledge transfer, sharing feature-importance knowledge while searching independently, and provides the multi-task PSO framework on which KDS-EL builds in Section 4.

2.6. Gene-Expression Feature Selection

High-dimensional gene expression data present a particularly challenging testbed: the number of features routinely exceeds the number of samples by two to three orders of magnitude—a "small N, large D" regime where conventional statistical methods break down and overfitting is pervasive [43, 44]. Representative applications include cancer staging [45], Alzheimer's classification from neuroimaging features [46], biomarker gene identification [8], multi-cancer risk module discovery [19, 20], and surrogate-assisted multi-objective gene selection from large-scale single-cell RNA sequencing data [47].

Two structural properties amplify the difficulty. First, feature redundancy is exceptionally high: co-regulated genes exhibit correlated expression, so distinct subsets frequently yield near-identical accuracy. Second, a non-trivial fraction of gene datasets reach accuracy ceilings—exceeding 0.998 under multiple methods—rendering accuracy alone insufficient to distinguish algorithmic quality; on these datasets the selected feature count becomes the primary discriminatory metric. Evaluation must therefore jointly consider classification accuracy and subset compactness.

Recent methodological development in this area is extensive; surveys of evolutionary computation for biomarker identification in complex diseases [3] and instance-selection-assisted evolutionary feature selection [48] both target the small-sample, ultra-high-dimensional regime central to gene-expression analysis.

2.7. Positioning of KDS-EL

Across the evolutionary feature selection literature reviewed above, three gaps recur in high-dimensional feature selection. First, initialization is typically random, wasting a large fraction of the evaluation budget in a sparse search space. Second, search operators treat all dimensions uniformly and discard the cross-iteration evidence of which features have improved fitness. Third, the objective weights balancing accuracy against subset size remain fixed throughout search, although feature selection has a natural temporal structure: noise elimination is best performed early, and accuracy refinement late.

KDS-EL addresses these gaps as a closed-loop knowledge-guided evolutionary framework built upon the multi-task PSO architecture of MEL [12]. Mutual-information-guided initialization with adaptive re-initialization injects prior feature-relevance knowledge at the outset; the KER strategy accumulates a feature-importance matrix from the fitness comparisons already performed during search, discounts obsolete evidence by exponential decay, and flips elite solutions toward effective features within a bounded refinement budget; and Two-Phase Weight Scheduling aligns the objective weights with a reduction-then-refinement schedule. The fixed-length encoding inherited from MEL keeps the framework efficient at tens of thousands of dimensions, and each mechanism is independently verifiable through the component contribution analysis in Section 5.8.

3. Preliminaries

3.1. Particle Swarm Optimization and Its Binary Adaptation

Particle swarm optimization, introduced by Kennedy and Eberhart [5], maintains a population of candidate solutions that traverse the search space under the guidance of personal and collective historical experience. Each particle i holds a continuous position $\mathbf{x}_i \in \mathbb{R}^D$ and a velocity $\mathbf{v}_i \in \mathbb{R}^D$. At iteration t, the velocity and position are updated as

$$\mathbf{v}_i^{(t+1)} = \omega \mathbf{v}_i^{(t)} + c_1 r_1 (\mathbf{p}_{\text{best},i} - \mathbf{x}_i^{(t)}) + c_2 r_2 (\mathbf{g}_{\text{best}} - \mathbf{x}_i^{(t)}), \tag{1}$$

where ω is the inertia weight introduced by Shi and Eberhart [6], c_1 and c_2 are acceleration coefficients, r_1 and r_2 are independent random scalars drawn uniformly from $[0, 1]$, and $\mathbf{p}_{\text{best},i}$ and $\mathbf{g}_{\text{best}}$ denote the personal and global best positions, respectively.

For feature selection, each candidate solution is a binary mask indicating whether a feature is selected. The binary adaptation proposed by Kennedy and Eberhart [7] maps the continuous position to a binary decision through a sigmoid transfer function:

$$S(v_{i,d}) = \frac{1}{1 + e^{-v_{i,d}}}, \qquad x_{i,d} = \begin{cases} 1 & \text{if rand}() < S(v_{i,d}), \\ 0 & \text{otherwise}, \end{cases} \tag{2}$$

KDS-EL, in contrast, binarizes continuous positions with the fixed threshold $\theta = 0.6$ described in Section 4.2. In the $p \gg n$ regime, standard binary PSO converges slowly and tends toward premature stagnation, because the velocity update provides no domain-grounded gradient toward sparse, high-accuracy regions.

3.2. Mutual-Information-Based Feature Filtering

Mutual information measures the reduction in label uncertainty provided by a feature and is a standard filter score in high-dimensional feature selection [49, 28, 15]. For a discrete feature f_j and the class label Y, the mutual information is defined as

$$I(f_j; Y) = \sum_{b \in \{0,1\}} \sum_{c \in \mathcal{C}} p(f_j = b, Y = c) \, \log_2 \frac{p(f_j = b, Y = c)}{p(f_j = b)\, p(Y = c)}, \tag{3}$$

where the sums run over the binary feature values and the class labels. Filtering by mutual information computes the score of every feature, ranks the features in descending order, and retains only a small fraction of the top-ranked dimensions before a classifier is trained. Algorithm 1 summarizes this procedure with a retention ratio k [49]; in high-dimensional gene-expression settings, retaining 2–5% of the features is typical.

Algorithm 1 Feature filtering based on mutual information

Require: Training dataset Ω, feature count p, class labels Y, retention ratio k
Ensure: Filtered feature set S_{MI}
1: $S_{\text{MI}} \leftarrow \emptyset$
2: **for** $i = 1$ to p **do**
3: Compute the mutual information $I(f_i; Y)$ ▷ (3)
4: **end for**
5: Sort features by descending MI score and store the ranked indices in IDX
6: $S_{\text{MI}} \leftarrow$ the top $\lfloor p \cdot k \rfloor$ features of IDX
7: **return** S_{MI}

KDS-EL does not hard-prune the search space with this filter. Instead, Section 4.2 reuses the same MI ranking to seed the initial population, preserving the full feature space while transferring the filter's knowledge into the evolutionary loop.

3.3. Multi-Task Evolutionary Learning

MEL [12] is a PSO-based multi-task evolutionary learning method for high-dimensional feature selection. Its population of N particles is divided into two equal-sized subpopulations, Sub 1 and Sub 2, each searching for an optimal feature subset with a distinct strategy. All particles share the same fixed-length real-valued representation, and a threshold θ converts each position into the binary decision of whether a feature is selected. The two subpopulations share a feature-importance matrix $W \in \mathbb{R}^D$ as the knowledge state of the search. Initially zero, W is updated after each generation from the accuracy change of every particle: when the accuracy of a particle improves, the importance scores of features newly selected by that particle increase and the importance scores of discarded features decrease by the same amount, with the opposite adjustment when accuracy declines.

Knowledge transfer connects the two tasks in complementary directions. Sub 1 follows the standard PSO update with an additional learning term, namely the best solution of Sub 2, so the experience of the second subpopulation guides the first. Sub 2 exploits the learned weights directly, excluding features with non-positive weights and selecting each remaining feature with probability proportional to its weight. The two subpopulations therefore search independently while continuously sharing feature-importance knowledge across tasks.

KDS-EL inherits three elements of MEL. First, the fixed-length binary decision vector is shared by all components of Section 4. Second, the scalar fitness with which MEL maintains personal, subpopulation, and global bests is the ranking criterion reused by the KER strategy to select elite particles for refinement. Third, the fitness function of MEL is a weighted sum of classification error and subset sparsity:

$$f(\mathbf{x}) = \alpha \cdot (1 - \text{Acc}(\mathbf{x})) + \beta \cdot \frac{\|\mathbf{x}\|_0}{D}, \tag{4}$$

where $\text{Acc}(\mathbf{x})$ is the k-nearest-neighbor accuracy of the selected features, $\|\mathbf{x}\|_0$ is the number of selected features, and $\alpha = 0.9$ and $\beta = 0.1$ are fixed weights [12]. Since both weights remain constant throughout search, MEL cannot adapt its objective emphasis to the temporal structure of feature selection, and Section 4 replaces this static weighting with Two-Phase Weight Scheduling. The dominant per-generation cost of MEL is the evaluation of the N particles, each requiring KNN-based five-fold cross-validation, giving an overall complexity of $O(T \cdot N \cdot T_{\text{eval}})$.

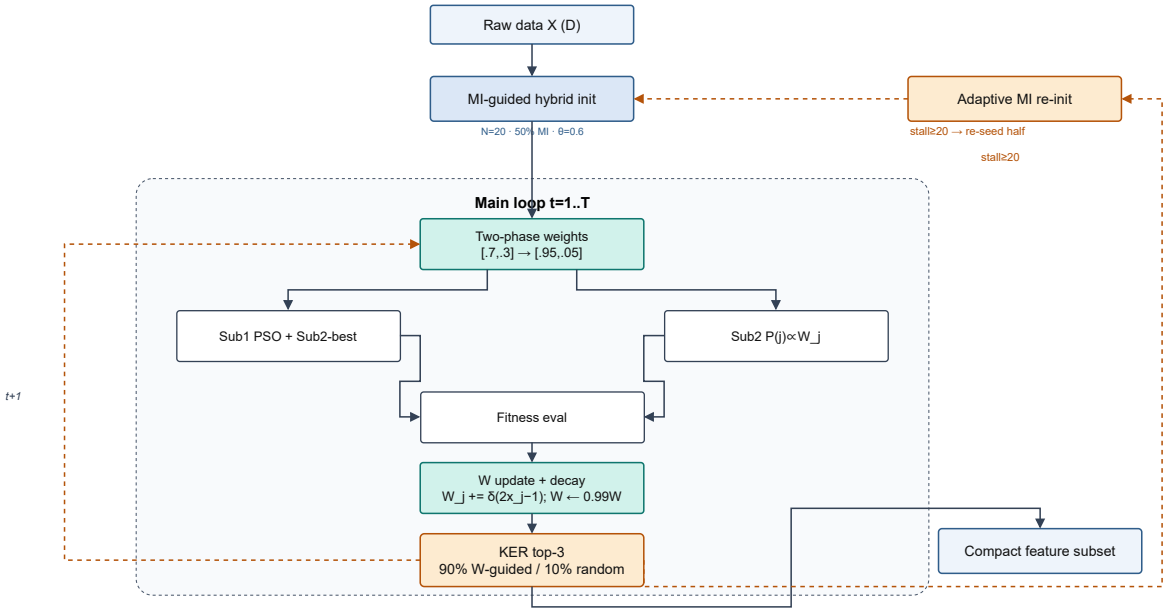


Figure 1: Overall framework of KDS-EL: raw high-dimensional data is transformed into a compact feature subset through mutual-information-guided initialization, the two-phase weighted main loop with cumulative feature-importance accumulation, and the KER strategy.

4. Proposed KDS-EL Algorithm

4.1. Overall Framework

KDS-EL is a dual-subpopulation evolutionary algorithm integrating three complementary mechanisms within the multi-task particle swarm optimization framework of MEL [12]. Figure 1 shows the complete pipeline transforming raw high-dimensional data into a compact, high-accuracy feature subset through knowledge-guided operations.

The algorithm takes as input a dataset with D features and C classes. In the preprocessing stage, mutual information between each feature and the class label is computed to produce a quality ranking over all D dimensions, as detailed in Section 4.2. An initial population of N fixed-length binary individuals is then constructed by mixing MI-ranked seeds with randomly initialized individuals at a controlled ratio MI_init_ratio = 0.5, followed by binarization with the fixed threshold $\theta = 0.6$. This hybrid scheme supplies a quality-informed starting point while preserving the diversity necessary for effective exploration.

The population enters the main iterative loop. Each generation opens with dual-objective fitness evaluation under the Two-Phase Weight Scheduling weights defined in Section 4.5. Observed improvements update the cumulative feature-importance matrix $W \in \mathbb{R}^D$ with exponential decay $\gamma = 0.99$, maintained as in Section 4.3. The top elite_k = 3 individuals are refined by the KER strategy of Section 4.4—a knowledge-guided flip mechanism with a 90%/10% knowledge-to-random split—and the population then undergoes the framework's dual-subpopulation update: Sub 1 follows the PSO rule with an additional attractor toward Sub 2's best, while Sub 2 samples features according to the cumulative importance matrix W. Each refined elite replaces its pre-refinement counterpart in place when fitness improves, avoiding the $2N$-candidate pool of earlier designs. On stagnation for stall_patience = 20 generations, adaptive MI re-initialization refreshes the lower-performing half of the population. After T generations, the best individual is decoded as the final feature subset.

KDS-EL builds upon the multi-task evolutionary learning framework of [12], inheriting its fixed-length binary encoding and cooperative multi-task mechanism while adding three knowledge-driven components. These operate at independent layers—feature-quality input in Section 4.2, knowledge accumulation and exploitation in Sections 4.3 and 4.4, and objective scheduling in Section 4.5—so each contribution is independently verifiable.

4.2. MI-guided Initialization and Adaptive Re-initialization

To overcome the cold-start inefficiency of random initialization in sparse high-dimensional search spaces, we propose an MI-guided initialization with adaptive re-initialization that seeds the initial population with feature-quality information derived from mutual information. Random initialization typically produces individuals with approximately half of the features active, far from the sparse configurations that characterize effective feature subsets, wasting a substantial fraction of the fixed evaluation budget on low-quality regions; the adaptive re-initialization restarts stagnated individuals with informed seeds.

Following the filter formulation of Section 3.2, all D features are ranked by descending MI score. With MI_init_ratio = 0.5, half of the initial population is seeded from the normalized MI score vector perturbed by small Gaussian noise, while the other half is initialized with uniform random positions. The resulting continuous position values are binarized with the fixed threshold $\theta = 0.6$, i.e., a feature is selected when its position value exceeds θ. This balanced design ensures that half the population inherits prior feature-quality knowledge while the other half maintains the exploratory coverage of random seeding. Algorithm 2 summarizes this hybrid initialization procedure.

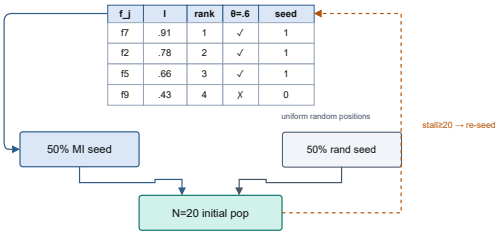


Figure 2: MI-guided initialization and adaptive re-initialization: normalized MI scores are ranked, thresholded at $\theta = 0.6$, and blended with uniform random seeds to form the $N = 20$ initial population; stagnation triggers re-seeding with MI-informed positions.

Algorithm 2 MI-guided Population Initialization

Require: MI score vector $\mathbf{m} \in [0,1]^D$, population size N, ratio ρ = MI_init_ratio = 0.5, threshold $\theta = 0.6$, noise scale $\sigma = 0.1$
Ensure: Hybrid binary population $\{P_i\}_{i=1}^N$

for $i = 1$ to N **do**
  if $i \le \rho N$ **then**
    $P_i \leftarrow \text{clip}(\mathbf{m} + \mathcal{N}(0, \sigma^2 I),\ 0,\ 1)$ ▷ MI-informed continuous seed
  else
    $P_i \leftarrow$ random continuous position in $[0,1]^D$
  end if
  Binarize the continuous position of P_i with threshold θ
end for
return $\{P_i\}_{i=1}^N$

To further guard against premature convergence, KDS-EL incorporates an adaptive MI re-initialization loop. When the global-best fitness stagnates for stall_patience = 20 consecutive generations, the lower-performing half of the population is re-seeded with MI-informed positions—the precomputed feature-quality scores perturbed by small Gaussian noise—and re-evaluated under the current objective weights. This restores diversity with feature-quality guidance, allowing the population to escape local optima that trap static-initialization algorithms. Algorithm 3 summarizes this procedure.

Algorithm 3 Adaptive MI Re-initialization

Require: Population $\{P_i, f_i^{\text{pbest}}\}_{i=1}^N$ with personal-best fitness values, MI scores $\mathbf{m}$, threshold θ, noise scale $\sigma = 0.1$
Ensure: Refreshed lower-performing half of the population

Sort individuals by personal-best fitness in ascending order
for each individual i in the lower-performing half **do**
  $P_i \leftarrow \text{clip}(\mathbf{m} + \mathcal{N}(0, \sigma^2 I),\ 0,\ 1)$
  Binarize P_i via threshold θ and re-evaluate its fitness under current weights
end for
Reset the stagnation counter

MI-guided initialization is a well-established filter-wrapper technique in the feature selection literature [28, 15]. It provides the foundational quality layer upon which the subsequent knowledge-accumulation and scheduling mechanisms build.

4.3. Cumulative Feature-Importance Matrix with Exponential Decay

To preserve the cross-iteration feature-importance signals discarded by uniform mutation, KDS-EL maintains a cumulative feature-importance matrix $W \in \mathbb{R}^D$, initialized to the zero vector, that aggregates these signals across iterations. This matrix serves as the knowledge substrate for the KER strategy described in Section 4.4.

After each generation's fitness evaluation, the importance matrix is updated from every evaluated candidate whose fitness improves over its personal best. Let C_i be the set of features whose selection status changed for individual i, and let $\delta_i = f_i^{\text{pbest}} - f_i^{\text{new}} > 0$ be the improvement magnitude. The update rule is

$$W_j \leftarrow W_j + \delta_i \cdot (2x_{ij} - 1), \quad j \in C_i, \tag{5}$$

where $x_{ij} \in \{0,1\}$ is the selection status of feature j in the improved candidate. Features newly selected by the candidate gain importance scores and discarded features lose importance scores, with the magnitude of the update equal to the fitness improvement; features whose selection status did not change receive no update, and candidates that do not improve leave W unchanged. After processing all individuals, the entire matrix undergoes exponential decay:

$$W \leftarrow \gamma \cdot W, \quad \gamma = 0.99. \tag{6}$$

Algorithm 4 consolidates the per-generation update-and-decay procedure.

Algorithm 4 Cumulative Feature-Importance Matrix Update with Decay

Require: Population with personal-best positions and evaluated candidates $\{P_i, \mathbf{p}_{\text{best},i}\}_{i=1}^N$, candidate fitness values $\{f_i^{\text{new}}\}_{i=1}^N$, personal-best fitness values $\{f_i^{\text{pbest}}\}_{i=1}^N$, feature-importance matrix W, decay factor $\gamma = 0.99$
Ensure: Updated feature-importance matrix W

for each individual i **do**
  if $f_i^{\text{new}} < f_i^{\text{pbest}}$ **then**
    Compute the changed-feature set C_i by comparing the candidate with $\mathbf{p}_{\text{best},i}$
    Compute $\delta_i = f_i^{\text{pbest}} - f_i^{\text{new}}$
    for each feature $j \in C_i$ **do**
      $W_j \leftarrow W_j + \delta_i \cdot (2x_{ij} - 1)$ ▷ (5)
    end for
  end if
end for
Apply exponential decay $W \leftarrow \gamma \cdot W$ ▷ (6)
return W

The decay factor $\gamma < 1$ serves two purposes: it prevents indefinite accumulation of stale signals (features useful early but later dominated lose their inflated importance scores once improvement feedback ceases), and it curbs feature inflation—without decay the matrix would grow monotonically and saturate the KER flip decision, erasing the relative ranking among features. $\gamma = 0.99$ retains long-range memory while providing gradual discounting, analogous to the exponential forgetting mechanisms employed in adaptive evolutionary algorithms [9]. The importance matrix W is the sole knowledge state maintained by KDS-EL across iterations; all KER decisions rely exclusively on its relative ordering. This distinguishes W from the feature-importance matrix of MEL: MEL learns importance scores from per-generation accuracy changes and uses them to bias the probabilistic selection of one subpopulation, whereas KDS-EL discounts the accumulated evidence with exponential decay and re-deploys it through deterministic, budget-bounded flips on elite solutions.

4.4. Knowledge-guided Elite Refinement Strategy

To exploit the cross-iteration feature-importance signals discarded by uniform mutation, we propose the KER strategy, which injects population-accumulated feature importance into elite solutions through a bounded, knowledge-guided perturbation. It operates once per generation, after fitness evaluation and importance-matrix update, and targets only the top elite_k = 3 individuals selected by fitness rank; each refined elite is re-evaluated immediately and retained only if it improves fitness.

For each elite individual $\mathbf{x} \in \{0,1\}^D$, a probabilistic branching decision is made. With probability rand_flip_p = 0.10, the individual enters a pure random exploration branch: $\lfloor \phi D \rfloor$ randomly selected features are toggled, where ϕ = flip_frac = 0.01. This branch guarantees ergodicity and prevents the search from locking onto stale or misleading importance accumulations. With probability $1 -$ rand_flip_p = 0.90, the individual enters the knowledge-driven refinement branch. Each feature j is assigned the flip score

$$\sigma_j = \begin{cases} W_j, & \text{if } x_j = 0, \\ -W_j, & \text{if } x_j = 1, \end{cases}$$

so that absent high-importance features and present low-importance features receive the largest scores. Let $\mathcal{S}$ contain the $\lfloor \phi D \rfloor$ features with the largest scores. The refined binary vector $\mathbf{x}'$ is then:

$$x'_j = \begin{cases} 1 - x_j, & \text{if } j \in \mathcal{S}, \\ x_j, & \text{otherwise.} \end{cases} \tag{7}$$

The two refinement branches are illustrated in Figure 3: the knowledge-driven branch flips the $\lfloor \phi D \rfloor$ highest-scoring features according to the importance matrix W, while the random branch flips the same number of uniformly selected bits; both branches re-evaluate the refined elite and retain it only when fitness improves.

The intuition behind (7) is straightforward: the cumulative importance matrix W encodes which features have been consistently correlated with fitness improvements. High-importance features absent from an elite are candidates for inclusion; low-importance features still present are candidates for removal. The flip fraction $\phi = 0.01$ is deliberately small—roughly one feature per hundred dimensions—so KER performs targeted local refinement rather than disruptive mutation. Crucially, acquiring the knowledge itself does not enlarge the evaluation budget: the importance matrix is updated from the fitness improvements already computed for the population. KER spends only a bounded refinement

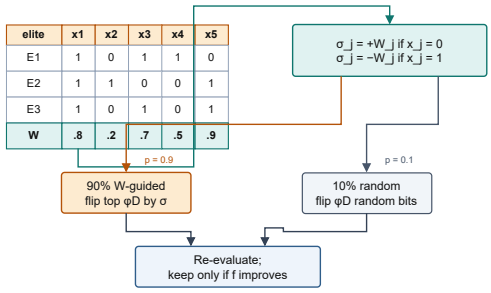


Figure 3: KER strategy: for each top-3 elite, the flip direction is determined by the cumulative importance matrix W with probability $p = 0.9$, or by uniform random bit selection with probability $p = 0.1$; refined elites are re-evaluated and kept only on improvement.

budget—re-evaluating the top elite_k = 3 elites per generation—which does not alter the dominant per-iteration complexity, since KNN cross-validation dominates and elite_k $\ll N$. This reuse of accumulated feature-level knowledge at bounded cost distinguishes KER from recent knowledge-guided metaheuristics [50]; the full refinement procedure is consolidated in Algorithm 5.

Algorithm 5 Knowledge-guided Elite Refinement

Require: Elite solution $\mathbf{x}_e$, feature-importance matrix W, flip fraction ϕ = flip_frac, random-flip probability p = rand_flip_p, fitness function f
Ensure: Refined elite $\mathbf{x}_e^\star$ with $f(\mathbf{x}_e^\star) \le f(\mathbf{x}_e)$
if rand() $< p$ **then**
  Select $\lfloor \phi D \rfloor$ features uniformly at random
  $\mathbf{x}' \leftarrow$ flip the selected bits of $\mathbf{x}_e$
else
  Assign each feature j the score $\sigma_j = W_j$ if $x_{e,j} = 0$, else $\sigma_j = -W_j$
  Select the $\lfloor \phi D \rfloor$ highest-scoring features
  $\mathbf{x}' \leftarrow$ flip the selected bits of $\mathbf{x}_e$
end if
Evaluate $f(\mathbf{x}')$ under current weights ▷ (9)
if $f(\mathbf{x}') < f(\mathbf{x}_e)$ **then**
  $\mathbf{x}_e^\star \leftarrow \mathbf{x}'$
else
  $\mathbf{x}_e^\star \leftarrow \mathbf{x}_e$
end if
return $\mathbf{x}_e^\star$

4.5. Two-Phase Weight Scheduling Strategy

To align objective weights with the temporal structure of feature selection, we propose a Two-Phase Weight Scheduling strategy that replaces fixed objective weights with an early feature-reduction phase and a late accuracy-refinement phase. Conventional PSO-based feature selection methods employ fixed objective weights throughout the entire search, typically balancing classification error and feature ratio with a static weighting such as $\mathbf{w} = [0.9, 0.1]$. This rigid scheduling fails to exploit a structural property of the feature selection problem: noise-feature elimination is a one-time cost best handled aggressively at the outset, whereas accuracy refinement benefits from sustained fine-tuning over a compact, already-pruned feature set. A single fixed weight vector must compromise between these two temporally distinct demands.

KDS-EL resolves this tension through a Two-Phase Weight Scheduling strategy. The total T iterations are partitioned by a split point at αT, where α = PhaseSplit = 0.7. The dual-objective weight vector $\mathbf{w}(t) = [w_a(t), w_f(t)]$ follows a step schedule:

$$\mathbf{w}(t) = \begin{cases} [0.7,\ 0.3], & t \le \alpha T \quad \text{(Phase 1: aggressive feature reduction)}, \\ [0.95,\ 0.05], & t > \alpha T \quad \text{(Phase 2: accuracy refinement)}. \end{cases} \tag{8}$$

The fitness of individual $\mathbf{x}$ at generation t is then:

$$f(\mathbf{x}, t) = w_a(t) \cdot (1 - \mathrm{Acc}(\mathbf{x})) + w_f(t) \cdot \frac{\|\mathbf{x}\|_0}{D}, \tag{9}$$

where $\mathrm{Acc}(\mathbf{x})$ is the KNN-based classification accuracy and $\|\mathbf{x}\|_0$ counts the selected features.

Algorithm 6 summarizes this schedule.

Algorithm 7 KDS-EL

Require: Dataset (X, Y) with D features and C classes; population $N = 20$ ($N/2$ per subpopulation), $T = 100$
Require: $\theta = 0.6$; MI_init_ratio = 0.5; WeightDecay = 0.99; elite_k = 3; flip_frac = 0.01
Require: rand_flip_p = 0.10; PhaseSplit = 0.7; stall_patience = 20
Ensure: Best feature subset S_{best}
Compute MI scores $I_j = I(f_j; Y)$ for $j = 1, \ldots, D$ ▷ (3)
Sort features by descending I_j
Normalize MI scores to $[0, 1]$ and store as $\mathbf{m}$
Initialize $W \leftarrow (0, \ldots, 0) \in \mathbb{R}^D$
for $i = 1$ to N **do**
  if $i \le N/2$ **then**
    Seed individual i from perturbed MI vector $\mathbf{m}$ ▷ Algorithm 2
  else
    Initialize individual i uniformly at random in $[0, 1]^D$
  end if
  Binarize position with threshold θ
end for
for $t = 1$ to T **do**
  Set dual-objective weights per Two-Phase schedule ▷ Algorithm 6
  if $t \le \alpha T$ **then** $\mathbf{w} \leftarrow [0.7, 0.3]$
  else $\mathbf{w} \leftarrow [0.95, 0.05]$
  end if
  Evaluate fitness of all individuals
  Update importance matrix from fitness improvements ▷ Algorithm 4
  Apply decay $W \leftarrow \gamma W$
  Apply KER refinement to the top elite_k individuals ▷ Algorithm 5
  Update Sub 1 by the PSO rule with the Sub 2 best attractor
  Update Sub 2 by importance-guided sampling from W
  if global best stagnates for stall_patience generations **then**
    Apply adaptive MI re-initialization ▷ Algorithm 3
  end if
end for
return best binary individual as feature subset S_{best}

Algorithm 6 Two-Phase Weight Scheduling

Require: Generation index t, total generations T, split ratio α = PhaseSplit = 0.7
Ensure: Dual-objective weight vector $\mathbf{w}(t)$
if $t \le \alpha T$ **then**
  $\mathbf{w}(t) \leftarrow [0.7, 0.3]$ ▷ Phase 1: aggressive feature reduction
else
  $\mathbf{w}(t) \leftarrow [0.95, 0.05]$ ▷ Phase 2: accuracy refinement
end if
return $\mathbf{w}(t)$

In Phase 1 ($t \le 0.7T$), the heavier feature-ratio penalty $w_f = 0.3$ drives rapid elimination of irrelevant and redundant features, reducing the search from thousands of dimensions to dozens. In Phase 2 ($t > 0.7T$), the weight prioritizes classification accuracy ($w_a = 0.95$), enabling fine-grained refinement over the surviving compact subset. This coarse-to-fine design contrasts with the fixed $\mathbf{w} = [0.9, 0.1]$ of KDS-EL w/o TPS: a constant weight cannot simultaneously deliver aggressive early feature reduction and precise late-stage accuracy tuning. Time-varying search parameters are a general PSO principle, most notably the linearly decreasing inertia weight of [6] [5].

KDS-EL employs an in-place update: after KER refinement, each elite's refined vector overwrites its original within the population when it improves fitness; no auxiliary pool is maintained, so the candidate set never expands to $2N$. This simplifies the architecture and yields a measurable speed improvement, as analyzed in Section 4.7, without sacrificing quality—KDS-EL w/o TPS already surpasses MEL by +6.20 pp.

4.6. Main Algorithm and Pseudocode

Algorithm 7 presents the complete KDS-EL procedure, consolidating the initialization and main loop; Algorithms 2, 3, 4, 5, and 6 detail the initialization, re-initialization, knowledge-update, refinement, and scheduling procedures it invokes.

The algorithm computes MI scores and generates a hybrid initial population, then runs $T = 100$ generations of phase-dependent fitness evaluation, importance-matrix update with decay, KER refinement of the top three elites via Algorithm 5, and the dual-subpopulation update in which Sub 1 follows the

PSO rule and Sub 2 samples features by importance-guided selection. Stagnation triggers adaptive MI re-initialization via Algorithm 3, and the best encoding is returned as the feature subset.

4.7. Complexity Analysis

The computational cost of KDS-EL decomposes into a one-time initialization cost and per-iteration costs. The initial MI ranking requires computing $I(f_j; Y)$ for all D features and sorting the scores, costing $O(D \log D)$. Each generation incurs: (a) fitness evaluation over N individuals with KNN classification, $O(N \cdot T_{\text{eval}})$, where T_{eval} dominates all other per-iteration operations; (b) importance-matrix update from fitness improvements, $O(N \cdot D)$; (c) elite selection and KER refinement (sorting D importance scores, flipping ϕD features, and re-evaluating the elite_k refined elites), $O(\text{elite_k} \cdot (D + T_{\text{eval}}))$. The adaptive MI re-initialization executes only on stagnation, contributing an amortized $O(D)$ per iteration. Collecting terms, the total complexity is:

$$O(D \log D + T \cdot (N \cdot T_{\text{eval}} + N \cdot D + \text{elite_k} \cdot (D + T_{\text{eval}}))). \tag{10}$$

MEL [12] has complexity $O(T \cdot N \cdot T_{\text{eval}})$. The added $O(N \cdot D + \text{elite_k} \cdot (D + T_{\text{eval}}))$ per generation is bounded relative to $O(N \cdot T_{\text{eval}})$, since elite_k = $3 \ll N = 20$ and KNN cross-validation T_{eval} dominates. Empirically, KDS-EL completes in 51.4 s versus 175.0 s for MEL (3.4× speedup), confirming that the sparser subsets promoted by the knowledge-guided mechanisms shorten each evaluation and more than offset the marginal overhead.

5. Experimental Results and Analysis

This section reports the experimental outcomes across six complementary perspectives: overall classification accuracy, feature-selection compactness, statistical significance, ceiling-effect behavior, convergence dynamics, and computational efficiency. All aggregate numbers are computed over the 20 benchmark datasets described in Section 5.1. We then quantify the contribution of each mechanism in Section 5.8 and assess the sensitivity of KDS-EL to its hyper-parameters in Section 5.9.

5.1. Experimental Setup

This section details the experimental framework. The evaluation employs 20 benchmark datasets drawn from five domains: 15 gene expression microarrays (ALL1, ALL4, ALL_AML, Adenoma, Colon, DLBCL, Gastric, Gastric1, Gastric2, Leukaemia, Lymphoma, Myeloma, Prostate, SRBCT, and T1D), two text corpora (BASEHOCK and PCMAC), one image dataset (COIL20), one speech dataset (ISOLET), and one synthetic dataset (madelon). Table 2 summarizes their key characteristics. The datasets span feature dimensionalities from 500 (madelon) to 54,675 (T1D) and sample sizes from 36 (Adenoma) to 2,600 (madelon), capturing the hallmark $p \gg n$ configuration: 17 of 20 datasets possess more features than instances. Class structure ranges from 17 binary problems to SRBCT (4 classes), COIL20 (20), and ISOLET (26). The collection was selected to span low-, medium-, and ultra-high-dimensional regimes and to include eight gene-expression datasets—ALL1, ALL_AML, Adenoma, DLBCL, Gastric2, Leukaemia, Lymphoma, and SRBCT—on which methods routinely reach near-perfect accuracy (above 0.998). On these ceiling datasets accuracy alone ceases to discriminate among methods, so the number of selected features becomes the decisive metric; Section 5.5 pursues this evaluation.

Fourteen fixed-length feature selection methods are compared, organized into three groups. The first comprises the complete KDS-EL and the configurations obtained by successively removing its mechanisms, where TPS denotes Two-Phase Weight Scheduling: KDS-EL w/o TPS (without Two-Phase scheduling), KDS-EL w/o TPS&KER (additionally without the KER strategy), and MEL [12]—which constitute the ablation chain analyzed in Section 5.8. The second consists of five single-objective evolutionary methods: FWPSO [8], PSO-CSM [11], SaWDE [9], IGPSO [10], and the classic BPSO [7]. The third contains five MOFS methods jointly optimizing accuracy and subset size: W-MOPSO-FS (a self-built Wrapper–Filter hybrid), fMOPSO-FS [51], NSGA-II-FS [32], MOEA-D-FS [33], and MOPSO-FS [34]. Table 1 maps methods to references. All comparison methods use the hyper-parameters recommended in their original publications, retaining default values where a parameter is unspecified.

All methods share a unified configuration. A K-nearest neighbor (KNN) classifier with $k = 3$ serves as the base evaluator, with accuracy estimated via stratified 5-fold cross-validation with random shuffling. All algorithms use $N = 20$ candidate solutions and $T = 100$ iterations, and methods relying on continuous-to-binary mapping use threshold $\theta = 0.6$. Every experiment is repeated across 10 independent runs without fixing random seeds, letting the random-number generator flow naturally between runs. The KDS-EL parameter set is: knowledge weight decay WeightDecay = 0.99; phase split PhaseSplit = 0.7 with objective weights $\mathbf{w}_s = [0.7, 0.3]$ (Phase 1, feature reduction) and $\mathbf{w}_s = [0.95, 0.05]$ (Phase 2, refinement); KER flip fraction flip_frac = 0.01; random exploration probability rand_flip_p = 0.10; MI-init ratio MI_init_ratio = 0.5; elite count elite_k = 3; and stall patience stall_patience = 20 for adaptive MI re-initialization. Table 3 summarizes these settings. Experiments run on a desktop workstation with an Intel i5-13700KF CPU and 32 GB RAM.

For each method and dataset we report classification accuracy (mean and standard deviation over 10 runs), the number of selected features with its selection ratio, and runtime in seconds. Statistical significance follows the protocol of Demšar [52]: the Friedman rank test [53] on the 20-dataset accuracy matrix detects overall differences; pairwise one-sided Wilcoxon signed-rank tests ($p < 0.05$) compare KDS-EL with each of the 13 methods in the direction of superiority; and the resulting p-values are Holm–Bonferroni corrected [54] to control the family-wise error rate. This standard protocol establishes whether observed differences are statistically meaningful rather than artifacts of finite experimental repetition.

Following the table convention of [49], Table 4 additionally annotates every competitor cell with +, ≈, or − from a two-sided Wilcoxon rank-sum test ($p < 0.05$) over the 10 independent runs of that dataset. The symbols indicate that KDS-EL is significantly better than, statistically indistinguishable from, or significantly worse than the column method, respectively, on the corresponding dataset.

Table 1: The 14 fixed-length methods compared in this study.

Method	Category	Reference
KDS-EL	Proposed	This work
KDS-EL w/o TPS	Ablation	This work
KDS-EL w/o TPS&KER	Ablation	This work
MEL	Baseline	[12]
FWPSO	Single-objective	[8]
PSO-CSM	Single-objective	[11]
SaWDE	Single-objective	[9]
IGPSO	Single-objective	[10]
BPSO	Single-objective	[7]
W-MOPSO-FS	Multi-objective	Self-built
fMOPSO-FS	Multi-objective	[51]
NSGA-II-FS	Multi-objective	[32]
MOEA-D-FS	Multi-objective	[33]
MOPSO-FS	Multi-objective	[34]

Table 2: Dataset Description (20 Benchmark Datasets)

Dataset	Samples	Features	Classes	Domain
Adenoma	36	7,457	2	Gene expression
ALL_AML	72	7,129	2	Gene expression
ALL1	128	12,625	2	Gene expression
ALL4	93	12,625	2	Gene expression
Colon	62	2,000	2	Gene expression
DLBCL	77	7,129	2	Gene expression
Gastric	65	22,645	2	Gene expression
Gastric1	144	22,283	2	Gene expression
Gastric2	124	22,283	2	Gene expression
Leukaemia	72	7,129	2	Gene expression
Lymphoma	45	4,026	2	Gene expression
Myeloma	173	12,625	2	Gene expression
Prostate	102	12,625	2	Gene expression
SRBCT	83	2,308	4	Gene expression
T1D	101	54,675	2	Gene expression
BASEHOCK	1,993	4,862	2	Text
COIL20	1,440	1,024	20	Image
ISOLET	1,560	617	26	Speech
madelon	2,600	500	2	Synthetic
PCMAC	1,943	3,289	2	Text

Range: $n \in [36, 2600]$, $d \in [500, 54675]$, $c \in [2, 26]$

Table 3: Parameter Settings

Common Experimental Setup		
Population size (N)	=	20
Maximum iterations (T)	=	100
KNN neighbors (k)	=	3
Real-to-binary threshold (θ)	=	0.6
Cross-validation	=	5-fold stratified, shuffled
Number of independent runs	=	10
KDS-EL (full framework)		
Inertia weight (w)	=	0.9
Acceleration coefficients	=	$c_1 = c_2 = c_3 = 2.0$
Knowledge weight decay	=	0.99
Elite flip fraction	=	0.01
Elite count (k_e)	=	3
Random flip probability	=	0.10
MI-guided init ratio	=	0.5
Stall patience (P_s)	=	20
Phase split ratio	=	0.7
Phase-1 objective weights	=	$\mathbf{w}_s$ = [0.7, 0.3]
Phase-2 objective weights	=	$\mathbf{w}_s$ = [0.95, 0.05]
KDS-EL w/o TPS (no two-phase scheduling)		
Fixed objective weights	=	$\mathbf{w}_s$ = [0.9, 0.1]
MEL and KDS-EL w/o TPS&KER		
Inertia weight (w)	=	0.9
Acceleration coefficients	=	$c_1 = c_2 = c_3 = 2.0$
MI-guided init ratio (KDS-EL w/o TPS&KER)	=	0.5
Stall patience (KDS-EL w/o TPS&KER)	=	20
Competitor Methods		
Competitors (BPSO, FWPSO, IGPSO, PSO-CSM, SaWDE, NSGA-II-FS, MOPSO-FS, MOEA-D-FS, fMOPSO-FS, W-MOPSO-FS) use the hyper-parameters recommended in their original publications, with $N = 20$ and $T = 100$ unless the original paper specifies otherwise.		

5.2. *Overall Classification Accuracy*

Table 4 presents per-dataset and aggregate accuracy of the 14 methods. KDS-EL ranks first with a mean of 0.9574, leading KDS-EL w/o TPS by +0.86 pp, the strongest single-objective competitor FWPSO [8] by +3.97 pp, and the strongest multi-objective competitor W-MOPSO-FS by +4.42 pp. Against MEL [12]—the framework stripped of all three mechanisms—the margin is +7.06 pp.

The per-cell symbols in Table 4 make the comparison structure explicit: KDS-EL is significantly better than every competitor on at least 14 of the 20 datasets. The ≈ ties concentrate on ceiling-effect gene-expression problems and on the strongest ablation variant (KDS-EL w/o TPS), while the only − entries occur against multi-objective methods on ALL1, COIL20, SRBCT, BASEHOCK, and ISOLET, where those methods reach perfect or highest accuracy at the cost of substantially larger feature subsets, as detailed in Section 5.3.

KDS-EL attains the highest mean accuracy on 15 of 20 datasets, being the sole best method on 13 of these (ALL4, ALL_AML, Colon, DLBCL, Gastric, Gastric1, Gastric2, Leukaemia, Myeloma, PCMAC, Prostate, T1D, and madelon) and tying KDS-EL w/o TPS on Adenoma and Lymphoma. The margin is largest on low-sample gene-expression problems: on ALL4 (93 samples, 12,625 features) KDS-EL reaches 0.9522 versus 0.9418 for the next-best method; on Colon, 0.9500 versus 0.9347 for KDS-EL w/o TPS and 0.9221 for FWPSO.

The five datasets on which another method registers higher accuracy warrant honest discussion. On ALL1 and SRBCT, several multi-objective methods reach a perfect 1.0 while KDS-EL records 0.9998 and 0.9986; Section 5.5 shows these differences reflect saturation, not ranking reversal. On BASEHOCK, COIL20, and ISOLET, NSGA-II-FS [32] attains the highest accuracy (0.9238, 0.9971, and 0.9090 versus 0.9113, 0.9939, and 0.8894 for KDS-EL); these moderate-dimensionality, larger-sample problems favor non-dominated-sorting methods, which retain a structural advantage at the cost of 2,073, 283, and 185 features respectively—several times to an order of magnitude more than KDS-EL, as detailed in Section 5.3. KDS-EL also exhibits narrower accuracy dispersion across runs, as the standard deviations in Table 4 indicate, showing that MI-guided initialization and KER refinement suppress run-to-run variance alongside raising central tendency; its advantage is consistent rather than the artifact of a few lucky runs.

Table 4: Classification accuracy (mean ± std) of 14 fixed-length methods on 20 benchmark datasets. Bold indicates the best per dataset; +, ≈, and − in parentheses report two-sided Wilcoxon rank-sum tests ($p < 0.05$) over 10 independent runs between KDS-EL and the column method: + KDS-EL significantly better, ≈ no significant difference, − KDS-EL significantly worse. The Average row reports the mean across the 20 datasets.

Dataset	KDS-EL	KDS-EL w/o TPS	KDS-EL w/o TPS&KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MOPSO-FS	fMOPSO-FS	NSGA-II-FS	MOEA-D-FS	MOPSO-FS
ALL1	0.9998±0.0001	0.9996±0.0000(+)	0.9996±0.0000(+)	0.9664±0.0023(+)	0.9952±0.0015(+)	0.9900±0.0002(+)	0.9710±0.0011(+)	0.9802±0.0002(+)	0.9534±0.0003(+)	**1.0000±0.0000**(−)	**1.0000±0.0000**(−)	**1.0000±0.0000**(−)	0.9985±0.0031(+)	**1.0000±0.0000**(−)
ALL4	**0.9522±0.0061**	0.9418±0.0113(+)	0.9415±0.0112(+)	0.8704±0.0661(+)	0.8617±0.0304(+)	0.8363±0.0135(+)	0.8062±0.0217(+)	0.7914±0.0179(+)	0.7569±0.0069(+)	0.8068±0.0047(+)	0.8015±0.0089(+)	0.7694±0.0118(+)	0.7737±0.0184(+)	0.7654±0.0074(+)
ALL_AML	**0.9995±0.0002**	0.9990±0.0002(+)	0.9986±0.0016(+)	0.9539±0.0008(+)	0.9850±0.0067(+)	0.9748±0.0039(+)	0.9591±0.0012(+)	0.9669±0.0003(+)	0.9426±0.0034(+)	0.9943±0.0070(≈)	0.9900±0.0065(≈)	0.9857±0.0000(+)	0.9857±0.0000(+)	0.9857±0.0000(+)
Adenoma	**0.9996±0.0002**	**0.9996±0.0001**(≈)	0.9958±0.0081(+)	0.9410±0.0011(+)	0.9783±0.0143(+)	0.9665±0.0108(+)	0.9461±0.0014(+)	0.9527±0.0003(+)	0.9287±0.0002(+)	0.9800±0.0131(≈)	0.9743±0.0086(+)	0.9714±0.0000(+)	0.9714±0.0000(+)	0.9714±0.0000(+)
BASEHOCK	0.9113±0.0120	0.8997±0.0114(+)	0.8597±0.0062(+)	0.8518±0.0063(+)	0.8715±0.0095(+)	0.8637±0.0050(+)	0.8860±0.0065(+)	0.8479±0.0062(+)	0.8597±0.0036(+)	0.8847±0.0066(+)	0.8861±0.0045(+)	**0.9238±0.0082**(−)	0.9150±0.0086(≈)	0.8920±0.0034(+)
COIL20	0.9939±0.0011	0.9919±0.0008(+)	0.9619±0.0097(+)	0.9634±0.0017(+)	0.9894±0.0018(+)	0.9782±0.0006(+)	0.9763±0.0017(+)	0.9707±0.0010(+)	0.9573±0.0006(+)	0.9960±0.0009(−)	0.9960±0.0005(−)	**0.9971±0.0007**(−)	0.9958±0.0010(−)	0.9964±0.0009(−)
Colon	**0.9500±0.0075**	0.9347±0.0094(+)	0.9211±0.0092(+)	0.8548±0.0276(+)	0.9221±0.0203(+)	0.8880±0.0090(+)	0.8867±0.0110(+)	0.8636±0.0101(+)	0.8529±0.0059(+)	0.9033±0.0072(+)	0.8988±0.0071(+)	0.8731±0.0097(+)	0.8745±0.0145(+)	0.8806±0.0104(+)
DLBCL	**0.9995±0.0002**	0.9993±0.0004(≈)	0.9902±0.0016(+)	0.9483±0.0058(+)	0.9819±0.0092(+)	0.9819±0.0048(+)	0.9629±0.0055(+)	0.9655±0.0063(+)	0.9435±0.0069(+)	0.9974±0.0052(+)	0.9960±0.0061(≈)	0.9781±0.0060(+)	0.9781±0.0060(+)	0.9872±0.0058(+)
Gastric	**0.9782±0.0034**	0.9677±0.0093(+)	0.9748±0.0088(≈)	0.8955±0.0786(+)	0.9191±0.0340(+)	0.8874±0.0087(+)	0.8234±0.0191(+)	0.8564±0.0087(+)	0.8046±0.0093(+)	0.8908±0.0046(+)	0.8846±0.0077(+)	0.8154±0.0069(+)	0.8169±0.0083(+)	0.8215±0.0102(+)
Gastric1	**0.9800±0.0003**	0.9738±0.0004(+)	0.9642±0.0025(+)	0.9310±0.0155(+)	0.9597±0.0032(+)	0.9496±0.0024(+)	0.9348±0.0030(+)	0.9402±0.0018(+)	0.9145±0.0002(+)	0.9588±0.0021(+)	0.9581±0.0000(+)	0.9588±0.0021(+)	0.9554±0.0034(+)	0.9588±0.0021(+)
Gastric2	**0.9981±0.0029**	0.9936±0.0029(+)	0.9872±0.0038(+)	0.9522±0.0170(+)	0.9837±0.0049(+)	0.9710±0.0002(+)	0.9532±0.0038(+)	0.9623±0.0001(+)	0.9357±0.0034(+)	0.9837±0.0000(+)	0.9838±0.0002(+)	0.9765±0.0025(+)	0.9765±0.0025(+)	0.9832±0.0025(+)
ISOLET	0.8894±0.0056	0.8898±0.0090(≈)	0.8557±0.0063(+)	0.8501±0.0067(+)	0.8822±0.0109(≈)	0.8584±0.0043(+)	0.8903±0.0076(≈)	0.8468±0.0070(+)	0.8607±0.0027(+)	0.8883±0.0032(≈)	0.8874±0.0066(≈)	**0.9090±0.0059**(−)	0.8963±0.0059(−)	0.8828±0.0036(+)
Leukaemia	**0.9995±0.0002**	0.9993±0.0002(+)	0.9895±0.0017(+)	0.9534±0.0015(+)	0.9886±0.0059(+)	0.9775±0.0059(+)	0.9587±0.0017(+)	0.9682±0.0037(+)	0.9416±0.0002(+)	0.9886±0.0057(+)	0.9886±0.0057(+)	0.9857±0.0000(+)	0.9857±0.0000(+)	0.9857±0.0000(+)
Lymphoma	**0.9994±0.0003**	**0.9994±0.0004**(≈)	0.9868±0.0102(+)	0.8731±0.0239(+)	0.9404±0.0306(+)	0.9206±0.0098(+)	0.8841±0.0092(+)	0.8861±0.0139(+)	0.8585±0.0081(+)	0.9178±0.0102(+)	0.9200±0.0109(+)	0.8844±0.0133(+)	0.8867±0.0067(+)	0.8911±0.0067(+)
Myeloma	**0.9155±0.0101**	0.8971±0.0072(+)	0.8825±0.0084(+)	0.8450±0.0244(+)	0.8702±0.0069(+)	0.8550±0.0079(+)	0.8522±0.0071(+)	0.8304±0.0077(+)	0.8186±0.0056(+)	0.8535±0.0037(+)	0.8529±0.0060(+)	0.8594±0.0070(+)	0.8482±0.0115(+)	0.8425±0.0035(+)
PCMAC	**0.8705±0.0139**	0.8424±0.0163(+)	0.8007±0.0051(+)	0.7870±0.0080(+)	0.8103±0.0090(+)	0.8036±0.0095(+)	0.8256±0.0095(+)	0.7840±0.0058(+)	0.7982±0.0062(+)	0.8175±0.0056(+)	0.8105±0.0086(+)	0.8523±0.0083(+)	0.8505±0.0117(+)	0.8169±0.0035(+)
Prostate	**0.9701±0.0032**	0.9595±0.0045(+)	0.9638±0.0049(+)	0.8458±0.0083(+)	0.8887±0.0092(+)	0.8721±0.0088(+)	0.8711±0.0091(+)	0.8526±0.0102(+)	0.8485±0.0038(+)	0.8840±0.0031(+)	0.8839±0.0054(+)	0.8704±0.0120(+)	0.8703±0.0107(+)	0.8743±0.0052(+)
SRBCT	0.9986±0.0004	0.9984±0.0006(≈)	0.9944±0.0043(+)	0.9321±0.0085(+)	0.9692±0.0148(+)	0.9695±0.0091(+)	0.9659±0.0048(+)	0.9515±0.0116(+)	0.9488±0.0050(+)	**1.0000±0.0000**(−)	**1.0000±0.0000**(−)	0.9853±0.0074(+)	0.9829±0.0099(+)	0.9829±0.0080(+)
T1D	**0.8598±0.0094**	0.8232±0.0126(+)	0.8246±0.0084(+)	0.7679±0.0461(+)	0.7610±0.0169(+)	0.7318±0.0065(+)	0.7237±0.0049(+)	0.7113±0.0064(+)	0.7063±0.0041(+)	0.7256±0.0049(+)	0.7227±0.0053(+)	0.7118±0.0063(+)	0.7088±0.0088(+)	0.7179±0.0048(+)
madelon	**0.8822±0.0180**	0.8659±0.0246(≈)	0.7679±0.0121(+)	0.7521±0.0095(+)	0.7963±0.0165(+)	0.7669±0.0229(+)	0.8106±0.0089(+)	0.7497±0.0169(+)	0.7652±0.0071(+)	0.7920±0.0049(+)	0.7845±0.0048(+)	0.8214±0.0097(+)	0.8122±0.0055(+)	0.7808±0.0071(+)
Average	**0.9574**	0.9488	0.9325	0.8868	0.9177	0.9021	0.8944	0.8839	0.8698	0.9132	0.9110	0.9065	0.9042	0.9009

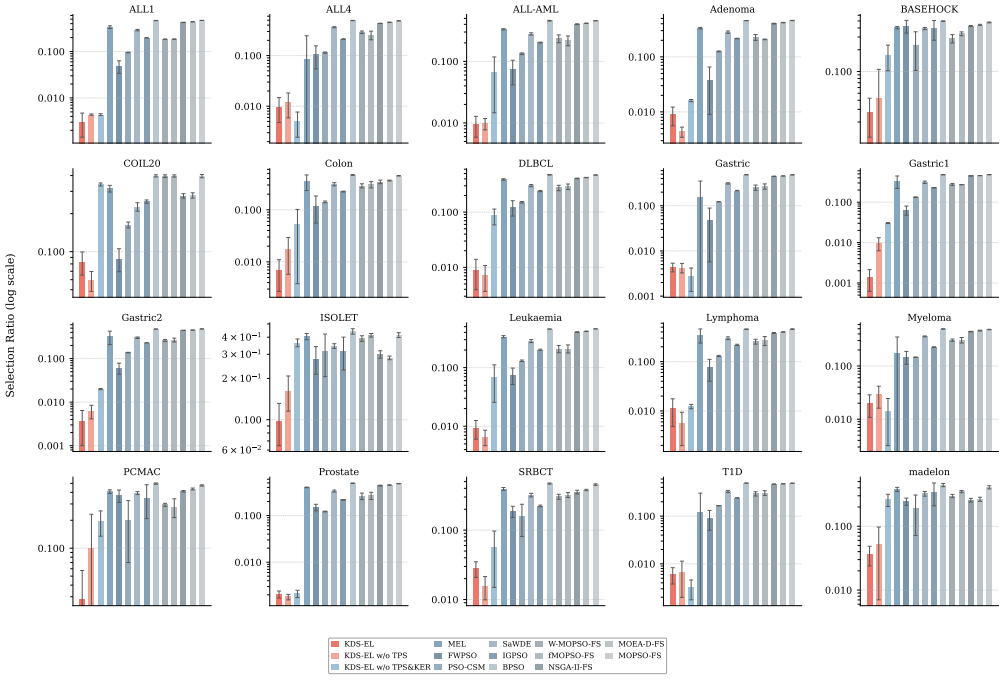


Figure 4: Selection ratio (mean over 10 runs) of the 14 fixed-length methods across the 20 benchmark datasets; the log-scale ordinate highlights the compactness of KDS-EL relative to all competitors.

5.3. *Feature-Selection Compactness*

KDS-EL selects, on average, 87.8 features across the 20-dataset benchmark, corresponding to an effective selection ratio of approximately 2–3%. For comparison, MEL retains 2,722 features, and all multi-objective methods (W-MOPSO-FS, fMOPSO-FS, NSGA-II-FS, MOEA-D-FS, MOPSO-FS) select between 2,960 and 5,263 features on average. Even the most compact single-objective competitor, FWPSO, selects 1,069 features—more than an order of magnitude above KDS-EL. Table 5 reports the per-dataset feature counts, and Fig. 4 reports the per-dataset selection ratio of all 14 methods on a log scale, underscoring the consistent compactness of KDS-EL relative to every competitor.

This combined finding—equal or higher accuracy with 20- to 50-fold compression—has a clear Pareto interpretation: in the accuracy-versus-features plane, KDS-EL occupies the extreme top-left region no competing method can dominate. The compression is most pronounced on gene-expression data: on ALL1, 0.9998 accuracy with 39 features (0.31% of 12,625) versus 4,248 for MEL; on Gastric1 (22,283 dimensions), 0.9800 with 31 features (0.14%) versus 6,044 for the most compact multi-objective method; on Colon, 14 features—a 140:1 compression—at the highest accuracy (0.9500); on Prostate, 26 features at 0.9701. Across all 20 datasets the feature count ascends monotonically along the ablation chain: 87.8 (KDS-EL) → 122 (KDS-EL w/o TPS) → 291 (KDS-EL w/o TPS&KER) → 2,722 (MEL). These figures illustrate KDS-EL's central empirical contribution: diagnostic-quality classification on high-dimensional biomedical data using a fraction of the original feature budget.

5.4. *Statistical Significance*

The ranking is underpinned by the Demšar [52] framework for classifier comparison over multiple datasets.

Pairwise Wilcoxon tests. One-sided Wilcoxon signed-rank tests ($p < 0.05$) in the direction of superiority confirm that KDS-EL significantly outperforms MEL, PSO-CSM, IGPSO, and BPSO on all 20 datasets (20/20); KDS-EL w/o TPS&KER, FWPSO, and SaWDE on 19/20; KDS-EL w/o TPS on 13/20; and the five multi-objective methods on 14–18 of 20. We adopt the one-sided formulation because each comparison tests the explicit claim that the row method outperforms the column method, following the Demšar [52] protocol for superiority claims. The lower counts against MOFS-class methods (14 vs. W-MOPSO-FS, 15 vs. fMOPSO-FS, 16 vs. NSGA-II-FS and MOEA-D-FS) do not imply weaker performance: they reflect ceiling-effect datasets where accuracy is saturated and statistical ties prevail, while KDS-EL decisively dominates on feature count, as discussed in Section 5.5. Table 6 reports the pairwise win counts, confirming that no competitor outperforms KDS-EL on a majority of datasets.

Table 4 provides the complementary per-dataset view: its +, ≈, and − symbols come from two-sided rank-sum tests within each dataset, whereas Table 6 pools one-sided signed-rank tests across datasets. Both protocols agree that no competitor holds a significant advantage over KDS-EL on a majority of datasets.

Holm-Bonferroni correction. Table 7 presents the Holm-Bonferroni step-down correction [54] applied to the 13 pairwise comparisons between KDS-EL and every other method. All 13 null hypotheses of no superiority are rejected at corrected significance levels; the one-sided Wilcoxon p-values are so small that both the Fisher combined p-values and the Holm-adjusted p-values are below 10^{-15} and are reported as such in Table 7. Even the most conservative comparison, against KDS-EL w/o TPS, rejects H_0 at $\alpha = 0.05$. These results confirm that removing any individual component yields a statistically significant degradation in accuracy, consistent with the component-wise analysis of Section 5.8.

Friedman ranking. The Friedman mean ranks [53] over the 20-dataset accuracy matrix place KDS-EL first (1.90), KDS-EL w/o TPS second (2.95), W-MOPSO-FS third (5.25), and KDS-EL w/o TPS&KER fourth (5.55); Fig. 5 reports the complete ranking. The complete method therefore holds the best Friedman rank, while the pairwise Wilcoxon and Holm-Bonferroni tests in Tables 6 and 7 establish significance over all 13 competitors.

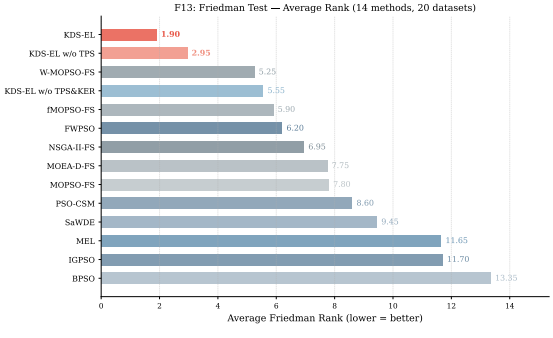


Figure 5: Average Friedman ranks of the 14 methods over the 20-dataset accuracy matrix (lower is better), computed from the same per-dataset mean accuracies used in Table 4.

5.5. *Ceiling-Effect Analysis on Gene-Expression Data*

Eight datasets in the benchmark—ALL1, ALL_AML, Adenoma, DLBCL, Gastric2, Leukaemia, Lymphoma, and SRBCT—exhibit mean accuracy at or

Table 5: Average number of selected features (over 10 runs) of 14 methods on 20 datasets. Bold indicates the minimum per dataset.

Dataset	KDS-EL	KDS-EL w/o TPS	KDS-EL w/o TPS&KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MOPSO-FS	fMOPSO-FS	NSGA-II-FS	MOEA-D-FS	MOPSO-FS
ALL1	**39**	55	55	4248	612	1206	3667	2501	5886	2333	2339	5368	5548	5920
ALL4	124	152	**64**	1082	1326	1447	4581	2677	6217	3645	3213	5461	5736	6126
ALL_AML	**67**	70	473	2369	521	949	1996	1444	3266	1678	1562	2872	2968	3266
Adenoma	66	**33**	120	2483	278	937	2102	1608	3399	1691	1551	3000	3127	3420
BASEHOCK	**134**	208	812	1965	2066	1113	1912	1894	2419	1390	1617	2073	2150	2319
COIL20	84	**60**	348	323	90	167	232	255	407	405	407	283	286	405
Colon	**14**	35	105	702	240	284	621	448	935	580	612	690	730	903
DLBCL	64	**52**	615	2784	876	1074	2155	1718	3379	1963	2067	2904	3010	3349
Gastric	100	97	**62**	3604	1070	2782	7072	4916	11004	5744	6082	10165	10342	10938
Gastric1	**31**	217	678	7393	1420	2966	6963	5074	10661	6148	6044	9954	10128	10706
Gastric2	**83**	140	444	7084	1362	3062	6708	5134	10627	5803	5956	9919	10062	10701
ISOLET	**60**	100	225	250	172	194	213	195	272	242	255	185	174	256
Leukaemia	66	**47**	488	2409	534	934	2026	1440	3246	1492	1503	2893	2973	3272
Lymphoma	45	**23**	50	1404	304	524	1204	880	1832	1027	1076	1536	1609	1834
Myeloma	250	365	**175**	2137	1863	1847	4543	2859	6259	3854	3866	5585	5839	6113
PCMAC	**91**	329	637	1339	1209	651	1292	1137	1635	966	912	1352	1431	1560
Prostate	26	**23**	27	5058	1898	1541	4229	2729	6211	3294	3384	5532	5676	6100
SRBCT	64	**36**	129	906	431	366	739	515	1083	700	742	818	874	1044
T1D	330	366	**174**	6715	4984	9047	18009	13170	26971	16096	16774	25649	25939	26830
madelon	**18**	26	130	190	122	95	162	171	221	149	174	128	132	204

Table 6: One-sided Wilcoxon signed-rank test results ($p < 0.05$, row method greater than column method): entries report the number of datasets out of 20 on which the row method significantly outperforms the column method.

Method	KDS-EL	KDS-EL w/o TPS	KDS-EL w/o TPS&KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MOPSO-FS	fMOPSO-FS	NSGA-II-FS	MOEA-D-FS	MOPSO-FS
KDS-EL	—	13/20	19/20	20/20	19/20	20/20	19/20	20/20	20/20	14/20	15/20	16/20	16/20	18/20
KDS-EL w/o TPS	**0/20**	—	15/20	20/20	20/20	20/20	19/20	20/20	20/20	14/20	15/20	15/20	15/20	17/20
KDS-EL w/o TPS&KER	**0/20**	0/20	—	18/20	11/20	15/20	15/20	19/20	16/20	10/20	9/20	14/20	14/20	13/20
MEL	**0/20**	0/20	0/20	—	0/20	1/20	3/20	3/20	11/20	2/20	2/20	3/20	3/20	3/20
FWPSO	**0/20**	0/20	0/20	0/20	—	16/20	15/20	20/20	20/20	6/20	7/20	9/20	10/20	9/20
PSO-CSM	**0/20**	0/20	0/20	0/20	0/20	—	12/20	20/20	16/20	2/20	2/20	5/20	5/20	6/20
SaWDE	**0/20**	0/20	0/20	0/20	0/20	0/20	—	10/20	20/20	2/20	2/20	3/20	3/20	6/20
IGPSO	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	—	14/20	0/20	0/20	2/20	1/20	2/20
BPSO	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	0/20	0/20	0/20	0/20	0/20
W-MOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	2/20	10/20	12/20	12/20
fMOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	10/20	10/20	8/20
NSGA-II-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	6/20	5/20
MOEA-D-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	4/20
MOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—

Table 7: Holm–Bonferroni corrected pairwise comparison between KDS-EL and each of the 13 comparison methods over the 20 datasets. ΔAcc is the mean accuracy advantage of KDS-EL; Combined p is the unadjusted Fisher combined one-sided Wilcoxon p-value and Holm p is its Holm–Bonferroni adjustment; H_0 of no superiority is rejected at the Holm threshold in every case.

Method	ΔAcc	Combined p	Holm p	Reject H_0
KDS-EL w/o TPS	+0.0086	$< 10^{-15}$	$< 10^{-15}$	✓
KDS-EL w/o TPS&KER	+0.0248	$< 10^{-15}$	$< 10^{-15}$	✓
MEL	+0.0706	$< 10^{-15}$	$< 10^{-15}$	✓
FWPSO	+0.0396	$< 10^{-15}$	$< 10^{-15}$	✓
PSO-CSM	+0.0552	$< 10^{-15}$	$< 10^{-15}$	✓
SaWDE	+0.0630	$< 10^{-15}$	$< 10^{-15}$	✓
IGPSO	+0.0734	$< 10^{-15}$	$< 10^{-15}$	✓
BPSO	+0.0875	$< 10^{-15}$	$< 10^{-15}$	✓
W-MOPSO-FS	+0.0442	$< 10^{-15}$	$< 10^{-15}$	✓
fMOPSO-FS	+0.0464	$< 10^{-15}$	$< 10^{-15}$	✓
NSGA-II-FS	+0.0509	$< 10^{-15}$	$< 10^{-15}$	✓
MOEA-D-FS	+0.0532	$< 10^{-15}$	$< 10^{-15}$	✓
MOPSO-FS	+0.0565	$< 10^{-15}$	$< 10^{-15}$	✓

above 0.998 across the top-performing methods. On these datasets, the accuracy differences among the leading methods are typically narrower than 0.002, which corresponds to at most one or two misclassified instances across ten independent five-fold cross-validation runs. At this saturation level, classification accuracy ceases to be a meaningful discriminator; instead, the number of features selected becomes the primary figure of merit.

KDS-EL selects between 39 (ALL1) and 83 (Gastric2) features across these eight datasets. In contrast, the best single-objective competitor, FWPSO, retains 278–1,362 features, and the multi-objective methods select between 700 and 10,701 features on the same data. This gap carries practical significance for gene-expression analysis: a compact signature of 40–80 genes is amenable to targeted biological validation, pathway enrichment analysis, and potential clinical deployment, whereas signatures spanning several hundred to several thousand genes offer little interpretability and are vulnerable to overfitting in the $p \gg n$ regime.

That KDS-EL ties or marginally trails methods achieving perfect 1.0 on ALL1 and SRBCT does not indicate inferior discrimination: all methods misclassify fewer than two samples on average, differing only in which samples are misclassified—stochastic variation expected in small-sample data. A similar pattern holds on Lymphoma (45 samples), where KDS-EL and KDS-EL w/o TPS share the best accuracy of 0.9994 with a difference of zero, yet KDS-EL achieves this with 45 features versus 1,027 for W-MOPSO-FS. We therefore interpret the subset of ceiling datasets as evidence that KDS-EL attains the performance frontier while employing a substantially leaner feature representation.

5.6. Convergence Behavior

Fig. 6 shows the mean convergence trajectories of KDS-EL, KDS-EL w/o TPS, KDS-EL w/o TPS&KER, and MEL averaged over the 20 datasets. KDS-EL enters the high-accuracy region within the first 10–15 iterations thanks to MI-guided initialization, and KER then drives a prolonged refinement that lifts accuracy past iteration 30 where competitors have plateaued. The Two-Phase schedule yields a characteristic "fast-drop → plateau → micro-rise" trajectory: after aggressive feature reduction under Phase 1 weights, accuracy plateaus until the Phase 2 switch triggers a final upward correction, whereas KDS-EL w/o TPS plateaus earlier at lower terminal accuracy. MEL, initialized randomly, needs 40–60 iterations to exit the low-accuracy regime; MOFS-class methods stagnate late because Pareto-front maintenance retains a diverse archive that dilutes search pressure, flattening their curves after iteration 60.

5.7. Computational Efficiency

KDS-EL records a mean runtime of 51.4 seconds, ranking second among all methods and trailing FWPSO by only 1.7 seconds (49.7 s). Relative to MEL (175.0 s), KDS-EL achieves a 3.4× speedup. All five multi-objective methods exceed 110 seconds on average—more than double the runtime of KDS-EL—owing to the overhead of non-dominated sorting and Pareto-archive maintenance at each generation. KDS-EL thereby occupies the high-accuracy, low-runtime corner of the accuracy–runtime trade-off; the complete per-dataset runtime table is provided in the Supplementary Material.

The efficiency advantage is most pronounced at high dimensionality: on T1D (54,675 features) KDS-EL finishes in 118.8 s versus 256.6 s for PSO-CSM (2.2×), and on BASEHOCK it is 4.7× faster than SaWDE (130.1 s vs. 616.0 s). Three factors explain this: MI-guided initialization shortens the path to a high-quality region; KER refines elites within a bounded refinement budget that does not alter the dominant per-iteration complexity; and Two-Phase scheduling concentrates feature reduction early so search later operates on compact, cheaper subsets. FWPSO's marginal 1.7 s lead stems from its SA-driven local search, which terminates early on small-sample problems at the expense of a 3.97-percentage-point accuracy deficit.

5.8. Component Contribution Analysis

5.8.1. Aggregate Contributions

The ablation chain follows a remove-in-sequence paradigm and reads subtractively: each row of Table 8 strips one mechanism from the complete KDS-EL (0.9574 accuracy, 87.8 features on average), and the accuracy loss at each step quantifies the removed component's contribution. These four configurations are exactly the methods compared in Section 5: KDS-EL w/o TPS is KDS-EL without Two-Phase scheduling, KDS-EL w/o TPS&KER is KDS-EL without both the KER strategy and Two-Phase scheduling, and MEL [12] is the framework stripped of all three mechanisms—the residual that remains once every component has been removed.

Removing MI-guided initialization and adaptive re-initialization inflicts the largest single accuracy loss: -4.58pp, dropping the mean accuracy from 0.9325 to 0.8868 while the feature count expands from 291 to 2,722, roughly an order of magnitude. This dominant role stems from informed starting positions and stagnation-triggered re-initialization: the MI-informed half provides a high-quality search origin, while adaptive re-initialization counters premature convergence.

Removing the KER strategy costs a further -1.63pp, dropping the mean accuracy from 0.9488 to 0.9325 and expanding the feature set from 122 to 291. Because the accumulated importance matrix is derived from fitness comparisons already performed during search, KER's contribution is obtained at a bounded refinement cost that does not alter the dominant per-iteration complexity; this gain is also achieved on top of the importance-matrix knowledge already present in MEL, confirming that KER re-deploys rather than duplicates it.

Removing the Two-Phase Weight Scheduling strategy from the complete KDS-EL costs -0.86pp, dropping the mean accuracy from 0.9574 to 0.9488 and expanding the feature set from 87.8 to 122. Runtime also rises from 51.4 s to 53.6 s during this removal.

Each component contributes along a distinct dimension of the search process, and its removal damages that dimension. MI-guided initialization operates in the start-and-reduction dimension, providing the bulk of feature reduction alongside an informed search origin. The KER strategy occupies the accuracy dimension, refining elite solutions through cross-iteration accumulated knowledge at a bounded refinement cost. The Two-Phase strategy delivers simultaneous gains across all three axes—accuracy, sparsity, and efficiency—by aligning objective weights with the structural phases of the search. The mechanisms are complementary and independent—none subsumes another, and each is verifiable in isolation.

5.8.2. Per-Dataset Component Contribution

Representative rows of Table 8 reveal two characteristic patterns across the 20 datasets.

On the eight ceiling datasets—where classification accuracy saturates above 0.998—the MI and KER components rapidly approach the accuracy upper bound, after which the ablative configurations differentiate through feature counts alone. On ALL1, KDS-EL attains 0.9998 accuracy with 39 features; KDS-EL w/o TPS holds 55 features at 0.9996; removing KER in addition leaves KDS-EL w/o TPS&KER, also at 55 features; and MEL selects 4248 features at 0.9664. On Adenoma, MEL selects 2483 features at 0.9410 accuracy, KDS-EL w/o TPS&KER compresses to 120 at 0.9958, KDS-EL w/o TPS to 33, and the complete KDS-EL to 66, with the latter two tying at the saturated accuracy of 0.9996. DLBCL, Leukaemia, Lymphoma, and SRBCT follow the same trajectory: accuracy saturates early, and the configurations differentiate through increasingly compact feature subsets with no accuracy penalty.

On medium-difficulty datasets, accuracy degrades in a stair-step progression as mechanisms are removed. On BASEHOCK, the accuracy sequence reads 0.9113 (KDS-EL), 0.8997 (KDS-EL w/o TPS), 0.8597 (KDS-EL w/o TPS&KER), and 0.8518 (MEL). On madelon, a comparable trajectory emerges: 0.8822, 0.8659, 0.7679, and 0.7521. The sharpest feature inflation uniformly occurs when MI-guided initialization is removed (the KDS-EL w/o TPS&KER-to-MEL step), where feature counts expand by an order of magnitude on most high-dimensional gene-expression datasets. On Gastric1, features rise from 678 to 7,393; on T1D, from 174 to 6,715; on Prostate, from 27 to 5,058. These patterns confirm that the initial investment in MI-informed initialization delivers both the dominant accuracy contribution and the foundational feature reduction on which subsequent mechanisms build.

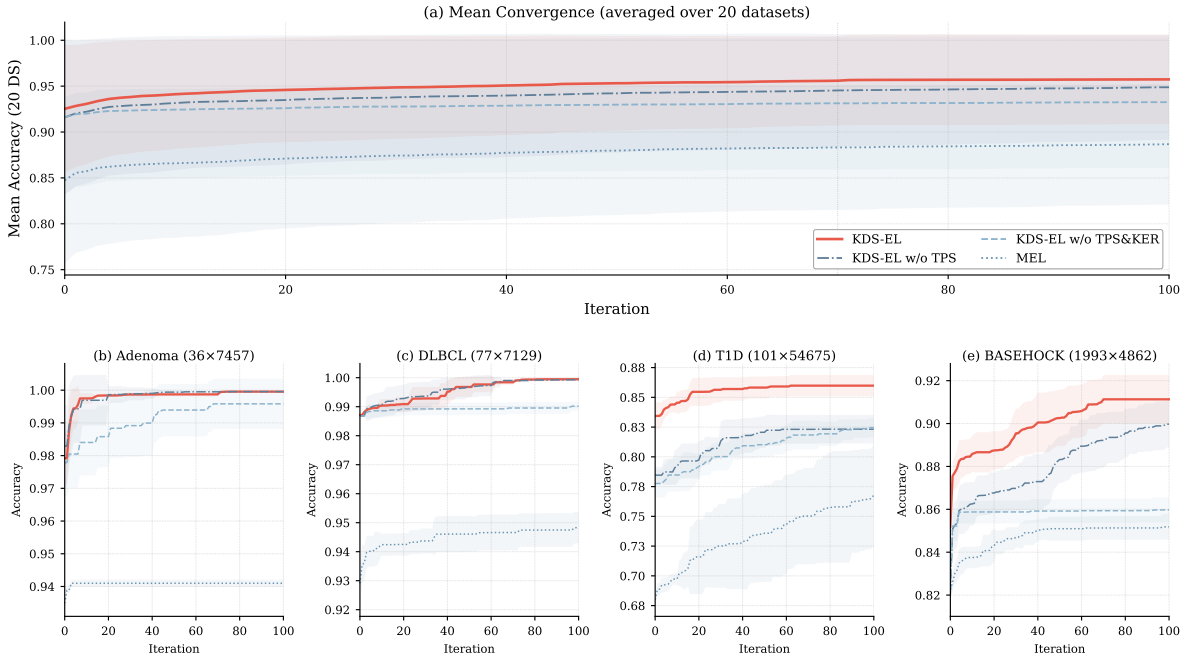


Figure 6: Mean convergence trajectories of KDS-EL, KDS-EL w/o TPS, KDS-EL w/o TPS&KER, and MEL averaged over the 20 benchmark datasets.

5.8.3. *Two-Phase Scheduling and In-place Update Verification*

The Two-Phase Weight Scheduling strategy is isolated by comparing the complete KDS-EL against its fixed-weight counterpart (fixed ws=[0.9, 0.1] with in-place update) on the same 20-dataset batch. Removing the schedule lowers accuracy from 0.9574 to 0.9488 (-0.86pp), expands average features from 87.8 to 122, and increases runtime from 51.4 s to 53.6 s. Following the Phase-2 weight switch in the full method, accuracy continues to rise while features undergo further compression, confirming that the coarse-then-fine schedule outperforms a fixed-weight regime. On a per-dataset basis, Table 4 shows that KDS-EL is not inferior to the fixed-weight configuration on the overwhelming majority of datasets.

Table 8: Component contribution analysis: each row strips one mechanism from the complete KDS-EL, and ΔAcc reports the resulting accuracy loss relative to the row above. "MEL" denotes the framework stripped of all three mechanisms—the residual that remains once every component has been removed, corresponding to MEL [12]. Feature counts are dataset means.

Variant	Accuracy	Features	ΔAcc (pp)
KDS-EL	**0.9574**	**87.8**	–
KDS-EL w/o TPS	0.9488	122	-0.86
KDS-EL w/o TPS&KER	0.9325	291	-1.63
MEL	0.8868	2722	-4.58

5.9. *Parameter Sensitivity Analysis*

To assess the sensitivity of KDS-EL to its key hyper-parameters, we conduct a one-factor-at-a-time analysis on four representative datasets: Adenoma (36 samples, 7,457 features), Colon (62, 2,000), ALL4 (93, 12,625), and BASEHOCK (1,993, 4,862). The subset spans small-sample gene expression and larger-sample text domains, with 2,000–12,625 features. Five parameters—WeightDecay, PhaseSplit, flip_frac, MI_init_ratio, and elite_k—are varied individually while the remaining four stay fixed at their defaults listed in Table 3. For each configuration, the mean accuracy, feature count, and runtime across the four datasets are recorded; Table 9 summarizes the tested ranges, optimal values, stable bands, and maximum accuracy deviations, and the complete tables are provided in the Supplementary Material. These four-dataset aggregates serve only to isolate single-parameter effects and are therefore not directly comparable with the 20-dataset global means reported in Section 5. Across all five parameters, the maximum accuracy deviation from the default configuration never exceeds ±1.5 pp.

Weight decay (WeightDecay). This parameter controls the discounting rate of historical feature importance scores. At the default 0.99, the four-dataset mean is 0.944 accuracy, 107 features, and 51.3 s. Reducing the decay to 0.90 erodes accuracy to 0.932 (−1.2 pp), as the importance matrix forgets useful patterns too rapidly. Disabling decay (1.00) inflates the feature count by approximately 40% (to 151) with no accuracy benefit—stale scores accumulate without bound. Performance plateaus across [0.95, 0.995], with accuracy fluctuation under ±1.2 pp.

Phase split ratio (PhaseSplit). This parameter governs the proportion of iterations allocated to aggressive feature reduction before accuracy-oriented refinement. The default 0.70 gives 70% of iterations to Phase 1. Values below 0.50 yield insufficient reduction time, leaving elevated feature counts (146–189), while 0.90 starves the refinement phase and reduces accuracy by 0.7 pp. The stable band [0.60, 0.80] keeps variation below 0.5 pp, confirming the schedule tolerates imprecise split points.

Elite flip fraction (flip_frac). At the default 0.01, roughly one feature per 100 dimensions is flipped per elite. Reducing this to 0.005 slows convergence (runtime +3.5 s). Increasing it beyond 0.02 injects excessive randomness: at 0.10, accuracy drops by 1.3 pp and features nearly double (to 203), indicating random perturbation overwhelms knowledge-guided refinement. Only the narrow window [0.005, 0.02] preserves the intended knowledge–randomness balance.

MI initialization ratio (MI_init_ratio). This is the most influential parameter. At the default 0.50, half the population is seeded with MI-ranked features. Eliminating MI guidance (0.00) imposes a cold start—accuracy falls by 1.2 pp and runtime rises by 21% (to 62.3 s) as the search rediscovers informative features from scratch. Full MI initialization (1.00) causes premature convergence (−0.8 pp) through reduced diversity. The stable band [0.25, 0.75] keeps accuracy within 0.4 pp of the optimum.

Elite count (elite_k). With $N = 20$, the default of three elites represents 15% of the swarm. A single elite restricts knowledge propagation (0.938 accuracy), while 10 elites dilute the quality signal of the KER strategy (0.935 accuracy, 141 features). The stable range [2, 5] maintains accuracy within 0.3 pp of the optimum, confirming that a moderate elite subset concentrates computational effort on the most promising candidates without dilution.

The defaults therefore transfer directly to new datasets without per-dataset tuning; practitioners should preserve MI_init_ratio and flip_frac, which exhibit the largest individual sensitivity and jointly govern the balance between injected prior knowledge and exploratory diversity.

6. Conclusion

This paper presented KDS-EL, a closed-loop knowledge-guided evolutionary framework addressing three structural limitations of general-purpose PSO in high-dimensional feature selection: random initialization provides insufficient guidance in sparse high-dimensional spaces, uniform mutation discards cross-iteration feature-importance signals, and static objective preferences cannot adapt to different evolutionary phases. MI-guided initialization

Table 9: One-factor-at-a-time parameter sensitivity analysis over four representative datasets (Adenoma, Colon, ALL4, BASEHOCK). The maximum accuracy deviation from the default configuration never exceeds ±1.5 pp.

Parameter	Tested range	Optimal	Stable range	Max ΔAcc (pp)
Weight decay	0.90–1.00	0.99	[0.95, 0.995]	±1.2
Phase split	0.3–0.9	0.70	[0.60, 0.80]	±0.7
Flip fraction	0.005–0.10	0.01	[0.005, 0.02]	±1.3
MI init ratio	0.0–1.0	0.50	[0.25, 0.75]	±1.2
Elite count	1–10	3	[2, 5]	±0.9

and adaptive re-initialization supply informed starting points and stagnation escape; removing them costs the largest single loss of 4.58pp and lets the feature set expand from 291 to 2,722. The KER strategy accumulates cross-iteration feature-importance scores from fitness comparisons already performed during search and conducts knowledge-guided elite flipping with a bounded refinement budget; removing it costs a further 1.63pp and expands the feature set from 122 to 291. Two-Phase Weight Scheduling aligns objective weights with search phases—aggressive feature reduction followed by accuracy refinement; removing it costs 0.86pp and expands the feature set from 87.8 to 122. On 20 high-dimensional datasets spanning gene expression, text, image, and speech domains, KDS-EL ranks first among 14 fixed-length methods, surpassing the strongest single-objective and multi-objective competitors by +3.97pp and +4.42pp. Wilcoxon tests confirm superiority over MEL in 20/20 datasets; Holm-Bonferroni correction establishes significance against all 13 competitors. KDS-EL delivers a 3.4× speedup over MEL (51.4 s vs. 175.0 s) with parameter sensitivity within ±1.5pp.

Several limitations suggest future directions. MI-based feature ranking incurs $O(D \log D)$ preprocessing cost, a bottleneck when D exceeds 10^5, where GPU-accelerated MI estimation may be required. The method currently targets single-label classification; extensions to multi-label learning, streaming feature selection, class-imbalanced data, and deep embedding representations remain open. Promising directions include GPU-accelerated MI approximation, multi-label and imbalanced extensions, transformer-based feature representations, and cross-dataset transfer via multi-task evolutionary extensions of the proposed framework.

CRediT Author Statement

Anonymous Author: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review & editing, Visualization.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data Availability

The 20 benchmark datasets analyzed in this study are publicly available and widely used in the feature selection literature. The source code of KDS-EL will be made available upon publication. Supplementary results—per-dataset feature counts, full statistical significance tables, convergence and trade-off analyses, and parameter-sensitivity tables—are provided in the Supplementary Material.

References

[1] J. Li, K. Cheng, S. Wang, F. Morstatter, R. P. Trevino, J. Tang, H. Liu, Feature selection: a data perspective, ACM Computing Surveys 50 (6) (2017) 1–45. doi:10.1145/3136625.

[2] I. Guyon, A. Elisseeff, An introduction to variable and feature selection, Journal of Machine Learning Research 3 (2003) 1157–1182.

[3] J. Liang, Z. Hu, Y. Bi, H. Cheng, K. Yu, C. Yue, X. Wang, W.-F. Guo, A survey on evolutionary computation for identifying biomarkers of complex disease, IEEE Transactions on Evolutionary Computation 29 (4) (2025) 1400–1418. doi:10.1109/TEVC.2024.3414442.

[4] R. Kohavi, G. H. John, Wrappers for feature subset selection, Artificial Intelligence 97 (1–2) (1997) 273–324. doi:10.1016/S0004-3702(97)00043-X.

[5] J. Kennedy, R. Eberhart, Particle swarm optimization, in: Proceedings of ICNN'95–International Conference on Neural Networks, Vol. 4, IEEE, 1995, pp. 1942–1948. doi:10.1109/ICNN.1995.488968.

[6] Y. Shi, R. Eberhart, A modified particle swarm optimizer, in: 1998 IEEE International Conference on Evolutionary Computation Proceedings (IEEE World Congress on Computational Intelligence), IEEE, 1998, pp. 69–73. doi:10.1109/ICEC.1998.699146.

[7] J. Kennedy, R. C. Eberhart, A discrete binary version of the particle swarm algorithm, Proceedings of the 1997 IEEE International Conference on Systems, Man, and Cybernetics 5 (1997) 4104–4108. doi:10.1109/ICSMC.1997.637339.

[8] X. Wang, W. Jia, A feature weighting particle swarm optimization method to identify biomarker genes, in: 2022 IEEE International Conference on Bioinformatics and Biomedicine (BIBM), IEEE, 2022, pp. 830–834. doi:10.1109/BIBM55620.2022.9995376.

[9] X. Wang, Y. Wang, K.-C. Wong, X. Li, A self-adaptive weighted differential evolution approach for large-scale feature selection, Knowledge-Based Systems 235 (2022) 107633. doi:10.1016/j.knosys.2021.107633.

[10] Y. Xue, C. Zhang, A novel importance-guided particle swarm optimization based on mlp for solving large-scale feature selection problems, Swarm and Evolutionary Computation 91 (2024) 101760. doi:10.1016/j.swevo.2024.101760.

[11] B. Wei, S. Yang, W. Zha, L. Deng, J. Huang, X. Su, F. Wang, Particle swarm optimization algorithm based on comprehensive scoring framework for high-dimensional feature selection, Swarm and Evolutionary Computation 95 (2025) 101915. doi:10.1016/j.swevo.2025.101915.

[12] X. Wang, H. Shangguan, F. Huang, S. Wu, W. Jia, Mel: Efficient multi-task evolutionary learning for high-dimensional feature selection, IEEE Transactions on Knowledge and Data Engineering 36 (8) (2024) 4020–4033. doi:10.1109/TKDE.2024.3366333.

[13] M. A. Kramer, Nonlinear principal component analysis using autoassociative neural networks, AIChE Journal 37 (2) (1991) 233–243.

[14] G. Roffo, S. Melzi, Features selection via eigenvector centrality, in: Proceedings of New Frontiers in Mining Complex Patterns, 2016.

[15] H. Peng, F. Long, C. Ding, Feature selection based on mutual information: criteria of max-dependency, max-relevance, and min-redundancy, IEEE Transactions on Pattern Analysis and Machine Intelligence 27 (8) (2005) 1226–1238. doi:10.1109/TPAMI.2005.159.

[16] L. Yu, H. Liu, Feature selection for high-dimensional data: A fast correlation-based filter solution, in: Proceedings of the 20th International Conference on Machine Learning (ICML-03), 2003, pp. 856–863.

[17] I.-S. Oh, J.-S. Lee, B.-R. Moon, Hybrid genetic algorithms for feature selection, IEEE Transactions on Pattern Analysis and Machine Intelligence 26 (11) (2004) 1424–1437. doi:10.1109/TPAMI.2004.105.

[18] F. R. Bach, Bolasso: Model consistent Lasso estimation through the bootstrap, in: Proceedings of the 25th International Conference on Machine Learning, 2008, pp. 33–40.

[19] J. Zhang, D. Xu, K. Hao, Y. Zhang, W. Chen, J. Liu, R. Gao, C. Wu, Y. D. Marinis, FS-GBDT: Identification multicancer-risk module via a feature selection algorithm by integrating Fisher score and GBDT, Briefings in Bioinformatics 22 (3) (2021) bbaa189. doi:10.1093/bib/bbaa189.

[20] K. Yu, W. Xie, L. Wang, W. Li, ILRC: A hybrid biomarker discovery algorithm based on improved l_1 regularization and clustering in microarray data, BMC Bioinformatics 22 (1) (2021) 1–19. doi:10.1186/s12859-021-04223-3.

[21] M. H. Aghdam, N. Ghasem-Aghaee, M. E. Basiri, Text feature selection using ant colony optimization, Expert Systems with Applications 36 (3) (2009) 6843–6853.

[22] D. Karaboga, B. Basturk, A powerful and efficient algorithm for numerical function optimization: Artificial bee colony (ABC) algorithm, Journal of Global Optimization 39 (3) (2007) 459–471.

[23] C.-L. Huang, C.-J. Wang, A GA-based feature selection and parameters optimization for support vector machines, Expert Systems with Applications 31 (2) (2006) 231–240.

[24] Y. Xue, B. Xue, M. Zhang, Self-adaptive particle swarm optimization for large-scale feature selection in classification, ACM Transactions on Knowledge Discovery from Data 13 (5) (2019) 1–27. doi:10.1145/3340848.

[25] Y. Prasad, K. Biswas, Gene selection in microarray datasets using progressively refined PSO scheme, in: Proceedings of the AAAI Conference on Artificial Intelligence, Vol. 29, 2015.

[26] B. Wei, J. Huang, L. Deng, S. Yang, J. Zheng, Y. Huang, Reinforcement learning-based particle swarm optimization with adaptive scoring mechanism for high-dimensional feature selection, Swarm and Evolutionary Computation (2025) 102104doi:10.1016/j.swevo.2025.102104.

[27] K. Chen, K. Liu, P. Wang, J. Liang, Z. Yang, K. Yu, An efficient evolutionary feature selection algorithm with divide-and-conquer strategy for classification, Expert Systems with Applications 302 (2025) 130416. doi:10.1016/j.eswa.2025.130416.

[28] T. M. Cover, J. A. Thomas, Elements of Information Theory, John Wiley & Sons, New York, NY, USA, 1991.

[29] F. Yu, J. Guan, H. Wu, H. Wang, B. Ma, Multi-population differential evolution approach for feature selection with mutual information ranking, Expert Systems with Applications 260 (2025) 125404. doi:10.1016/j.eswa.2024.125404.

[30] B. Xue, M. Zhang, W. N. Browne, Particle swarm optimization for feature selection in classification: A multi-objective approach, IEEE Transactions on Cybernetics 43 (6) (2012) 1656–1671.

[31] Y. Hu, Y. Zhang, D. Gong, Multiobjective particle swarm optimization for feature selection with fuzzy cost, IEEE Transactions on Cybernetics 51 (2) (2020) 874–888. doi:10.1109/TCYB.2019.2946898.

[32] K. Deb, A. Pratap, S. Agarwal, T. Meyarivan, A fast and elitist multiobjective genetic algorithm: NSGA-II, IEEE Transactions on Evolutionary Computation 6 (2) (2002) 182–197. doi:10.1109/4235.996017.

[33] Q. Zhang, H. Li, MOEA/D: A multiobjective evolutionary algorithm based on decomposition, IEEE Transactions on Evolutionary Computation 11 (6) (2007) 712–731. doi:10.1109/TEVC.2007.892759.

[34] C. A. C. Coello, G. T. Pulido, M. S. Lechuga, Handling multiple objectives with particle swarm optimization, IEEE Transactions on Evolutionary Computation 8 (3) (2004) 256–279. doi:10.1109/TEVC.2004.826067.

[35] E. Hancer, B. Xue, M. Zhang, An evolutionary filter approach to feature selection in classification for both single- and multi-objective scenarios, Knowledge-Based Systems 280 (2023) 111008. doi:10.1016/j.knosys.2023.111008.

[36] S. Liu, H. Wang, W. Peng, W. Yao, Surrogate-assisted evolutionary algorithms for expensive combinatorial optimization: a survey, Complex & Intelligent Systems 10 (4) (2024) 5933–5949. doi:10.1007/s40747-024-01465-5.

[37] Y. Zhang, Q. Yang, A survey on multi-task learning, IEEE Transactions on Knowledge and Data Engineering 34 (12) (2021) 5586–5609. doi:10.1109/TKDE.2021.3070203.

[38] N. Zhang, A. Gupta, Z. Chen, Y.-S. Ong, Evolutionary machine learning with minions: A case study in feature selection, IEEE Transactions on Evolutionary Computation 26 (1) (2021) 130–144.

[39] K. Chen, B. Xue, M. Zhang, F. Zhou, An evolutionary multitasking-based feature selection method for high-dimensional classification, IEEE Transactions on Cybernetics 52 (7) (2020) 7172–7186. doi:10.1109/TCYB.2020.3042243.

[40] K. Chen, B. Xue, M. Zhang, F. Zhou, Evolutionary multitasking for feature selection in high-dimensional classification via particle swarm optimization, IEEE Transactions on Evolutionary Computation 26 (3) (2022) 446–460. doi:10.1109/TEVC.2021.3100056.

[41] C.-L. He, Y. Zhang, D. Gong, X. Song, X. Sun, A multitask bee colony band selection algorithm with variable-size clustering for hyperspectral images, IEEE Transactions on Evolutionary Computation 26 (6) (2022) 1566–1580. doi:10.1109/TEVC.2022.3159253.

[42] S. Liu, Q. Lin, L. Feng, K.-C. Wong, K. C. Tan, Evolutionary multitasking for large-scale multiobjective optimization, IEEE Transactions on Evolutionary Computation 27 (4) (2022) 863–877. doi:10.1109/TEVC.2022.3166482.

[43] M. L. Bermingham, R. Pong-Wong, A. Spiliopoulou, C. Hayward, I. Rudan, H. Campbell, A. F. Wright, J. F. Wilson, F. Agakov, P. Navarro, C. S. Haley, Application of high-dimensional feature selection: Evaluation for genomic prediction in man, Scientific Reports 5 (1) (2015) 1–12.

[44] T. M. Phuong, Z. Lin, R. B. Altman, Choosing SNPs using feature selection, in: 2005 IEEE Computational Systems Bioinformatics Conference (CSB'05), IEEE, 2005, pp. 301–309.

[45] E. Saghapour, S. Kermani, M. Sehhati, A novel feature ranking method for prediction of cancer stages using proteomics data, PLoS ONE 12 (9) (2017) e0184203. doi:10.1371/journal.pone.0184203.

[46] Y. Zhang, Z. Dong, P. Phillips, S. Wang, G. Ji, J. Yang, T. F. Yuan, Detection of subjects and brain regions related to Alzheimer's disease using 3D MRI scans based on eigenbrain and machine learning, Frontiers in Computational Neuroscience 9 (2015) 66. doi:10.3389/fncom.2015.00066.

[47] J. Lin, C. He, H. Jiang, Y. Huang, Y. Jin, Surrogate-assisted multiobjective gene selection for cell classification from large-scale single-cell RNA sequencing data, IEEE Transactions on Evolutionary Computation 29 (2025) 601–615. doi:10.1109/TEVC.2025.3533490.

[48] M. Xia, K. Li, J. Wan, Z. Huang, F. Cheng, An instance selection assisted evolutionary method for high-dimensional feature selection, IEEE Transactions on Evolutionary Computation 30 (3) (2026) 1225–1239. doi:10.1109/TEVC.2025.3600815.

[49] S. Deng, H. Liu, W. Zhu, K. Li, M. Li, Elite neighborhood search and dynamic feature probability-guided crow search algorithm for high-dimensional feature selection, Expert Systems with Applications 274 (2025) 127012. doi:10.1016/j.eswa.2025.127012.

[50] Z. Ye, Y. Zeng, T. Cai, J. Shen, W. Zhou, Q. He, M. Mei, Y. Cen, KMHBO: A knowledge-guided multi-niche hybrid breeding optimization algorithm for high-dimensional multimodal feature selection, Expert Systems with Applications 298 (2026) 129911. doi:10.1016/j.eswa.2025.129911.

[51] F. Han, F. Li, Q. Ling, H. Han, T. Lu, Z. Jiao, H. Zhang, A feature selection method based on feature-label correlation information and self-adaptive MOPSO, Neural Processing Letters 56 (2024) 110. doi:10.1007/s11063-024-11553-9.

[52] J. Demšar, Statistical comparisons of classifiers over multiple data sets, Journal of Machine Learning Research 7 (2006) 1–30.

[53] M. Friedman, A comparison of alternative tests of significance for the problem of m rankings, The Annals of Mathematical Statistics 11 (1) (1940) 86–92. doi:10.1214/aoms/1177731944.

[54] S. Holm, A simple sequentially rejective multiple test procedure, Scandinavian Journal of Statistics 6 (2) (1979) 65–70.