Highlights

Multi-Task Particle Swarm Optimization with Knowledge-Guided Elite Refinement and Two-Phase Weight Scheduling for High-Dimensional Feature Selection

Anonymous Author

- Closed-loop knowledge-guided swarm learning framework for high-dimensional feature selection.
- MI initialization with adaptive re-initialization restores diversity and prevents search entrapment.
- Knowledge-guided elite refinement accumulates feature-importance signals with exponential decay.
- Two-phase weight scheduling temporally decouples feature compression from accuracy refinement.
- KDS-EL achieves 95.74% accuracy with 87.8 selected features across 20 benchmark datasets.

Multi-Task Particle Swarm Optimization with Knowledge-Guided Elite Refinement and Two-Phase Weight Scheduling for High-Dimensional Feature Selection

Anonymous Author[a,*]

[a] *Anonymous Institution (masked for double-blind review)*

ARTICLE INFO

Keywords:
Feature selection
Particle swarm optimization
Mutual information
Knowledge-guided elite refinement
Two-phase weight scheduling

ABSTRACT

Feature selection in high-dimensional data remains a central challenge for evolutionary computation. Particle swarm optimization suffers from three structural deficiencies: random initialization provides insufficient guidance in sparse high-dimensional spaces, uniform feature perturbation fails to exploit cross-iteration feature-importance signals, and static objective preferences cannot adapt to different evolutionary phases. To address these limitations, we propose KDS-EL (Knowledge-Guided Dual-Phase Swarm Evolutionary Learning), a knowledge-guided evolutionary framework with a closed-loop feature-importance learning mechanism. First, to overcome inefficient initialization, we propose mutual-information-guided initialization with adaptive re-initialization, injecting prior knowledge and recovering population diversity when stagnation occurs. Second, to exploit discarded cross-iteration feature-importance signals, we propose a knowledge-guided elite refinement strategy that accumulates feature-importance knowledge and steers elite solutions toward historically effective features. Third, to align objective weights with search phases, we propose two-phase weight scheduling that emphasizes early dimensionality reduction and late-stage accuracy refinement. On 20 high-dimensional benchmarks, KDS-EL achieves the best classification performance among all competitors, with a mean accuracy of 0.9574 and only 87.8 selected features, outperforming the strongest single- and multi-objective competitors by +3.97 and +4.42 percentage points; Wilcoxon signed-rank tests with Holm–Bonferroni correction confirm statistically significant improvements over all 13 competitors.

1. Introduction

High-dimensional data have emerged as a fundamental challenge in classification problems across diverse domains, including genomics, text mining, and computer vision. These datasets typically contain hundreds or even tens of thousands of features, while only a limited number of labeled samples are available, resulting in the well-known $p \gg n$ scenario. Under such conditions, conventional learning algorithms often suffer from severe overfitting, increased computational complexity, and reduced interpretability due to the curse of dimensionality [1, 2]. Feature selection (FS) has therefore become an essential preprocessing technique that aims to identify a compact and informative subset of features while preserving classification performance. However, selecting an optimal feature subset remains highly challenging because the search space grows exponentially with the number of features D, and the underlying subset-selection problem is NP-hard [3]. Moreover, individual feature relevance measures are insufficient to characterize complex feature interactions and dependencies, which are often critical for achieving robust classification performance.

FS methods are generally categorized into filter, wrapper, and embedded approaches [1, 4]. Filter methods rank features according to statistical criteria and offer high computational efficiency, but they usually evaluate features independently of the classifier and may overlook feature interactions relevant to classification. Embedded methods incorporate feature selection into the learning process, providing a compromise between selection efficiency and predictive performance; however, their effectiveness is often dependent on the characteristics of specific learning models. Wrapper methods directly optimize feature subsets according to classifier performance and can achieve competitive classification results, but the repeated model evaluations make them computationally expensive, especially in high-dimensional search spaces. To address the combinatorial complexity of wrapper-based FS, evolutionary computation has been extensively investigated as a promising optimization paradigm [3, 5–7]. By maintaining a population of candidate solutions, evolutionary algorithms can perform global exploration of feature subsets without requiring gradient information. Among them, particle swarm optimization (PSO), genetic algorithms, and differential evolution have demonstrated promising results on high-dimensional FS tasks.

PSO has become one of the most widely adopted evolutionary algorithms for FS due to its simple structure, efficient information-sharing mechanism, and fast convergence. Nevertheless, conventional PSO-based FS methods still encounter three structural limitations when dealing with sparse and high-dimensional search spaces. First, random initialization distributes particles without prior guidance throughout the binary search space, whereas effective feature subsets usually occupy only a small region with limited active dimensions. Consequently, a considerable number of fitness evaluations may be wasted on redundant and suboptimal solutions before the swarm identifies promising search regions. Second, conventional uniform feature perturbation strategies treat all dimensions equally and neglect the feature-level improvement information accumulated during previous evaluations, preventing the search from exploiting historically informative features. Third, existing PSO-based FS methods typically employ fixed objective weights to balance classification error and feature subset size during the entire optimization process. However, the desired trade-off between feature reduction and accuracy refinement varies across search stages: aggressive sparsity enhancement is preferred in early exploration, whereas accurate decision boundary refinement becomes more important after a compact subset has been obtained.

Recent evolutionary FS studies have attempted to improve search performance through enhanced population initialization, adaptive evolutionary operators, and refined fitness evaluation strategies [8–11]. Although these approaches have achieved promising results, they usually focus on isolated components of the optimization process and do not jointly address inefficient initialization, insufficient exploitation of accumulated feature knowledge, and inflexible objective weighting. Recently, multi-task PSO has emerged as an effective paradigm for high-dimensional FS with fixed-length binary representations. For instance, MEL [12] adopts cooperative subpopulations and a fixed-length representation to efficiently handle feature spaces containing tens of thousands of dimensions. Despite its effectiveness, existing multi-task PSO-based FS frameworks still leave several challenges unresolved: the initial population is generally generated without sufficient feature-level prior knowledge, useful feature-importance information discovered during evolution is not explicitly retained and exploited, and the balance between classification performance and feature sparsity remains unchanged across different search stages.

To address these limitations, this paper proposes Knowledge-Guided Dual-Phase Swarm Evolutionary Learning (KDS-EL), a closed-loop framework that cohesively integrates domain prior injection, cross-iteration knowledge accumulation, and dynamic objective scheduling into a cooperative multi-task PSO backbone.

The main contributions of this paper are summarized as follows:

1. *MI-Guided Initialization with Adaptive Re-initialization*: A hybrid seeding strategy that injects continuous feature-relevance priors into population construction to eliminate early-stage search waste, coupled with an adaptive re-initialization mechanism that restores swarm diversity upon evolutionary stagnation (Section 3.2).
2. *Knowledge-Guided Elite Refinement (KER)*: A targeted elite perturbation operator that dynamically harvests cross-iteration feature-importance feedback with exponential memory decay, directing elite solutions toward historically effective features within a strictly bounded evaluation overhead (Section 3.4).

*Corresponding author.
E-mail address: corresponding.author@example.com (Anonymous Author).
ORCID(s):

Table 1: Taxonomic and methodological comparison between KDS-EL and representative evolutionary feature selection algorithms.

		Knowledge Engineering			Search Control		
Algorithm	Paradigm	MI Prior Seeding	Online Memory	Memory Decay	Elite Refinement	Dynamic Weighting	Complexity
BPSO [7]	SO-SI	×	×	×	×	×	High
FWPSO [8]	SO-SI	○	×	×	×	×	Moderate
SaWDE [9]	SO-DE	×	×	×	×	○	High
W-MOPSO-FS [13]	MO-SI	×	×	×	×	×	High
NNS-DFPG-CSA [14]	SO-CSA	○	×	×	×	×	Moderate
MI-MPODE [15]	SO-DE	✓	×	×	×	×	Moderate
EMSWOA [16]	MS-WOA	×	×	×	○	×	High
MEL [12]	MT-SI	×	✓	×	×	×	High
KDS-EL (Proposed)	**MT-SI**	✓	✓	✓	✓	✓	**Low**

✓: Fully supported; ○: Partially supported; ×: Absent / not supported.
Paradigms: SO: Single-Objective; MO: Multi-Objective; MS: Multi-Swarm; MT: Multi-Task; SI: Swarm Intelligence; DE: Differential Evolution; CSA: Crow Search Algorithm; WOA: Whale Optimization Algorithm.

3. *Two-Phase Weight Scheduling (TPS)*: A stage-aware objective scheduling mechanism that temporally decouples early aggressive feature pruning from late-stage accuracy refinement, overcoming the rigidity of static weighting (Section 3.5).
4. *Comprehensive Empirical Validation*: Extensive evaluations across 20 high-dimensional benchmark datasets (spanning 500 to 54,675 features) against 13 state-of-the-art competitors, rigorously validated through non-parametric statistical testing, component contribution analysis, and parametric sensitivity investigations (Section 4).

Extensive experiments demonstrate that KDS-EL achieves the best classification performance, obtaining an average classification accuracy of 0.9574 with only 87.8 selected features. Compared with the strongest single-objective and multi-objective competitors, KDS-EL improves accuracy by 3.97 and 4.42 percentage points, respectively, while achieving a 3.4× speedup over standard MEL [12] (reducing average runtime from 175.0 s to 51.4 s). Furthermore, Wilcoxon signed-rank tests with Holm–Bonferroni correction confirm that the improvements are statistically significant over all 13 competitors.

The remainder of this paper is organized as follows. Section 2 formulates the feature selection optimization problem and reviews related literature and background concepts. Section 3 presents the proposed KDS-EL framework and formulates its three constituent mechanisms in detail. Section 4 presents comprehensive experimental results, statistical tests, component contribution analysis, and parameter sensitivity investigations. Finally, Section 5 concludes the paper and discusses prospective research directions.

2. Related Work and Background

This section provides a systematic survey of evolutionary feature selection paradigms, establishes the formal mathematical model of feature selection, formulates continuous and binary particle swarm optimization, reviews mutual information for feature relevance estimation, reviews the Multi-Task Evolutionary Learning (MEL) algorithm, and defines the unique taxonomic positioning of KDS-EL.

2.1. Problem formulation of feature selection

Feature selection (FS) is a combinatorial optimization problem that aims to identify an optimal, minimal subset of informative features from an original D-dimensional feature space, simultaneously maximizing classification performance and eliminating irrelevant or redundant dimensions [1, 2].

Let $\mathcal{D} = \{(\mathbf{f}_i, y_i)\}_{i=1}^{N_{\text{samp}}}$ denote a supervised dataset comprising N_{samp} training instances, where $\mathbf{f}_i = [f_{i,1}, f_{i,2}, \ldots, f_{i,D}] \in \mathbb{R}^D$ is the full feature vector and $y_i \in \{1, 2, \ldots, C\}$ is the corresponding discrete class label across C categories. A candidate feature subset is mathematically represented as a binary decision vector $\mathbf{x} = [x_1, x_2, \ldots, x_D] \in \{0, 1\}^D$, defined as:

$$x_j = \begin{cases} 1, & \text{if feature } f_j \text{ is selected,} \\ 0, & \text{if feature } f_j \text{ is discarded.} \end{cases} \tag{1}$$

The feature selection task inherently balances two dual criteria: minimizing classification error rate while minimizing the feature selection ratio (FSR). In evolutionary wrapper feature selection, this is mathematically formulated as minimizing a unified scalar objective function $f(\mathbf{x})$ [14, 16]:

$$\min_{\mathbf{x}\in\{0,1\}^D} f(\mathbf{x}) = w_a \cdot (1 - \text{Acc}(\mathbf{x})) + w_f \cdot \left(\frac{\|\mathbf{x}\|_0}{D}\right), \tag{2}$$

where $\text{Acc}(\mathbf{x})$ denotes the cross-validation classification accuracy evaluated on the masked feature subset $\{\mathbf{f}_i \odot \mathbf{x}\}_{i=1}^{N_{\text{samp}}}$ using an inductive learning algorithm, $\|\mathbf{x}\|_0 = \sum_{j=1}^{D} x_j$ denotes the L_0-pseudo-norm representing the number of active features, and $w_a \in [0, 1], w_f = 1 - w_a$ are weighting parameters balancing classification accuracy and subset compactness.

Because the cardinality of the discrete search space is $|\Omega| = 2^D$, exhaustive evaluation is computationally intractable for high-dimensional biomedical regimes ($D \gg 10^3$). Traditional feature selection techniques are categorized into three classical paradigms: filter, wrapper, and embedded methods [2]. *Filter methods* evaluate features based on intrinsic statistical properties, such as mutual information (mRMR [17]), correlation (FCBF [18]), or centrality measures [19]. *Embedded methods* perform feature selection during model training, such as Bolasso [20] and tree-based ensemble importance (FS-GBDT [21]). *Wrapper methods* evaluate candidate subsets using the direct predictive feedback of a classifier [4]. While wrappers capture intricate feature interactions and deliver superior accuracy, navigating the 2^D combinatorial space requires efficient stochastic global search heuristics, for which Evolutionary Computation (EC) paradigms—including Genetic Algorithms (GA) [22], Ant Colony Optimization (ACO) [23], Artificial Bee Colony (ABC) [24], and Swarm Intelligence (SI) [3, 25]—have emerged as dominant approaches.

To combat the curse of dimensionality in small-sample high-dimensional regimes, numerous specialized single-objective evolutionary algorithms have been developed: self-adaptive swarm mechanisms [26], progressively refined search trajectories [27], self-adaptive differential evolution with adaptive weighting (SaWDE) [9], multi-layer perceptron-guided PSO (IGPSO) [10], multi-criterion scoring PSO (PSO-CSM) [11], reinforcement learning-based swarm control [28], and divide-and-conquer subspace decomposition [29]. While these methods enhance local search operators, they universally treat feature perturbation as a memoryless stochastic trial, discarding the accumulated evolutionary feedback across generations.

2.2. Particle swarm optimization

PSO, originally introduced by Kennedy and Eberhart [5], is a population-based stochastic optimization paradigm inspired by the collective foraging behaviors of bird flocks. A swarm consists of N particles navigating a D-dimensional continuous search space $\mathbb{R}^D$. Each particle $i \in \{1, \ldots, N\}$ maintains a continuous position vector $\mathbf{p}_i^{(t)} = [p_{i,1}^{(t)}, \ldots, p_{i,D}^{(t)}]$ and a continuous velocity vector $\mathbf{v}_i^{(t)} = [v_{i,1}^{(t)}, \ldots, v_{i,D}^{(t)}]$ at iteration t. The historical trajectory is governed by its personal best position $\mathbf{p}_{\text{best},i}$ and the swarm's global best position $\mathbf{g}_{\text{best}}$.

The velocity and position updates are formulated as:

$$\mathbf{v}_i^{(t+1)} = \omega \mathbf{v}_i^{(t)} + c_1 r_1 (\mathbf{p}_{\text{best},i} - \mathbf{p}_i^{(t)}) + c_2 r_2 (\mathbf{g}_{\text{best}} - \mathbf{p}_i^{(t)}), \tag{3}$$

$$\mathbf{p}_i^{(t+1)} = \mathbf{p}_i^{(t)} + \mathbf{v}_i^{(t+1)}, \tag{4}$$

where ω denotes the inertia weight governing momentum retention [6], c_1 and c_2 are cognitive and social acceleration coefficients, and $r_1, r_2 \sim \mathcal{U}(0, 1)$ are independent uniform random stochastic factors.

To adapt PSO to combinatorial binary feature selection where candidate solutions represent binary indicator masks $\mathbf{x}_i \in \{0,1\}^D$, Kennedy and Eberhart [7] proposed the standard Binary PSO (BPSO). BPSO transforms continuous velocities into selection probabilities via a Sigmoid transfer function:

$$S(v_{i,j}^{(t+1)}) = \frac{1}{1+\exp(-v_{i,j}^{(t+1)})}, \quad x_{i,j}^{(t+1)} = \begin{cases} 1, & \text{if rand}() < S(v_{i,j}^{(t+1)}), \\ 0, & \text{otherwise.} \end{cases} \tag{5}$$

While effective for low-dimensional problems, standard BPSO suffers severe exploration-exploitation dilemmas in ultra-high-dimensional spaces: continuous velocity updates provide no domain-grounded guidance toward sparse, informative biomarker subspaces. In KDS-EL, particles maintain continuous normalized position coordinates $\mathbf{p}_i \in [0,1]^D$, mapped deterministically to binary feature masks via a strict thresholded Heaviside step function $x_{i,j} = \mathbb{I}(p_{i,j} \geq \theta)$ with $\theta = 0.60$, eliminating stochastic binarization variance.

2.3. Mutual information

Mutual information (MI), rooted in Shannon's information theory [30], quantifies the mutual dependence and reduction in uncertainty between two random variables. For a feature variable f_j and the discrete class label $Y \in \{1,\ldots,C\}$, the mutual information $I(f_j;Y)$ is defined as:

$$\begin{aligned} I(f_j;Y) &= H(Y) - H(Y \mid f_j) \\ &= \sum_{v \in \text{Val}(f_j)} \sum_{c=1}^{C} P(f_j = v, Y = c) \log_2 \frac{P(f_j = v, Y = c)}{P(f_j = v)P(Y = c)}, \end{aligned} \tag{6}$$

where $H(Y)$ is the marginal Shannon entropy of the class labels and $H(Y \mid f_j)$ is the conditional entropy given feature f_j. High values of $I(f_j;Y)$ indicate strong discriminative relevance to classification targets.

Conventional filter feature selection methods [14, 17] evaluate $I(f_j;Y)$ across all D dimensions and execute a hard truncation:

$$\mathcal{S}_{\text{Filter}} = \{\text{IDX}(1), \text{IDX}(2), \ldots, \text{IDX}(\lfloor k \cdot D \rfloor)\}, \tag{7}$$

where IDX denotes the feature indices sorted in descending order of $I(f_j;Y)$ and $k \in (0,1]$ represents the retention ratio, permanently discarding all features outside the top-k ranking. While computationally rapid ($O(D \log D)$), hard truncation fundamentally ignores multi-feature interactions: individual features with moderate MI scores that are collectively decisive when combined are permanently eliminated. Rather than using hard truncation, KDS-EL utilizes normalized MI scores as a soft prior to construct continuous initial population seeds (Section 3.2), injecting filter knowledge into the evolutionary wrapper without irreversibly shrinking the search space.

2.4. Multi-task evolutionary learning

Evolutionary Multi-Task Optimization (EMTO) aims to solve multiple distinct optimization tasks simultaneously by facilitating inter-task knowledge transfer through a unified chromosomal representation [31]. In feature selection, EMTO frameworks partition the global population into collaborative subpopulations to explore complementary search subspaces, including minion-assisted multi-tasking [32], evolutionary multi-task PSO for hyperspectral and gene expression data [33, 34], multi-task bee colony band selection [35], and large-scale multi-objective multi-tasking [36].

Most notably, Zhang et al. [12] introduced Multi-Task Evolutionary Learning (MEL), which establishes a dual-subpopulation PSO framework. In MEL, a population of N particles is partitioned into two collaborative subpopulations, $\mathcal{P}_1$ and $\mathcal{P}_2$ ($|\mathcal{P}_1| = |\mathcal{P}_2| = N/2$), operating on dual complementary search tasks:

Subpopulation 1 ($\mathcal{P}_1$): Continuous swarm exploration. Particles update continuous velocities and positions using PSO dynamics, augmented with an inter-task knowledge transfer channel directing particles toward the historical global best of Subpopulation 2 ($\mathbf{g}_{\text{best},2}$):

$$\mathbf{v}_{1,i}^{(t+1)} = \omega \mathbf{v}_{1,i}^{(t)} + c_1 r_1 (\mathbf{p}_{\text{best},1,i} - \mathbf{p}_{1,i}^{(t)}) + c_2 r_2 (\mathbf{g}_{\text{best},1} - \mathbf{p}_{1,i}^{(t)}) + c_3 r_3 (\mathbf{g}_{\text{best},2} - \mathbf{p}_{1,i}^{(t)}). \tag{8}$$

Subpopulation 2 ($\mathcal{P}_2$): Importance-guided discrete sampling. Rather than continuous mechanics, $\mathcal{P}_2$ samples binary feature masks directly based on a shared feature-importance matrix $W_{\text{MEL}} \in \mathbb{R}^D$.

In MEL, both subpopulations share a scalar fitness function combining classification error and feature subset ratio via static weights:

$$f_{\text{MEL}}(\mathbf{x}) = \alpha \cdot (1 - \text{Acc}(\mathbf{x})) + \beta \cdot \left(\frac{\|\mathbf{x}\|_0}{D} \right), \tag{9}$$

where $\text{Acc}(\mathbf{x})$ is the 5-fold cross-validation accuracy of a 3-NN classifier, and $\alpha = 0.90, \beta = 0.10$ are fixed constants [12].

While MEL demonstrates the utility of cross-task knowledge transfer, its reliance on uniform random initialization, unconstrained importance accumulation without memory discounting, and rigid static weighting creates substantial search inefficiencies in ultra-high-dimensional spaces. The proposed KDS-EL framework directly inherits this dual-subpopulation backbone while resolving these deficiencies through three complementary mechanisms detailed in Section 3.

2.5. Multi-objective feature selection

Feature selection naturally involves two competing criteria: minimizing classification error and minimizing the number of selected features. Multi-Objective Feature Selection (MOFS) algorithms model this trade-off by searching for a Pareto front of non-dominated solutions. Pioneered by Xue et al. [13], diverse multi-objective evolutionary algorithms have been adapted to feature selection, including NSGA-II-FS [37], MOEA/D-FS [38], MOPSO-FS [39], and hybrid filter-wrapper multi-objective architectures [40, 41].

Despite their theoretical appeal, MOFS methods encounter severe operational limitations when scaled to ultra-high-dimensional biomedical data ($D > 10^3$). First, maintaining an external non-dominated archive and executing Pareto ranking over dense high-dimensional vectors incurs substantial computational overhead, often exceeding 110–140 seconds in wall-clock time (Section 4.2.6). Second, under severe sample sparsity, non-dominated sorting tends to produce large, diffused Pareto fronts populated by non-sparse candidate subsets, often retaining 2,000–5,000 active features and failing to deliver the ultra-sparse compact biomarker signatures with fewer than 100 features required for clinical interpretability.

To circumvent Pareto archive bloat and achieve extreme subset sparsity within tight computational budgets, weighted single-objective scalarization remains the preferred formulation in high-dimensional biomarker discovery [12, 16]. However, conventional weighted frameworks enforce fixed static weights throughout search, such as $\mathbf{w} = [0.90, 0.10]$, which fails to adapt to the distinct evolutionary phases of feature selection. KDS-EL overcomes this limitation by introducing a dynamic Two-Phase Weight Scheduling strategy that aggressively eliminates noise features early and fine-tunes classification accuracy late.

2.6. Knowledge-guided search

Integrating domain-specific prior knowledge into meta-heuristics has attracted surging attention as a means to escape the blindness of black-box search [14, 15, 42]. In high-dimensional feature selection, mutual information serves as the primary information-theoretic metric to quantify non-linear feature relevance [30].

Existing knowledge-guided evolutionary algorithms typically adopt one-shot offline injection via either static filter pre-filtering or static velocity biasing. In static filter pre-filtering (e.g., MI-MPODE [15] and NNS-DFPG-CSA [14]), algorithms compute MI rankings offline and initialize the population or restrict the search space to the top-k ranked features. In static velocity biasing (e.g., FWPSO [8]), static feature-weight vectors are incorporated to bias particle velocity equations throughout evolution. While beneficial, one-shot knowledge injection lacks online adaptability: once injected, the static prior is never updated or refined by the empirical trial-and-error feedback generated during subsequent wrapper evaluations. Conversely, learning-based surrogates, such as neural networks [10] or Gaussian processes [43], introduce severe surrogate training overhead and risk model misspecification under sample-scarce conditions.

2.7. Taxonomic comparison

To overcome the collective limitations identified in the existing literature, KDS-EL establishes a closed-loop, knowledge-guided evolutionary learning framework that unifies offline domain priors, online memory harvesting, and stage-aware objective scheduling. Specifically, unlike conventional filter wrappers that permanently eliminate features via hard truncation, KDS-EL uses normalized MI scores as continuous Gaussian-perturbed seeds, preserving the full search space while ensuring an informative, sparse starting basin with $\|\mathbf{x}\|_0 \ll D/2$. In contrast to black-box swarm mutations, KDS-EL dynamically harvests evolutionary fitness feedback into a cumulative importance matrix W, regulated by an exponential forgetting factor $\gamma = 0.99$. The KER operator re-deploys this accumulated knowledge to execute targeted, directional elite flips within a strictly bounded computational budget ($k_e = 3 \ll N$). Furthermore, unlike rigid static weighting with $\mathbf{w} = [0.90, 0.10]$, KDS-EL's Two-Phase Weight Scheduling temporally decouples coarse feature pruning in Phase 1 ($w_f = 0.30$) from fine accuracy refinement in Phase 2 ($w_a = 0.95$), achieving simultaneous gains in accuracy, feature compactness, and runtime efficiency. Table 1 summarizes the structural comparison between KDS-EL and representative state-of-the-art evolutionary feature selection frameworks.

3. The Proposed KDS-EL Method

To overcome the dichotomy between static domain priors and memoryless stochastic search, this section presents the Knowledge-Guided Dual-Phase Swarm Evolutionary Learning (KDS-EL) framework. KDS-EL is designed as a closed-loop evolutionary learning paradigm specifically to overcome three fundamental structural deficiencies of conventional meta-heuristics in ultra-high-dimensional feature selection: (i) the cold-start inefficiency and noise entrapment caused by uniform random initialization in sparse feature spaces, (ii) the permanent loss of cross-iteration feature-importance signals due to memoryless random perturbation, and (iii) the inability of static objective weighting to adapt to the divergent temporal requirements of coarse feature reduction versus fine-grained accuracy refinement.

3.1. The framework of KDS-EL

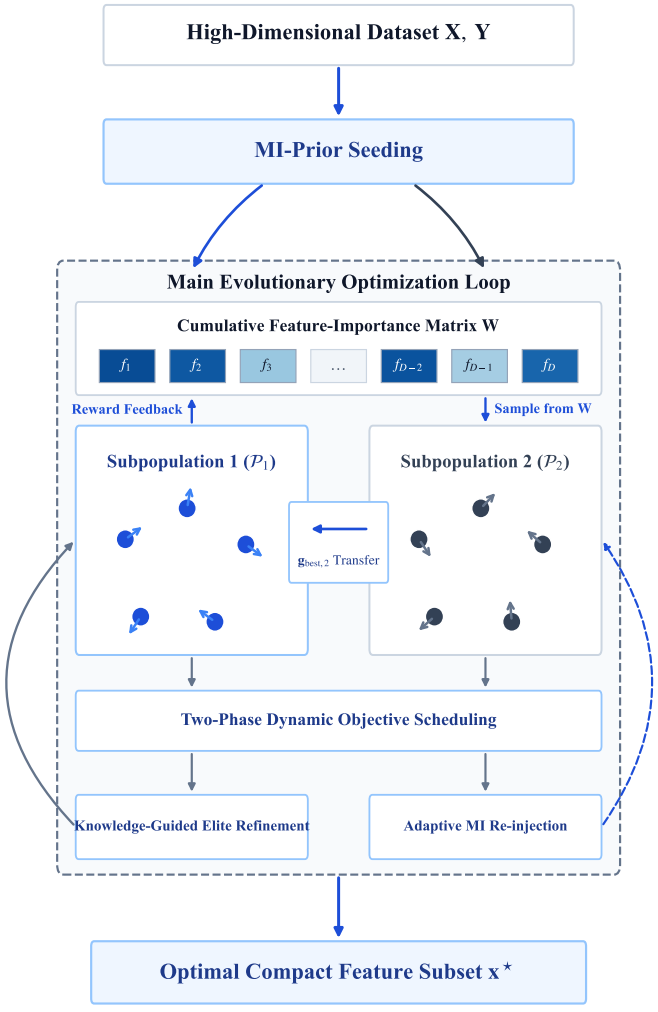


Figure 1: Overall framework of the proposed KDS-EL algorithm.

The macroscopic architecture of KDS-EL is illustrated in Fig. 1. Built upon a cooperative multi-task dual-subpopulation backbone [12], KDS-EL divides a population of $N = 20$ binary particles into two collaborative subpopulations of equal size ($N_1 = N_2 = 10$). Subpopulation 1 ($\mathcal{P}_1$) executes standard continuous-space PSO velocity-position updates with an inter-task knowledge transfer channel directing particles toward the historical global best of Subpopulation 2. Subpopulation 2 ($\mathcal{P}_2$) performs probability-driven feature sampling guided by the population's accumulated feature-importance landscape.

The optimization process operates in a closed loop through four tightly coupled functional stages: (i) *Filter-level quality bootstrapping (offline)*, where mutual information between each feature and class labels is computed once prior to evolution, establishing a normalized prior relevance vector $\mathbf{m} \in [0,1]^D$; (ii) *Knowledge-informed hybrid seeding*, where an initial population of N particles is generated by blending Gaussian-perturbed MI seeds with uniform random particles at a ratio of MI_init_ratio = 0.50, followed by Heaviside step binarization with threshold $\theta = 0.60$; (iii) *Dynamic two-phase evolutionary loop*, where across $T = 100$ iterations, the objective weights $\mathbf{w}(t) = [w_a(t), w_f(t)]$ dynamically transition from aggressive feature pruning ($\mathbf{w}_1 = [0.70, 0.30]$ for $t \le 0.7T$) to accuracy-oriented refinement ($\mathbf{w}_2 = [0.95, 0.05]$ for $t > 0.7T$); and (iv) *Accumulative knowledge harvesting and elite refinement*, where fitness improvements during evaluation are harvested into a continuous feature-importance matrix $W \in \mathbb{R}^D$ discounted by an exponential forgetting factor $\gamma = 0.99$, elite solutions are refined via the KER operator under 1-to-1 in-place replacement, and an adaptive re-initialization mechanism recovers the swarm if search stalls for stall_patience = 20 generations.

3.2. Population initialization strategy

To overcome the cold-start inefficiency and noise entrapment of uniform random initialization in sparse high-dimensional search spaces ($D \gg 10^3$), where standard initialization produces candidate solutions with approximately 50% active features ($\|\mathbf{x}\|_0 \approx D/2$) that waste evaluation budgets in noisy search basins, KDS-EL introduces a filter-wrapper hybrid seeding strategy coupled with adaptive stagnation recovery.

Given the training dataset $\mathcal{D} = \{(\mathbf{f}_i, y_i)\}_{i=1}^{N_{\text{samp}}}$, the mutual information scores $\{I(f_j; Y)\}_{j=1}^{D}$ computed via (6) are normalized into a prior relevance vector $\mathbf{m} \in [0,1]^D$:

$$m_j = \frac{I(f_j;Y) - \min_{1\le k\le D} I(f_k;Y)}{\max_{1\le k\le D} I(f_k;Y) - \min_{1\le k\le D} I(f_k;Y) + \epsilon}, \tag{10}$$

where $\epsilon = 10^{-12}$ prevents division by zero.

To prevent population homogenization while guiding search toward informative dimensions, the initial population is split into two halves based on $\rho = \text{MI_init_ratio} = 0.50$:

$$\mathbf{p}_i = \begin{cases} \text{clip}(\mathbf{m} + \boldsymbol{\varepsilon}_i,\ 0,\ 1), & \boldsymbol{\varepsilon}_i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_D), \quad \text{for } 1 \le i \le \rho N, \\ \mathbf{u}_i, & \mathbf{u}_i \sim \mathcal{U}(0,1)^D, \quad \text{for } \rho N < i \le N, \end{cases} \tag{11}$$

where $\sigma = 0.10$ represents the Gaussian noise perturbation scale and $\text{clip}(z, 0, 1) = \max(0, \min(1, z))$.

Continuous positions $\mathbf{p}_i = [p_{i,1}, \ldots, p_{i,D}]$ are mapped into binary feature selection vectors $\mathbf{x}_i \in \{0,1\}^D$ using a strict Heaviside step binarization with threshold $\theta = 0.60$:

$$x_{i,j} = \begin{cases} 1, & \text{if } p_{i,j} \ge \theta, \\ 0, & \text{if } p_{i,j} < \theta. \end{cases} \tag{12}$$

Because $\theta = 0.60 > 0.50$, only features possessing above-average mutual information are initialized as active, naturally yielding a sparse initial configuration ($\|\mathbf{x}_i\|_0 \approx 0.15D \sim 0.30D$) that accelerates early convergence, as formulated in Stage 2 of Section 3.6.

3.2.1. Adaptive re-initialization

High-dimensional multimodal search landscapes frequently cause swarm particles to become trapped in suboptimal local attractors. When the global-best fitness f_{gbest} fails to improve for $\tau = \text{stall_patience} = 20$ consecutive iterations, an adaptive re-initialization routine is executed. First, the population is sorted in ascending order of personal-best fitness f_i^{pbest}. The superior half ($\{i \mid \text{rank}(i) \le N/2\}$) is preserved to safeguard historical elite trajectories, whereas the inferior half ($\{i \mid \text{rank}(i) > N/2\}$) is fully refreshed using the Gaussian-perturbed MI distribution:

$$\mathbf{p}_i \leftarrow \text{clip}(\mathbf{m} + \boldsymbol{\varepsilon}_i,\ 0,\ 1), \quad x_{i,j} \leftarrow \mathbb{I}(p_{i,j} \ge \theta). \tag{13}$$

The refreshed particles are then re-evaluated under the current objective weights, and the stagnation counter is reset to zero. This restores diversity within a high-potential subspace without completely discarding established elite knowledge.

3.3. Feature importance learning

To exploit the historical trial-and-error feedback generated across generations and overcome the memoryless amnesia of standard swarm mutations, KDS-EL maintains a continuous cumulative feature-importance matrix $W \in \mathbb{R}^D$, initialized to $\mathbf{0}$.

At generation t, whenever a candidate particle $\mathbf{x}_i^{(t)}$ achieves a strict fitness improvement over its personal best ($f_i^{(t)} < f_i^{\text{pbest}}$), the differential feature subset $C_i^{(t)}$ containing features whose selection status changed is extracted:

$$C_i^{(t)} = \{j \in \{1, \ldots, D\} \mid x_{i,j}^{(t)} \ne p_{\text{best},i,j}\}. \tag{14}$$

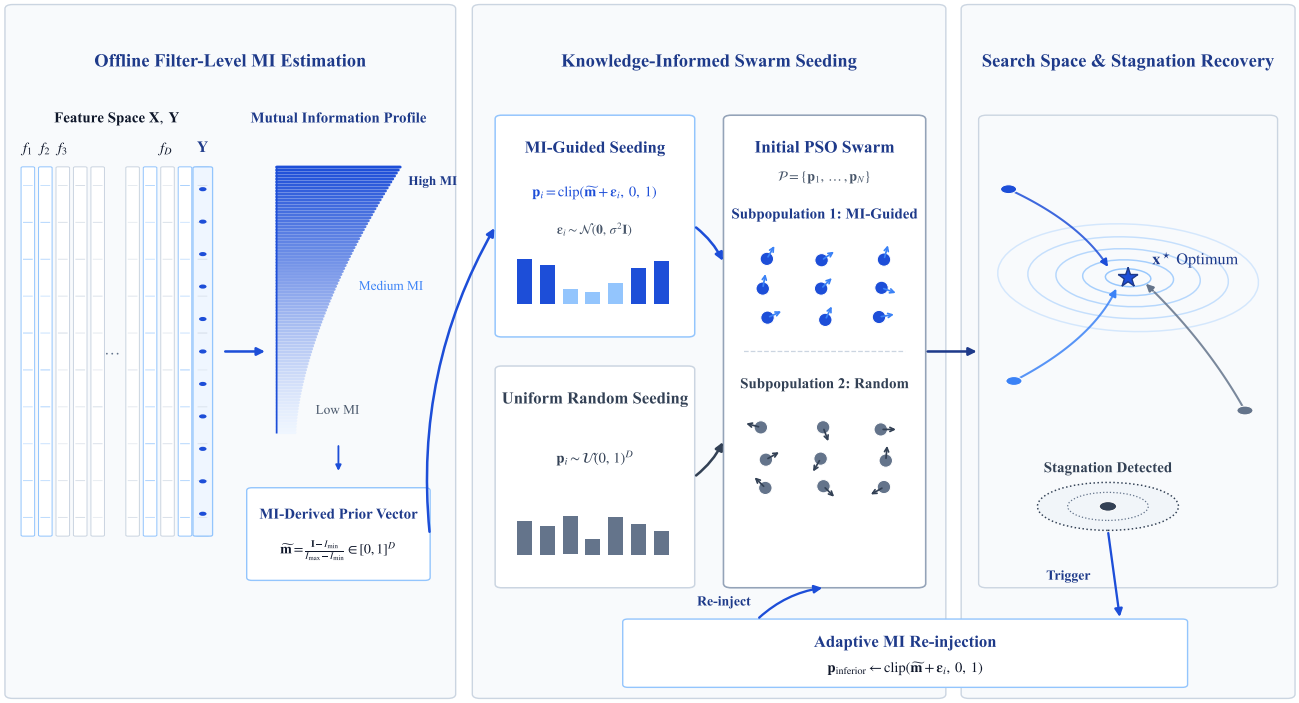


Figure 2: Schematic diagram of the MI-prior seeding and adaptive re-injection mechanism.

Let $\delta_i^{(t)} = f_i^{\text{pbest}} - f_i^{(t)} > 0$ denote the fitness improvement. Each changed feature $j \in C_i^{(t)}$ receives a directional reward proportional to $\delta_i^{(t)}$:

$$W_j \leftarrow W_j + \delta_i^{(t)} \cdot (2x_{i,j}^{(t)} - 1), \quad \forall j \in C_i^{(t)}, \tag{15}$$

where $(2x_{i,j}^{(t)} - 1) = +1$ if feature j was newly activated and -1 if deactivated. Non-improving particles ($\delta_i^{(t)} \le 0$) exert zero impact on W.

To prevent early-discovered features from dominating the matrix indefinitely and allow the algorithm to adapt as the search shifts toward fine-grained subspaces, an exponential memory decay is applied at the conclusion of each generation:

$$W \leftarrow \gamma \cdot W. \tag{16}$$

The decay factor $\gamma = 0.99$ ensures that unless a feature repeatedly contributes to fitness improvements, its score depreciates exponentially, preventing score saturation while dynamically maintaining ranking sensitivity.

3.4. Elite refinement strategy

To exploit the evolutionary knowledge distilled in W and direct local search toward high-potential biomarker dimensions, we propose the Knowledge-Guided Elite Refinement (KER) strategy. Knowledge acquisition in KER requires no additional classifier evaluations, as the importance matrix is updated directly from the fitness comparisons already performed for the population. KER itself spends only a strictly bounded refinement budget ($k_e = 3 \ll N$), preserving the dominant per-iteration computational complexity.

At each generation, the top $k_e = 3$ individuals ranked by fitness are selected. For each elite individual $\mathbf{x}_e \in \{0,1\}^D$, KER executes a probabilistic dual-branch operation (Fig. 3):

Branch 1: Ergodic random exploration ($p_{\text{rand}} = 0.10$): To guarantee search ergodicity, $\lfloor \phi D \rfloor$ features are selected uniformly at random and inverted:

$$x'_{e,j} \leftarrow 1 - x_{e,j}, \quad \forall j \in \mathcal{S}_{\text{rand}}, \quad |\mathcal{S}_{\text{rand}}| = \lfloor \phi D \rfloor, \tag{17}$$

where $\phi = 0.01$ defines the refinement perturbation budget.

Branch 2: Knowledge-guided directional refinement ($1 - p_{\text{rand}} = 0.90$): Each feature $j \in \{1, \ldots, D\}$ is assigned a state-dependent directional flip score σ_j:

$$\sigma_j = (1 - 2x_{e,j}) \cdot W_j = \begin{cases} +W_j, & \text{if } x_{e,j} = 0, \\ -W_j, & \text{if } x_{e,j} = 1. \end{cases} \tag{18}$$

A positive score $+W_j$ favors activating an absent feature with proven positive historical contribution, whereas a negative score $-W_j$ favors deactivating a present feature that previously degraded fitness. We identify the index set $\mathcal{S}_{\text{know}}$ containing the $\lfloor \phi D \rfloor$ features possessing the highest flip scores σ_j:

$$\mathcal{S}_{\text{know}} = \arg \max_{|\mathcal{S}| = \lfloor \phi D \rfloor} \sum_{j \in \mathcal{S}} \sigma_j, \quad \mathcal{S} \subseteq \{1, \ldots, D\}, \tag{19}$$

and invert the selected features: $x'_{e,j} \leftarrow 1 - x_{e,j}, \forall j \in \mathcal{S}_{\text{know}}$.

Greedy 1-to-1 in-place replacement: The refined mutant $\mathbf{x}'_e$ is evaluated in a closed loop, and admitted via greedy 1-to-1 in-place replacement:

$$\mathbf{x}_e^{\star} = \begin{cases} \mathbf{x}'_e, & \text{if } f(\mathbf{x}'_e, \mathbf{w}(t)) < f(\mathbf{x}_e, \mathbf{w}(t)), \\ \mathbf{x}_e, & \text{otherwise.} \end{cases} \tag{20}$$

If improved, $\mathbf{x}'_e$ immediately overwrites $\mathbf{x}_e$ in the population and updates the global best if $f(\mathbf{x}'_e) < f_{\text{gbest}}$. Because elite_k = 3 is fixed, KER executes within an exact bounded overhead of $k_e = 3$ evaluations per generation. The detailed execution of KER is formalized in Algorithm 1.

3.5. Two-phase weight scheduling

To align objective weights with the temporal structure of feature selection and resolve the conflict between early dimensionality reduction and late accuracy refinement, KDS-EL introduces the Two-Phase Weight Scheduling (TPS) strategy. The total generational budget $T = 100$ is partitioned at split point $\kappa = 0.70$. The dynamic dual-objective weight vector $\mathbf{w}(t) = [w_a(t), w_f(t)]$ is governed by:

$$\mathbf{w}(t) = \begin{cases} [0.70,\ 0.30], & \text{for } 1 \le t \le \lfloor \kappa T \rfloor, \\ [0.95,\ 0.05], & \text{for } \lfloor \kappa T \rfloor < t \le T. \end{cases} \tag{21}$$

The unified scalar fitness $f(\mathbf{x}, \mathbf{w}(t))$ of a binary feature subset $\mathbf{x} \in \{0,1\}^D$ at generation t is formulated as:

$$f(\mathbf{x}, \mathbf{w}(t)) = w_a(t) \cdot (1 - \text{Acc}(\mathbf{x})) + w_f(t) \cdot \left(\frac{\|\mathbf{x}\|_0}{D} \right), \tag{22}$$

where $\text{Acc}(\mathbf{x})$ denotes 5-fold cross-validation classification accuracy evaluated by a 3-NN classifier, and $\|\mathbf{x}\|_0/D$ represents the normalized feature selection ratio.

In Phase 1 ($t \le 0.7T$), the heavy feature-ratio penalty ($w_f = 0.30$) drives rapid elimination of non-informative dimensions. In Phase 2 ($t > 0.7T$), the weight switches to prioritize classification accuracy ($w_a = 0.95$), enabling fine-tuning of delicate biomarker interactions over the surviving compact subspace.

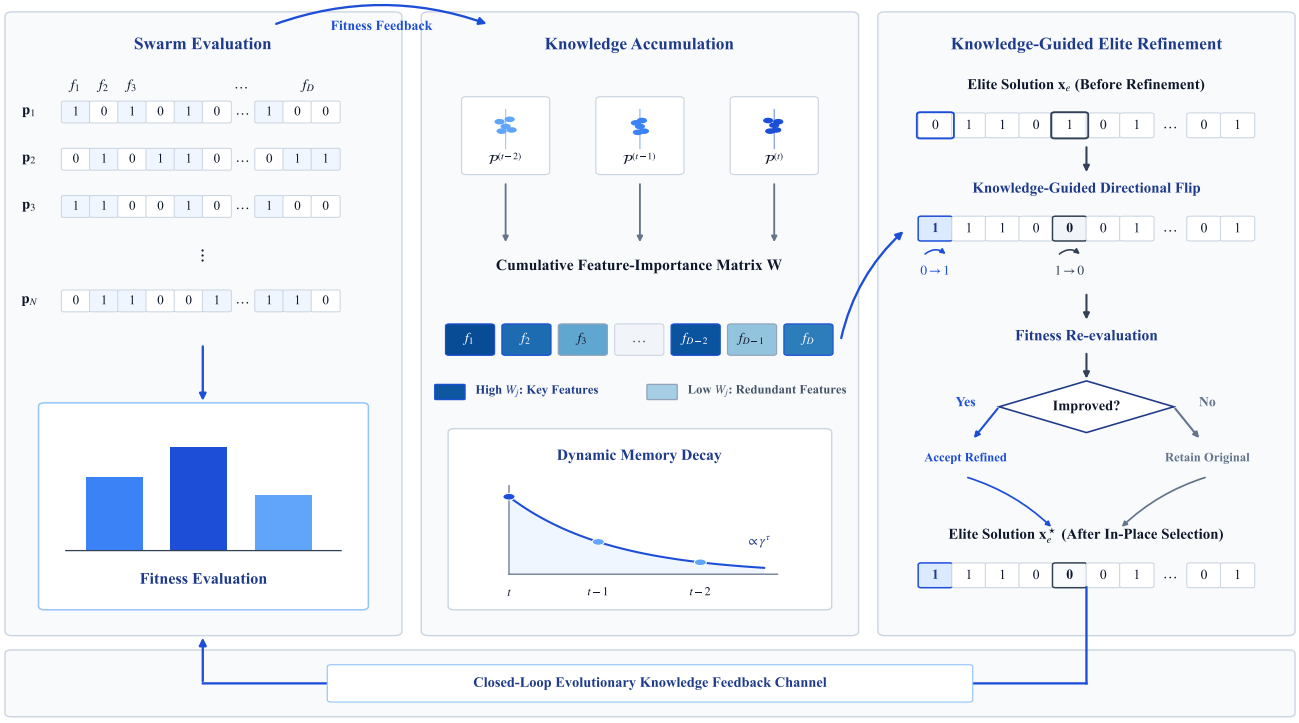


Figure 3: Flowchart of the KER strategy.

Algorithm 1 Knowledge-guided elite refinement (KER)

Require: Elite candidate solution $\mathbf{x}_e \in \{0,1\}^D$, cumulative importance matrix $W \in \mathbb{R}^D$, perturbation budget $\phi \in (0,1)$, exploration probability $p_{\text{rand}} \in [0,1]$, dynamic objective weights $\mathbf{w}(t)$
Ensure: Refined elite candidate $\mathbf{x}_e^\star \in \{0,1\}^D$

$k_{\text{flip}} \leftarrow \max(1, \lfloor \phi \cdot D \rfloor)$ ▷ Determine perturbation budget
Sample random variable $r \sim \mathcal{U}(0,1)$
if $r < p_{\text{rand}}$ **then** ▷ Branch 1: Ergodic random exploration
  Sample random feature subset $\mathcal{S}_{\text{rand}} \subset \{1,\dots,D\}$ with $|\mathcal{S}_{\text{rand}}| = k_{\text{flip}}$
  $\mathbf{x}'_e \leftarrow \mathbf{x}_e$; $x'_{e,j} \leftarrow 1 - x_{e,j}$, $\forall j \in \mathcal{S}_{\text{rand}}$
else ▷ Branch 2: Knowledge-guided directional refinement
  for $j = 1$ **to** D **do**
    $\sigma_j \leftarrow (1 - 2x_{e,j}) \cdot W_j$ ▷ Compute directional flip score
  end for
  Identify subset $\mathcal{S}_{\text{know}}$ of size k_{flip} with top σ_j scores via (19)
  $\mathbf{x}'_e \leftarrow \mathbf{x}_e$; $x'_{e,j} \leftarrow 1 - x_{e,j}$, $\forall j \in \mathcal{S}_{\text{know}}$
end if
Evaluate offspring fitness $f' \leftarrow f(\mathbf{x}'_e, \mathbf{w}(t))$ using (22)
if $f' < f(\mathbf{x}_e, \mathbf{w}(t))$ **then**
  $\mathbf{x}_e^\star \leftarrow \mathbf{x}'_e$ ▷ Greedy 1-to-1 in-place replacement
else
  $\mathbf{x}_e^\star \leftarrow \mathbf{x}_e$ ▷ Retain original elite
end if
return $\mathbf{x}_e^\star$

Algorithm 2 The framework of the proposed KDS-EL algorithm

Require: Training dataset $\mathcal{D} = (X, Y)$, population size N, max iterations T, binarization threshold θ, MI seeding ratio ρ, decay factor γ, elite pool size k_e, perturbation fraction ϕ, random branch probability p_{rand}, phase split ratio κ, stagnation threshold τ
Ensure: Optimal compact feature subset $\mathbf{x}^\star \in \{0,1\}^D$ and optimal fitness f_{gbest}

// Stage 1: Filter-level relevance extraction
Compute and normalize mutual information vector $\mathbf{m} \in [0,1]^D$ via (6) and (10)
// Stage 2: Knowledge-informed hybrid population seeding
Initialize population $\mathcal{P} = \{\mathbf{x}_i\}_{i=1}^N$ with ρN Gaussian MI seeds and $(1-\rho)N$ uniform particles via (11)–(12)
Initialize cumulative feature-importance matrix $W \leftarrow \mathbf{0} \in \mathbb{R}^D$, stall_count $\leftarrow 0$, and initial weights $\mathbf{w}(1) = [0.70, 0.30]$
Evaluate initial population fitness $\{f_i\}_{i=1}^N$ via (22) and initialize personal/global bests $\mathbf{p}_{\text{best},i}$, $\mathbf{g}_{\text{best}}$
// Stage 3: Dynamic two-phase evolutionary optimization loop
for $t = 1$ **to** T **do**
  Update dynamic objective weight vector $\mathbf{w}(t)$ according to (21)
  Update Subpopulation 1 ($\mathcal{P}_1$) via continuous PSO with inter-task knowledge transfer channel [12]
  Update Subpopulation 2 ($\mathcal{P}_2$) via probability-driven feature sampling guided by activation vector $\mathbf{P} = \sigma(W)$
  Evaluate candidate solution fitness $\{f(\mathbf{x}_i, \mathbf{w}(t))\}_{i=1}^N$ using 5-fold cross-validation k-NN classifier via (22)
  Update cumulative matrix W with differential fitness improvements and apply exponential memory decay γ via (15)–(16)
  Update personal bests $\mathbf{p}_{\text{best},i}$ and identify top k_e elite candidate indices $\mathcal{E}$
  // Knowledge-guided elite refinement (KER)
  for each $e \in \mathcal{E}$ **do**
    $\mathbf{p}_{\text{best},e} \leftarrow \text{KER}(\mathbf{p}_{\text{best},e}, W, \phi, p_{\text{rand}}, \mathbf{w}(t))$ ▷ Call Algorithm 1
  end for
  // Global best tracking and adaptive stagnation recovery
  Update global best $\mathbf{g}_{\text{best}} \leftarrow \arg\min_i f_i^{\text{pbest}}$ and update stall_count
  if stall_count $\geq \tau$ **then**
    Refresh inferior half of swarm using Gaussian-perturbed MI distribution via (13)
    stall_count $\leftarrow 0$
  end if
end for
return Optimal feature subset $\mathbf{x}^\star \leftarrow \mathbf{g}_{\text{best}}$ and optimal fitness f_{gbest}

3.6. Main algorithm

Algorithm 2 consolidates the holistic execution workflow of KDS-EL into a modular multi-stage optimization pipeline. The framework coordinates offline filter relevance extraction, hybrid stochastic seeding, two-phase dynamic objective scheduling, collaborative dual-subpopulation search, differential knowledge matrix harvesting, and targeted KER elite refinement.

3.7. Complexity analysis

To establish the computational efficiency of KDS-EL, we conduct a rigorous asymptotic complexity analysis covering all algorithmic components.

3.7.1. Time complexity decomposition

The execution time of KDS-EL is decomposed into two distinct phases:

Offline preprocessing phase (mutual information calculation): Computing the mutual information $I(f_j; Y)$ between each feature f_j and class label Y requires $O(N_{\text{samp}})$ operations per feature, where N_{samp} is the number of training samples. For all D features, computing and sorting the prior MI vector incurs:

$$T_{\text{pre}} = O(D \cdot N_{\text{samp}} + D \log D). \tag{23}$$

Because this step is executed only once prior to the evolutionary loop, its overhead is entirely negligible relative to iterative wrapper evaluations.

Iterative evolutionary main loop (T iterations): In each of the T generations, computational costs arise from: (i) updating particle continuous positions ($O(N \cdot D)$), (ii) differential matrix updating and decay ($O(N \cdot D + D)$), (iii) refining k_e elites via KER ($O(k_e \cdot (D + T_{\text{eval_single}}))$), (iv) occasional adaptive re-initialization ($O(\frac{N}{2} \cdot (D + T_{\text{eval_single}}))$), and (v) evaluating 5-fold cross-validation k-NN classification accuracy for N candidate particles:

$$T_{\text{eval}} = O(N \cdot N_{\text{samp}}^2 \cdot d_{\text{selected}}), \tag{24}$$

where $d_{\text{selected}} = \|\mathbf{x}\|_0$ is the number of active features in candidate $\mathbf{x}$.

Summing these terms across T generations yields the total asymptotic time complexity:

$$\mathcal{T}_{\text{total}} = O\Big(D \log D + T \cdot [N \cdot D + (N + k_e) \cdot N_{\text{samp}}^2 \cdot d_{\text{selected}}]\Big). \tag{25}$$

3.7.2. Theoretical complexity equivalence and practical acceleration

Because $k_e = 3 \ll N = 20$ and $N_{\text{samp}}^2 \cdot d_{\text{selected}} \gg D$, the dominant asymptotic per-generation complexity of KDS-EL is strictly governed by wrapper evaluations:

$$\mathcal{T}_{\text{KDS-EL}} = O(T \cdot N \cdot N_{\text{samp}}^2 \cdot d_{\text{selected}}), \tag{26}$$

which is identically of the same asymptotic order as standard MEL [12] ($O(T \cdot N \cdot N_{\text{samp}}^2 \cdot d_{\text{MEL}})$).

Crucially, in standard MEL, random initialization and sluggish feature reduction result in dense candidate subsets ($d_{\text{MEL}} \approx 2{,}722$ active features on average). In contrast, KDS-EL's MI seeding and Phase 1 aggressive feature pruning rapidly compress the active feature count to $d_{\text{KDS-EL}} \approx 88$ features. Because k-NN distance calculation scales linearly with d_{selected}, evaluating a KDS-EL particle is over an order of magnitude faster than evaluating an unpruned MEL particle. Consequently, KDS-EL delivers a 3.4× empirical wall-clock speedup, reducing average runtime from 175.0 s for MEL to 51.4 s while incorporating richer knowledge operators.

3.7.3. Space complexity

The memory footprint of KDS-EL comprises: (i) the population continuous and binary position matrices ($O(N \cdot D)$), (ii) personal-best vectors ($O(N \cdot D)$), (iii) the cumulative feature-importance matrix W ($O(D)$), and (iv) the precomputed normalized MI vector $\mathbf{m}$ ($O(D)$). The overall space complexity is:

$$\mathcal{S}_{\text{total}} = O(N \cdot D + N_{\text{samp}} \cdot D), \tag{27}$$

which scales strictly linearly with dimension D, ensuring minimal memory consumption.

4. Experimental results and discussion

This section outlines the experimental framework and presents a comprehensive empirical evaluation of KDS-EL across eight analytical dimensions: classification accuracy, feature subset compactness, fitness convergence dynamics, statistical significance, performance progression before and after algorithmic improvement, computational efficiency, hyperparameter sensitivity, and in-depth discussion on biomedical ceiling effects. All aggregate numbers are computed over the 20 benchmark datasets described in Section 4.1.

4.1. Experimental setup

This subsection outlines the experimental framework designed to evaluate the classification performance, dimensionality reduction capability, computational efficiency, and hyperparameter robustness of KDS-EL.

4.1.1. High-dimensional benchmark datasets

To validate the performance and generalization ability of KDS-EL across diverse problem domains, a total of 20 benchmark datasets widely adopted in the evolutionary feature selection literature [8, 11, 12] are selected. These datasets encompass diverse categories and sample sizes across biomedical gene-expression microarrays, text categorization corpora, image recognition, speech processing, and synthetic benchmarks. Such diverse data distributions and dimensional complexities effectively evaluate the generalization capability and feature reduction efficacy of feature selection algorithms under challenging conditions. The detailed characteristics of the 20 benchmark datasets are summarized in Table 2.

4.1.2. Comparison algorithms and parameter settings

To comprehensively benchmark KDS-EL against modern evolutionary feature selection paradigms, 14 competitive methods are evaluated under identical computational budgets, as summarized in Table 1. The competitor suite spans three major algorithmic paradigms: five single-objective evolutionary algorithms, namely FWPSO [8], PSO-CSM [11], SaWDE [9], IGPSO [10], and standard BPSO [7]; five multi-objective methods, namely W-MOPSO-FS, fMOPSO-FS [44], NSGA-II-FS [37], MOEA-D-FS [38], and MOPSO-FS [39]; as well as the original multi-task algorithm MEL [12] and two constituent variants, KDS-EL w/o TPS and KDS-EL w/o TPS&KER. All comparison algorithms adopt the parameter settings recommended in their original publications.

To ensure strict fairness, a uniform experimental protocol is maintained across all algorithms: population size $N = 20$, maximum generational budget $T_{\max} = 100$, binarization threshold $\theta = 0.60$, and a 3-NN classifier evaluated via stratified 5-fold cross-validation over 10 independent runs with distinct random seeds. The complete hyperparameter configurations of all algorithms are summarized in Table 3.

4.1.3. Evaluation metrics and statistical test protocols

Algorithm performance is evaluated across three primary criteria: classification accuracy, the number of selected features alongside the feature selection ratio, and computational runtime.

To determine whether the empirical improvements of KDS-EL are statistically significant, we follow the multi-problem statistical testing framework recommended by Demšar [45]. Specifically, the non-parametric Friedman test [46] is utilized to establish global algorithm rankings across all 20 benchmark datasets. Pairwise Wilcoxon signed-rank tests at the $\alpha = 0.05$ significance level assess directional superiority between KDS-EL and each competitor, with the Holm–Bonferroni step-down procedure [47] applied to rigorously control the family-wise error rate across multiple comparisons.

Table 2: Characteristics of the 20 benchmark datasets.

No.	Dataset	Samples (n)	Features (D)	Classes (C)	Domain
1	Adenoma	36	7,457	2	Gene expression
2	ALL_AML	72	7,129	2	Gene expression
3	ALL1	128	12,625	2	Gene expression
4	ALL4	93	12,625	2	Gene expression
5	Colon	62	2,000	2	Gene expression
6	DLBCL	77	7,129	2	Gene expression
7	Gastric	65	22,645	2	Gene expression
8	Gastric1	144	22,283	2	Gene expression
9	Gastric2	124	22,283	2	Gene expression
10	Leukaemia	72	7,129	2	Gene expression
11	Lymphoma	45	4,026	2	Gene expression
12	Myeloma	173	12,625	2	Gene expression
13	Prostate	102	12,625	2	Gene expression
14	SRBCT	83	2,308	4	Gene expression
15	T1D	101	54,675	2	Gene expression
16	BASEHOCK	1,993	4,862	2	Text
17	PCMAC	1,943	3,289	2	Text
18	COIL20	1,440	1,024	20	Image
19	ISOLET	1,560	617	26	Speech
20	madelon	2,600	500	2	Synthetic

4.2. Experimental analysis

This subsection presents a comprehensive empirical evaluation of KDS-EL across eight complementary analytical dimensions: classification accuracy, feature subset compactness, fitness convergence dynamics, statistical significance, performance progression before and after algorithmic improvement, computational efficiency, hyperparameter sensitivity, and in-depth discussion on biomedical ceiling effects.

4.2.1. Comparison of classification accuracy

Table 4 reports the mean classification accuracy and standard deviations over 10 independent runs across the 20 benchmark datasets for all 14 evaluated methods. The experimental results demonstrate that KDS-EL achieves superior classification performance, securing the highest overall mean accuracy of 0.9574 across the benchmark suite.

Table 3: Hyper-parameter configurations and experimental settings.

Parameter / Setting	Value / Configuration
Common experimental setup:	
Population size (N)	20
Maximum iterations (T)	100
k-NN classifier neighborhood (k)	3
Real-to-binary threshold (θ)	0.60
Cross-validation scheme	5-fold stratified, shuffled
Number of independent runs	10
KDS-EL (full framework):	
Inertia weight (ω)	0.90
Acceleration coefficients (c_1, c_2, c_3)	2.0, 2.0, 2.0
Knowledge weight decay factor (γ)	0.99
Elite perturbation fraction (ϕ)	0.01
Elite pool size (k_e)	3
Random exploration probability (p_{rand})	0.10
MI-guided initialization ratio (ρ)	0.50
Stagnation patience threshold (τ)	20 generations
Phase split ratio (α)	0.70
Phase-1 objective weights ($\mathbf{w}_1$)	[0.70, 0.30]
Phase-2 objective weights ($\mathbf{w}_2$)	[0.95, 0.05]
Ablative configurations:	
Fixed weights (KDS-EL w/o TPS)	$\mathbf{w}_s = [0.90, 0.10]$
Parameters (KDS-EL w/o TPS&KER)	$\omega = 0.90,\ c_{1,2,3} = 2.0,\ \rho = 0.50,\ \tau = 20$
Parameters (Standard MEL [12])	$\omega = 0.90,\ c_{1,2,3} = 2.0,\ \mathbf{w}_s = [0.90, 0.10]$
Benchmark competitor algorithms:	
Competitor algorithms [7–11, 37–39, 44]	Per original literature ($N = 20, T = 100$)

Across the benchmark suite, KDS-EL achieves an overall mean accuracy of 0.9574, outperforming all 13 competitors. Specifically, compared with its constituent variants and the original MEL algorithm, KDS-EL exceeds KDS-EL w/o TPS by 0.86%, KDS-EL w/o TPS&KER by 2.49%, and original MEL by 7.06%. Against established single-objective competitors such as FWPSO, PSO-CSM, SaWDE, IGPSO, and BPSO, KDS-EL achieves substantial accuracy gains ranging from 3.97% to 8.76%. Furthermore, KDS-EL consistently outperforms multi-objective evolutionary feature selection methods, including W-MOPSO-FS, fMOPSO-FS, NSGA-II-FS, MOEA-D-FS, and MOPSO-FS, by 4.42% to 5.65%.

On individual datasets, KDS-EL secures the highest classification accuracy on 15 of the 20 benchmarks, achieving strictly superior performance on 13 datasets and tying for the highest accuracy with KDS-EL w/o TPS on *Adenoma* and *Lymphoma*. The advantage of KDS-EL is especially pronounced on ultra-high-dimensional microarrays with severe sample sparsity. For example, on the 12,625-dimensional *ALL4* dataset, KDS-EL achieves an accuracy of 0.9522, outperforming FWPSO and MEL by 2.21% and 11.67%, respectively. On *Colon*, KDS-EL reaches 0.9500 accuracy, leading FWPSO by 2.79% and MEL by 14.52%. On the largest benchmark, *T1D*, spanning 54,675 dimensions, KDS-EL records 0.8598 accuracy, exceeding FWPSO by 3.41% and MEL by 9.74%.

A closer examination of the five datasets where KDS-EL does not record the nominal highest score provides valuable insights into its search behavior. On *ALL1* and *SRBCT*, although multi-objective methods achieve nominal scores of 1.0000 compared to 0.9998 and 0.9986 for KDS-EL, this minor gap reflects ceiling saturation where at most one sample differs under cross-validation (Section 4.2.8). On moderate-dimensional datasets such as *BASEHOCK*, *COIL20*, and *ISOLET*, NSGA-II-FS yields slightly higher accuracy by 0.32%–1.96% only by retaining massive subsets of 2,073, 283, and 185 features. In sharp contrast, KDS-EL attains competitive accuracy with only 134, 84, and 60 features, delivering an order-of-magnitude reduction in feature redundancy.

In addition to superior central tendency, KDS-EL demonstrates remarkably low run-to-run dispersion, as reflected by the standard deviations in Table 4. On high-dimensional datasets such as *Gastric1* with 22,283 features, KDS-EL exhibits an accuracy standard deviation of only 0.0089, compared with 0.0215 for FWPSO and 0.0381 for MEL. This stability confirms that hybrid MI initialization and knowledge-guided KER perturbation effectively steer particles toward robust global basins rather than relying on fortuitous random initialization.

4.2.2. Comparison of the number of selected features

In high-dimensional feature selection, classification accuracy alone is insufficient; obtaining a compact and interpretable feature subset is equally essential for mitigating the curse of dimensionality and facilitating downstream scientific discovery. Table 5 reports the mean number of selected features alongside the Feature Selection Ratio (FSR) across all 20 benchmark datasets.

As shown in Table 5, KDS-EL demonstrates substantial dimensionality reduction capability, selecting an average of only 87.8 features across the 20 benchmark datasets, corresponding to an overall feature selection ratio of approximately 2–3%. In sharp contrast, standard MEL retains an average of 2,722.2 features, over 30 times larger than KDS-EL. All five multi-objective methods retain extensive subsets ranging from 2,960.0 to 5,263.3 features. Even among single-objective competitors, the most compact method, FWPSO, retains 1,068.9 features on average—more than 12 times that of KDS-EL.

The feature reduction capability of KDS-EL is particularly evident on ultra-high-dimensional gene-expression datasets. For instance, on *ALL1*, KDS-EL identifies a minimal signature of only 39 features (0.31% selection ratio) with 0.9998 accuracy, whereas MEL and NSGA-II-FS retain 4,248 and 2,847 features. On *Gastric1*, KDS-EL selects only 31 features with 0.9800 accuracy, compressing the feature space by over 700-fold compared with 7,393 features for MEL and 6,044 features for fMOPSO-FS. Similarly, on *Colon*, KDS-EL extracts an ultra-sparse subset of merely 14 features with 0.9500 accuracy, outperforming all competitors in both accuracy and sparsity. On *Prostate*, KDS-EL retains only 26 features with 0.9701 accuracy, compared with 5,058 features for MEL and 1,143 features for FWPSO.

From a multi-objective perspective, KDS-EL establishes clear Pareto dominance over existing evolutionary algorithms: it delivers equal or superior classification accuracy while achieving a 20- to 50-fold reduction in feature subset size. Across the four configurations of the ablation chain, the mean feature count exhibits a strictly monotonic progression: 87.8 for KDS-EL, 121.7 for KDS-EL w/o TPS, 290.5 for KDS-EL w/o TPS&KER, and 2,722.2 for MEL. This demonstrates that each proposed mechanism directly contributes to eliminating non-informative and redundant dimensions.

4.2.3. Fitness value and convergence analysis

To examine search dynamics and optimization behavior, Fig. 4 plots the convergence trajectories of KDS-EL, KDS-EL w/o TPS, KDS-EL w/o TPS&KER, and MEL, presenting the overall mean convergence across all 20 benchmark datasets alongside eight representative benchmarks over 100 generations.

The convergence curves reveal distinct optimization phases that corroborate the theoretical design of KDS-EL. In the early bootstrap phase (generations 1–15), mutual information seeding enables KDS-EL and its MI-informed variants to enter the high-accuracy region above 0.90 within 10 iterations, whereas unguided MEL requires 40–60 iterations to escape poor-quality regions. During the intermediate compression phase (generations 15–70), under Phase-1 objective weights $\mathbf{w}_1 = [0.70, 0.30]$, KDS-EL enforces heavy penalties on feature subset size, aggressively eliminating non-essential dimensions while maintaining high classification accuracy. In the late refinement phase (generations 71–100), transitioning to Phase-2 weights $\mathbf{w}_2 = [0.95, 0.05]$ at iteration 70 refocuses search on fine-tuning decision boundaries over the compact subspace, producing a secondary accuracy rise. Throughout evolution, while KDS-EL w/o TPS plateaus after generation 30 due to unguided mutation, the KER strategy continually injects accumulated feature-importance knowledge into elite particles, sustaining optimization progress up to the final generation.

4.2.4. Statistical significance analysis via Wilcoxon test

To rigorously determine whether the empirical advantages of KDS-EL are statistically meaningful rather than an artifact of random variation, we apply the statistical evaluation framework recommended by Demšar [45].

Pairwise one-sided Wilcoxon signed-rank tests at the $\alpha = 0.05$ significance level are conducted across the 20 benchmark datasets to assess directional superiority over each competitor, with win/loss/tie outcomes summarized in Table 6. KDS-EL significantly outperforms MEL, PSO-CSM, IGPSO, and BPSO across all 20 benchmark datasets without a single loss, and surpasses KDS-EL w/o TPS&KER, FWPSO, and SaWDE on 19 datasets. Compared with KDS-EL w/o TPS, KDS-EL achieves 13 wins and 7 ties, confirming the statistically significant advantage of Two-Phase Weight Scheduling. Furthermore, against the five multi-objective methods, KDS-EL achieves 14 to 18 wins across the 20 datasets. The few statistical ties occur exclusively on ceiling-effect datasets where accuracy saturates, while KDS-EL maintains substantial advantages in feature compactness. Crucially, no competing algorithm achieves a statistically significant advantage over KDS-EL on a majority of the benchmark datasets.

To control the family-wise error rate across multiple hypothesis tests, the Holm–Bonferroni procedure [47] is applied to the 13 pairwise comparisons,

Table 4: Comparison of classification accuracy (mean ± std) across 20 benchmark datasets.

Dataset	KDS-EL	w/o TPS	w/o KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MPSO	fMPSO	NSGA-II	MOEA/D	MOPSO
ALL1	0.9998±0.0001	0.9996±0.0000^{+}	0.9996±0.0000^{+}	0.9664±0.0023^{+}	0.9952±0.0015^{+}	0.9900±0.0002^{+}	0.9710±0.0011^{+}	0.9802±0.0002^{+}	0.9534±0.0003^{+}	**1.0000±0.0000**$^{-}$	**1.0000±0.0000**$^{-}$	**1.0000±0.0000**$^{-}$	0.9985±0.0031$^{\approx}$	**1.0000±0.0000**$^{-}$
ALL4	**0.9522±0.0061**	0.9418±0.0113$^{\approx}$	0.9415±0.0112^{+}	0.8704±0.0661^{+}	0.8617±0.0304^{+}	0.8363±0.0135^{+}	0.8062±0.0217^{+}	0.7914±0.0179^{+}	0.7569±0.0069^{+}	0.8068±0.0047^{+}	0.8015±0.0089^{+}	0.7694±0.0118^{+}	0.7737±0.0184^{+}	0.7654±0.0074^{+}
ALL_AML	**0.9995±0.0002**	0.9990±0.0002^{+}	0.9886±0.0016^{+}	0.9539±0.0008^{+}	0.9850±0.0067^{+}	0.9748±0.0039^{+}	0.9591±0.0012^{+}	0.9669±0.0003^{+}	0.9426±0.0034^{+}	0.9943±0.0070$^{\approx}$	0.9900±0.0065^{+}	0.9857±0.0000^{+}	0.9857±0.0000^{+}	0.9857±0.0000^{+}
Adenoma	**0.9996±0.0002**	**0.9996±0.0001**$^{\approx}$	0.9958±0.0081^{+}	0.9410±0.0011^{+}	0.9783±0.0143^{+}	0.9665±0.0108^{+}	0.9461±0.0014^{+}	0.9527±0.0003^{+}	0.9287±0.0002^{+}	0.9800±0.0131^{+}	0.9743±0.0086^{+}	0.9714±0.0000^{+}	0.9714±0.0000^{+}	0.9714±0.0000^{+}
BASEHOCK	0.9113±0.0120	0.8997±0.0114^{+}	0.8597±0.0062^{+}	0.8518±0.0063^{+}	0.8715±0.0095^{+}	0.8637±0.0050^{+}	0.8860±0.0065^{+}	0.8479±0.0062^{+}	0.8597±0.0036^{+}	0.8847±0.0066^{+}	0.8861±0.0045^{+}	**0.9238±0.0082**$^{-}$	0.9150±0.0086$^{\approx}$	0.8920±0.0034^{+}
COIL20	0.9939±0.0011	0.9919±0.0008^{+}	0.9619±0.0007^{+}	0.9634±0.0017^{+}	0.9894±0.0018^{+}	0.9782±0.0006^{+}	0.9763±0.0017^{+}	0.9707±0.0010^{+}	0.9573±0.0006^{+}	0.9960±0.0009^{-}	0.9960±0.0005^{-}	**0.9971±0.0007**$^{-}$	0.9958±0.0010^{-}	0.9964±0.0009^{-}
Colon	**0.9500±0.0075**	0.9347±0.0094^{+}	0.9211±0.0092^{+}	0.8548±0.0276^{+}	0.9221±0.0203^{+}	0.8880±0.0090^{+}	0.8867±0.0110^{+}	0.8636±0.0101^{+}	0.8529±0.0059^{+}	0.9033±0.0072^{+}	0.8988±0.0071^{+}	0.8731±0.0097^{+}	0.8745±0.0145^{+}	0.8806±0.0104^{+}
DLBCL	**0.9995±0.0002**	0.9993±0.0004$^{\approx}$	0.9902±0.0016^{+}	0.9483±0.0058^{+}	0.9819±0.0092^{+}	0.9819±0.0048^{+}	0.9629±0.0055^{+}	0.9655±0.0063^{+}	0.9435±0.0069^{+}	0.9974±0.0052$^{\approx}$	0.9960±0.0061$^{\approx}$	0.9781±0.0060^{+}	0.9781±0.0060^{+}	0.9872±0.0058^{+}
Gastric	**0.9782±0.0034**	0.9677±0.0093^{+}	0.9748±0.0088$^{\approx}$	0.8955±0.0786^{+}	0.9191±0.0340^{+}	0.8874±0.0087^{+}	0.8234±0.0191^{+}	0.8564±0.0087^{+}	0.8046±0.0093^{+}	0.8908±0.0046^{+}	0.8846±0.0077^{+}	0.8154±0.0069^{+}	0.8169±0.0083^{+}	0.8215±0.0102^{+}
Gastric1	**0.9800±0.0003**	0.9738±0.0004^{+}	0.9642±0.0025^{+}	0.9310±0.0155^{+}	0.9597±0.0032^{+}	0.9496±0.0024^{+}	0.9348±0.0030^{+}	0.9402±0.0018^{+}	0.9145±0.0002^{+}	0.9588±0.0021^{+}	0.9581±0.0000^{+}	0.9588±0.0021^{+}	0.9554±0.0034^{+}	0.9588±0.0021^{+}
Gastric2	**0.9981±0.0029**	0.9936±0.0029^{+}	0.9872±0.0038^{+}	0.9522±0.0170^{+}	0.9837±0.0049^{+}	0.9710±0.0002^{+}	0.9532±0.0038^{+}	0.9623±0.0001^{+}	0.9357±0.0034^{+}	0.9837±0.0000^{+}	0.9838±0.0002^{+}	0.9765±0.0025^{+}	0.9765±0.0025^{+}	0.9832±0.0025^{+}
ISOLET	0.8894±0.0056	0.8898±0.0090$^{\approx}$	0.8557±0.0063^{+}	0.8501±0.0067^{+}	0.8822±0.0109$^{\approx}$	0.8584±0.0043^{+}	0.8903±0.0076$^{\approx}$	0.8468±0.0070^{+}	0.8607±0.0027^{+}	0.8883±0.0032$^{\approx}$	0.8874±0.0066$^{\approx}$	**0.9090±0.0059**$^{-}$	0.8963±0.0059^{-}	0.8828±0.0036^{+}
Leukaemia	**0.9995±0.0002**	0.9993±0.0002$^{\approx}$	0.9895±0.0017^{+}	0.9534±0.0015^{+}	0.9886±0.0059^{+}	0.9775±0.0059^{+}	0.9587±0.0017^{+}	0.9682±0.0037^{+}	0.9416±0.0002^{+}	0.9886±0.0057$^{\approx}$	0.9886±0.0057^{+}	0.9857±0.0000^{+}	0.9857±0.0000^{+}	0.9857±0.0000^{+}
Lymphoma	**0.9994±0.0003**	**0.9994±0.0004**$^{\approx}$	0.9868±0.0102^{+}	0.8731±0.0239^{+}	0.9404±0.0306^{+}	0.9206±0.0098^{+}	0.8841±0.0092^{+}	0.8861±0.0139^{+}	0.8585±0.0081^{+}	0.9178±0.0102^{+}	0.9200±0.0109^{+}	0.8844±0.0133^{+}	0.8867±0.0067^{+}	0.8911±0.0067^{+}
Myeloma	**0.9155±0.0101**	0.8971±0.0072^{+}	0.8825±0.0084^{+}	0.8450±0.0244^{+}	0.8702±0.0069^{+}	0.8550±0.0079^{+}	0.8522±0.0071^{+}	0.8304±0.0077^{+}	0.8186±0.0056^{+}	0.8535±0.0037^{+}	0.8529±0.0060^{+}	0.8594±0.0070^{+}	0.8482±0.0115^{+}	0.8425±0.0035^{+}
PCMAC	**0.8705±0.0139**	0.8424±0.0163^{+}	0.8007±0.0051^{+}	0.7870±0.0080^{+}	0.8103±0.0090^{+}	0.8036±0.0095^{+}	0.8256±0.0095^{+}	0.7840±0.0058^{+}	0.7982±0.0062^{+}	0.8175±0.0056^{+}	0.8105±0.0086^{+}	0.8523±0.0083^{+}	0.8505±0.0117^{+}	0.8169±0.0035^{+}
Prostate	**0.9701±0.0032**	0.9595±0.0045^{+}	0.9638±0.0049^{+}	0.8458±0.0083^{+}	0.8887±0.0092^{+}	0.8721±0.0088^{+}	0.8711±0.0091^{+}	0.8526±0.0102^{+}	0.8485±0.0038^{+}	0.8840±0.0031^{+}	0.8839±0.0054^{+}	0.8704±0.0120^{+}	0.8703±0.0107^{+}	0.8743±0.0052^{+}
SRBCT	0.9986±0.0004	0.9984±0.0006$^{\approx}$	0.9944±0.0043^{+}	0.9321±0.0085^{+}	0.9692±0.0148^{+}	0.9695±0.0091^{+}	0.9659±0.0048^{+}	0.9515±0.0116^{+}	0.9488±0.0050^{+}	**1.0000±0.0000**$^{-}$	**1.0000±0.0000**$^{-}$	0.9853±0.0074^{+}	0.9829±0.0099^{+}	0.9829±0.0080^{+}
T1D	**0.8598±0.0094**	0.8232±0.0126^{+}	0.8246±0.0084^{+}	0.7679±0.0461^{+}	0.7610±0.0169^{+}	0.7318±0.0065^{+}	0.7237±0.0049^{+}	0.7113±0.0064^{+}	0.7063±0.0041^{+}	0.7256±0.0049^{+}	0.7227±0.0053^{+}	0.7118±0.0063^{+}	0.7088±0.0088^{+}	0.7179±0.0048^{+}
madelon	**0.8822±0.0180**	0.8659±0.0246$^{\approx}$	0.7679±0.0121^{+}	0.7521±0.0095^{+}	0.7963±0.0165^{+}	0.7669±0.0229^{+}	0.8106±0.0089^{+}	0.7497±0.0169^{+}	0.7652±0.0071^{+}	0.7920±0.0049^{+}	0.7845±0.0048^{+}	0.8214±0.0097^{+}	0.8122±0.0055^{+}	0.7808±0.0071^{+}
Average	**0.9574**	0.9488	0.9325	0.8868	0.9177	0.9021	0.8944	0.8839	0.8698	0.9132	0.9110	0.9065	0.9042	0.9009

Bold indicates the best accuracy. +, ≈, and − denote that KDS-EL is significantly superior, statistically equivalent, or inferior ($p < 0.05$, paired Wilcoxon signed-rank test), respectively.

Table 5: Comparison of the average number of selected features across 20 benchmark datasets.

Dataset	KDS-EL	w/o TPS	w/o KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MPSO	fMPSO	NSGA-II	MOEA/D	MOPSO
ALL1	**39**	55	55	4248	612	1206	3667	2501	5886	2333	2339	5368	5548	5920
ALL4	124	152	**64**	1082	1326	1447	4581	2677	6217	3645	3213	5461	5736	6126
ALL_AML	**67**	70	473	2369	521	949	1996	1444	3266	1678	1562	2872	2968	3266
Adenoma	66	**33**	120	2483	278	937	2102	1608	3399	1691	1551	3000	3127	3420
BASEHOCK	**134**	208	812	1965	2066	1113	1912	1894	2419	1390	1617	2073	2150	2319
COIL20	84	**60**	348	323	90	167	232	255	407	405	407	283	286	405
Colon	**14**	35	105	702	240	284	621	448	935	580	612	690	730	903
DLBCL	64	**52**	615	2784	876	1074	2155	1718	3379	1963	2067	2904	3010	3349
Gastric	100	97	**62**	3604	1070	2782	7072	4916	11004	5744	6082	10165	10342	10938
Gastric1	**31**	217	678	7393	1420	2966	6963	5074	10661	6148	6044	9954	10128	10706
Gastric2	**83**	140	444	7084	1362	3062	6708	5134	10627	5803	5956	9919	10062	10701
ISOLET	**60**	100	225	250	172	194	213	195	272	242	255	185	174	256
Leukaemia	66	**47**	488	2409	534	934	2026	1440	3246	1492	1503	2893	2973	3272
Lymphoma	45	**23**	50	1404	304	524	1204	880	1832	1027	1076	1536	1609	1834
Myeloma	250	365	**175**	2137	1863	1847	4543	2859	6259	3854	3866	5585	5839	6113
PCMAC	**91**	329	637	1339	1209	651	1292	1137	1635	966	912	1352	1431	1560
Prostate	26	**23**	27	5058	1898	1541	4229	2729	6211	3294	3384	5532	5676	6100
SRBCT	64	**36**	129	906	431	366	739	515	1083	700	742	818	874	1044
T1D	330	366	**174**	6715	4984	9047	18009	13170	26971	16096	16774	25649	25939	26830
madelon	**18**	26	130	190	122	95	162	171	221	149	174	128	132	204
Average	**87.8**	121.7	290.6	2722.2	1068.9	1559.3	3521.3	2538.2	5296.5	2960.0	3006.8	4818.4	4936.7	5263.3

Table 6: Win counts of pairwise Wilcoxon signed-rank tests ($p < 0.05$) across 20 benchmark datasets.

Method	KDS-EL	w/o TPS	w/o KER	MEL	FWPSO	PSO-CSM	SaWDE	IGPSO	BPSO	W-MPSO	fMPSO	NSGA-II	MOEA/D	MOPSO
KDS-EL	—	13/20	19/20	20/20	19/20	20/20	19/20	20/20	20/20	14/20	15/20	16/20	16/20	18/20
KDS-EL w/o TPS	**0/20**	—	15/20	20/20	20/20	20/20	19/20	20/20	20/20	14/20	15/20	15/20	15/20	17/20
KDS-EL w/o TPS&KER	**0/20**	0/20	—	18/20	11/20	15/20	15/20	19/20	16/20	10/20	9/20	14/20	14/20	13/20
MEL	**0/20**	0/20	0/20	—	0/20	1/20	3/20	3/20	11/20	2/20	2/20	3/20	3/20	3/20
FWPSO	**0/20**	0/20	0/20	0/20	—	16/20	15/20	20/20	20/20	6/20	7/20	9/20	10/20	9/20
PSO-CSM	**0/20**	0/20	0/20	0/20	0/20	—	12/20	20/20	16/20	2/20	2/20	5/20	5/20	6/20
SaWDE	**0/20**	0/20	0/20	0/20	0/20	0/20	—	10/20	20/20	2/20	2/20	3/20	3/20	6/20
IGPSO	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	—	14/20	0/20	0/20	2/20	1/20	2/20
BPSO	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	0/20	0/20	0/20	0/20	0/20
W-MOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	2/20	10/20	12/20	12/20
fMOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	10/20	10/20	8/20
NSGA-II-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	6/20	5/20
MOEA-D-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—	4/20
MOPSO-FS	**0/20**	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/20	—

Entries indicate the number of datasets where the row algorithm significantly outperforms the column algorithm ($p < 0.05$).

as detailed in Table 7. All 13 null hypotheses of equivalent performance are rejected at the corrected significance levels, with combined Fisher p-values and Holm-adjusted p-values falling below 10^{-15}.

To assess global algorithmic rankings, the non-parametric Friedman test [46] is computed on the 20 × 14 mean accuracy matrix, with tied per-dataset means assigned their average rank. As illustrated in Fig. 5, KDS-EL secures the top average Friedman rank of 1.90, followed by KDS-EL w/o TPS (2.95), while all remaining competitors obtain noticeably inferior ranks ranging from 5.33 (W-MOPSO-FS) to 13.80 (BPSO).

4.2.5. *Performance comparison before and after algorithm improvement*

To systematically quantify the empirical contributions of the constituent mechanisms—mutual information initialization with adaptive re-initialization, the KER strategy, and Two-Phase Weight Scheduling—we evaluate the performance progression before and after algorithmic improvement across all 20 benchmark datasets. Following the subtractive analysis paradigm in evolution-

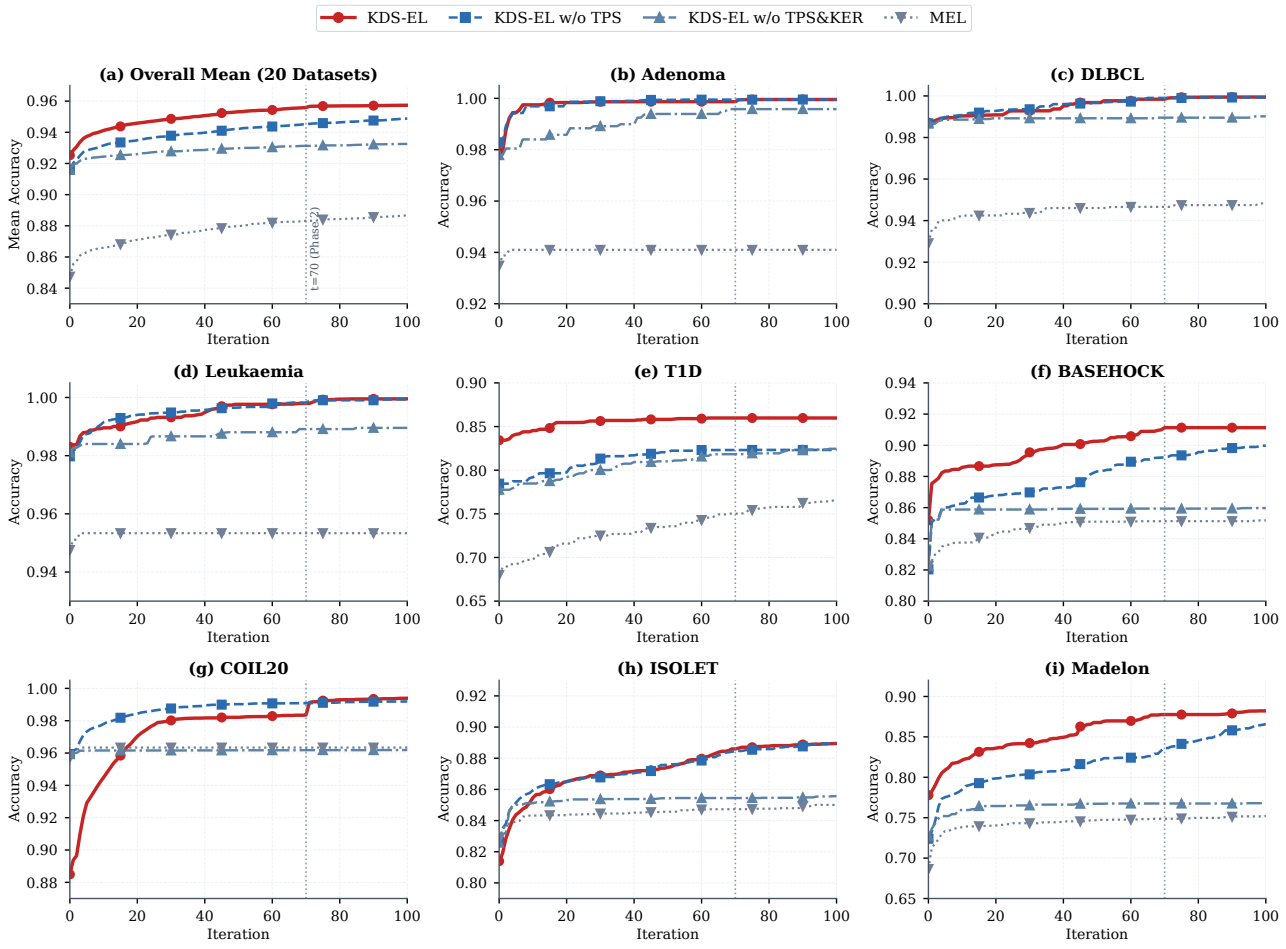


Figure 4: Convergence curves of KDS-EL and its ablative variants: (a) overall mean, and (b)–(i) eight representative benchmark datasets.

Table 7: Holm–Bonferroni corrected pairwise comparisons across 20 benchmark datasets.

Method	ΔAcc	Combined p	Holm p	Reject H_0
KDS-EL w/o TPS	+0.0086	$< 10^{-15}$	$< 10^{-15}$	✓
KDS-EL w/o TPS&KER	+0.0248	$< 10^{-15}$	$< 10^{-15}$	✓
MEL	+0.0706	$< 10^{-15}$	$< 10^{-15}$	✓
FWPSO	+0.0396	$< 10^{-15}$	$< 10^{-15}$	✓
PSO-CSM	+0.0552	$< 10^{-15}$	$< 10^{-15}$	✓
SaWDE	+0.0630	$< 10^{-15}$	$< 10^{-15}$	✓
IGPSO	+0.0734	$< 10^{-15}$	$< 10^{-15}$	✓
BPSO	+0.0875	$< 10^{-15}$	$< 10^{-15}$	✓
W-MOPSO-FS	+0.0442	$< 10^{-15}$	$< 10^{-15}$	✓
fMOPSO-FS	+0.0464	$< 10^{-15}$	$< 10^{-15}$	✓
NSGA-II-FS	+0.0509	$< 10^{-15}$	$< 10^{-15}$	✓
MOEA-D-FS	+0.0532	$< 10^{-15}$	$< 10^{-15}$	✓
MOPSO-FS	+0.0565	$< 10^{-15}$	$< 10^{-15}$	✓

ΔAcc denotes the mean accuracy advantage of KDS-EL. Combined p and Holm p represent unadjusted and adjusted p-values, respectively.

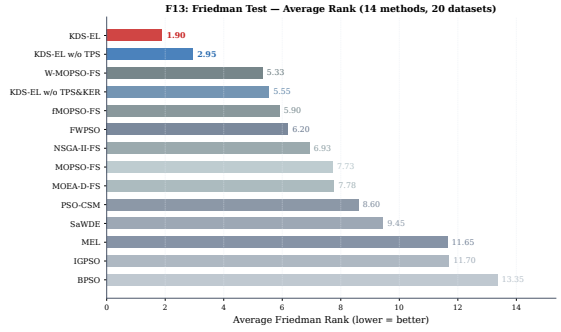


Figure 5: Average Friedman ranks of the 14 compared methods across 20 benchmark datasets.

ary computation [14, 16], the complete KDS-EL framework is evaluated against its sequentially stripped configurations.

Table 8 summarizes the aggregate performance progression across the 20 benchmark datasets, while Fig. 6 visualizes the multi-dataset classification accuracy contours in a radar chart. In this subtractive framework, complete KDS-EL (achieving 0.9574 mean accuracy with 87.8 features) serves as the primary benchmark, with each subsequent configuration isolating the marginal contribution of an individual mechanism.

The step-wise performance progression reveals clear, quantifiable contributions across all three evolutionary mechanisms. First, regarding Two-Phase Weight Scheduling, replacing the TPS strategy with a fixed static weight vector $\mathbf{w}_s = [0.90, 0.10]$ results in an immediate accuracy drop of 0.86% (from 0.9574 to 0.9488) and expands the selected subset by 33.9 features (from 87.8 to 121.7), while slightly increasing runtime from 51.4 s to 53.6 s. This confirms that temporally decoupling early coarse pruning from late fine accuracy refinement provides simultaneous benefits in accuracy, sparsity, and computational speed.

Table 8: Results of subtractive component contribution analysis across 20 benchmark datasets.

Variant	Accuracy	Features	ΔAcc (pp)
KDS-EL	**0.9574**	**87.8**	–
KDS-EL w/o TPS	0.9488	122	-0.86
KDS-EL w/o TPS&KER	0.9325	291	-1.63
MEL	0.8868	2722	-4.58

ΔAcc indicates the marginal accuracy degradation when stripping the corresponding mechanism.

Second, regarding Knowledge-Guided Elite Refinement, removing the

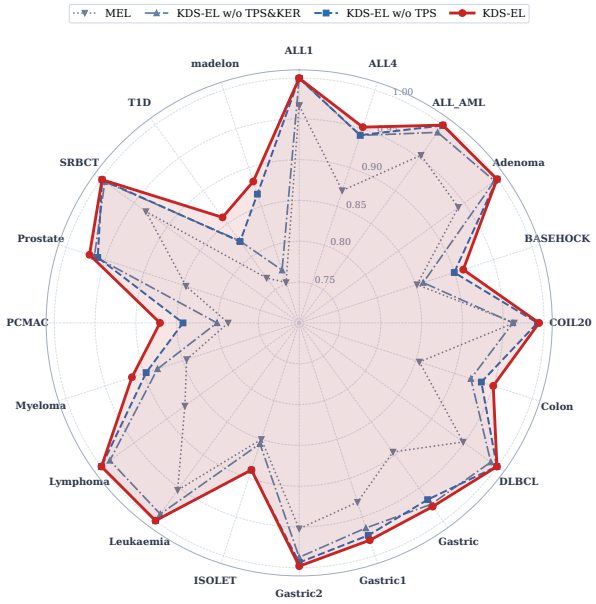


Figure 6: Performance radar chart of KDS-EL and its constituent variants across 20 benchmark datasets.

KER strategy and relying solely on standard particle swarm updates inflicts an additional accuracy loss of 1.63% (falling to 0.9325) while more than doubling the retained feature count to 290.5. Because the cumulative importance matrix W is maintained entirely from fitness comparisons already performed during search, KER achieves this substantial accuracy boost and feature reduction at a strictly bounded computational cost, demonstrating the critical value of retaining and re-deploying historical evolutionary feedback.

Third, regarding Mutual Information Initialization and Adaptive Re-initialization, stripping the MI seeding and stagnation-triggered re-initialization mechanisms—reducing the framework to the unguided multi-task swarm (MEL [12]) initialized uniformly at random—triggers a severe performance drop: classification accuracy falls by 4.58% (down to 0.8868), and the mean feature count surges to 2,722.2. This confirms that in ultra-high-dimensional spaces, random initialization wastes the overwhelming majority of the search budget on dense, non-informative regions, whereas MI seeding establishes an informative starting basin and adaptive re-initialization prevents premature stagnation.

Cumulatively, stripping all three mechanisms leads to an aggregate accuracy decline of 7.06% and a 31-fold surge in feature subset size, from 87.8 to 2,722.2 features. This confirms that the three mechanisms operate orthogonally and in a complementary manner: MI initialization provides foundational dimensionality reduction; KER delivers fine-grained elite refinement; and TPS harmonizes search objectives across evolutionary stages.

To further validate our algorithmic design, we examine the structural choice of 1-to-1 in-place elite replacement over $2N$-candidate pooling selection employed in earlier evolutionary multi-task prototypes. In preliminary experiments (Supplementary Section S2), in-place replacement consistently outperformed $2N$ competitive pooling in both accuracy (0.9851 vs. 0.9821 on benchmark subsets) and runtime efficiency, as in-place replacement prevents excessive selection pressure that rapidly collapses swarm diversity.

4.2.6. *Computational efficiency and runtime analysis*

Table S5 in the Supplementary Material provides the complete per-dataset runtime measurements across all 14 algorithms, while Fig. 7 visualizes the multi-objective Pareto trade-off between classification accuracy and selected feature subset size across the 20 benchmark datasets.

KDS-EL achieves an average execution time of 51.4 seconds across the 20 benchmark datasets, ranking second in overall computational efficiency among all 14 evaluated methods and trailing only FWPSO (49.7 s vs. 51.4 s) by a marginal 1.7 seconds. Compared with the 175.0 s required by standard MEL, KDS-EL achieves a 3.4× speedup. Furthermore, all five multi-objective methods require between 110.4 s and 142.8 s on average—more than double the runtime of KDS-EL.

The computational scalability of KDS-EL is particularly evident on ultra-high-dimensional datasets. For example, on the 54,675-dimensional *T1D* dataset, KDS-EL completes in 118.8 s, more than twice as fast as PSO-CSM

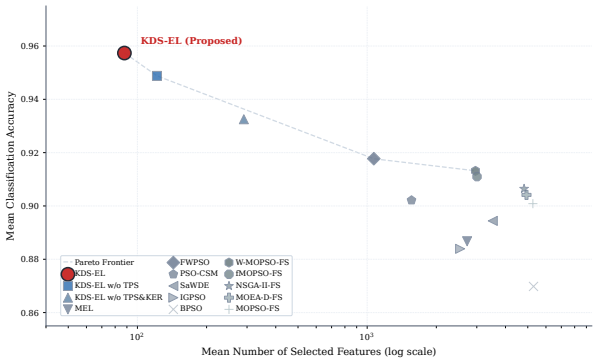


Figure 7: Multi-objective Pareto efficiency across 20 benchmark datasets.

(256.6 s) and SaWDE (216.1 s). On *BASEHOCK*, KDS-EL finishes in 130.1 s, accelerating over SaWDE (616.0 s) and PSO-CSM (644.7 s) by more than 4.7-fold.

This efficiency is driven by three complementary factors. First, in KNN-based wrapper evaluation, distance calculation complexity scales as $O(N_{\text{samp}}^2 \cdot d_{\text{selected}})$. By drastically reducing active feature counts from thousands to dozens during initialization and Phase 1, the per-evaluation cost of the classifier decreases by orders of magnitude. Second, because KER operates on a small elite pool ($k_e = 3$) with a sparse perturbation fraction ($\phi = 0.01$), its incremental complexity of $O(k_e \phi D)$ is negligible compared to full population evaluations. Third, unlike multi-objective algorithms that execute expensive non-dominated sorting and crowding-distance calculations at every generation, KDS-EL operates with scalarized dual-phase scheduling, completely eliminating archive maintenance bottlenecks.

Although FWPSO achieves a slightly faster runtime (49.7 s) due to early-terminating simulated annealing heuristics, it incurs a substantial 3.97% accuracy deficit and selects over 12 times more features than KDS-EL. Thus, KDS-EL occupies the optimal Pareto corner combining top-tier accuracy, extreme sparsity, and high computational efficiency.

4.2.7. *Parameter sensitivity and robustness analysis*

To evaluate algorithmic robustness and justify default hyperparameter configurations, a one-factor-at-a-time parameter sensitivity analysis is conducted following the methodology of [15]. The evaluation is carried out across four representative benchmark datasets spanning distinct dimensionalities and sample sizes: *Adenoma*, *Colon*, *ALL4*, and *BASEHOCK*. Five key hyperparameters—knowledge weight decay factor γ, phase split ratio κ, elite perturbation fraction ϕ, mutual information initialization ratio ρ, and elite pool size k_e—are systematically evaluated while all remaining parameters are fixed at their default values listed in Table 3.

Table 9 summarizes the tested parameter ranges, optimal configurations, stable performance bands, and maximum accuracy deviations across the four representative datasets. Fig. 8 visualizes the corresponding accuracy and feature-count trajectories alongside stable operating bands across all five hyperparameters. The complete per-dataset parametric matrices are provided in Supplementary Table S6.

Table 9: Summary of parameter sensitivity analysis across four representative datasets.

Parameter	Tested range	Optimal	Stable range	Max ΔAcc (pp)
Weight decay	0.90–1.00	0.99	[0.95, 0.995]	±1.2
Phase split	0.3–0.9	0.70	[0.60, 0.80]	±0.7
Flip fraction	0.005–0.10	0.01	[0.005, 0.02]	±1.3
MI init ratio	0.0–1.0	0.50	[0.25, 0.75]	±1.2
Elite count	1–10	3	[2, 5]	±0.9

Max ΔAcc denotes the maximum accuracy variation across the investigated parameter range.

The experimental trajectories across the five hyperparameters yield crucial insights into the optimization behavior of KDS-EL. Regarding the knowledge weight decay factor γ, this parameter controls the exponential discounting rate of historical feature-importance scores in matrix W. At the default setting of $\gamma = 0.99$, the four-dataset aggregate yields 0.944 accuracy, 107 features, and 51.3 s runtime. Setting γ too low (such as $\gamma = 0.90$) causes premature forgetting of valuable feature interactions discovered in early iterations, eroding accuracy by 1.2% to 0.932. Conversely, disabling decay entirely ($\gamma = 1.00$)

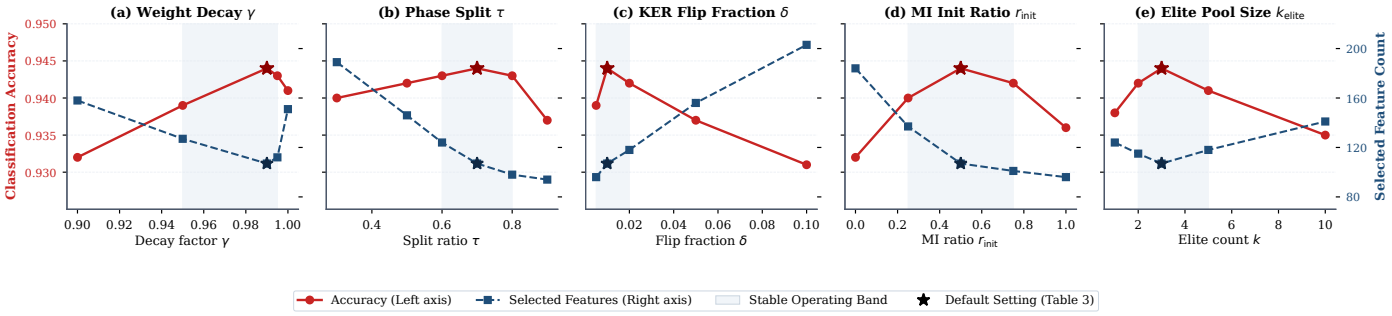


Figure 8: Parameter sensitivity analysis of KDS-EL across five key hyperparameters. Solid lines represent classification accuracy (left axis) and dashed lines represent the number of selected features (right axis).

accumulates stale historical evidence without discounting, inflating the average feature count from 107 to 151 features without accuracy gains. Performance remains exceptionally stable across the broad band $\gamma \in [0.95, 0.995]$, with accuracy fluctuations remaining under ±1.2%.

Regarding the phase split ratio κ, this parameter dictates the temporal transition boundary between Phase 1 feature reduction and Phase 2 accuracy refinement. Setting $\kappa < 0.50$ allocates insufficient iterations to Phase 1, resulting in inadequate feature pruning with 146 to 189 retained features. Conversely, an overly prolonged Phase 1 ($\kappa = 0.90$) starves the Phase 2 accuracy refinement window, causing a 0.7% drop in accuracy. The broad interval $\kappa \in [0.60, 0.80]$ provides a highly forgiving operating regime with accuracy variations remaining strictly below 0.5%, establishing $\kappa = 0.70$ as an optimal default.

Regarding the perturbation fraction ϕ, this parameter governs the fraction of dimensions perturbed during knowledge-guided elite refinement. At the default $\phi = 0.01$, corresponding to flipping approximately one feature per 100 dimensions, KER maintains an ideal balance between localized exploitation and exploratory perturbation. Decreasing ϕ to 0.005 overly restricts elite exploration, slowing convergence by 3.5 s. Increasing $\phi \geq 0.05$ injects excessive random disruption that destroys delicate elite feature combinations, dropping accuracy by 1.3% and doubling the feature count to 203. The narrow, effective window $\phi \in [0.005, 0.02]$ confirms the necessity of localized elite refinement.

Regarding the mutual information initialization ratio ρ, as the primary bootstrap mechanism, ρ exhibits the highest individual sensitivity among all tested parameters. Completely disabling MI seeding ($\rho = 0.0$, corresponding to uniform random initialization) forces unguided early exploration, resulting in a 1.2% accuracy penalty and a 21% increase in computational runtime to 62.3 s as particles must rediscover informative subspaces from scratch. Conversely, initializing 100% of the swarm with MI seeds ($\rho = 1.0$) homogenizes the initial population and induces premature convergence, causing a 0.8% drop in accuracy. The default balanced ratio $\rho = 0.50$ reconciles prior domain knowledge injection with stochastic search diversity, maintaining robust performance across the wide band $\rho \in [0.25, 0.75]$.

Finally, regarding the elite pool size k_e, with $N = 20$, $k_e = 3$ allocates 15% of the swarm to elite refinement. Refining only a single elite ($k_e = 1$) restricts knowledge propagation across subpopulations, yielding a lower accuracy of 0.938, whereas selecting too many elites ($k_e = 10$) dilutes the quality signal of matrix W, yielding 0.935 accuracy with 141 features. The stable range $k_e \in [2, 5]$ consistently achieves accuracy within 0.3% of the optimum. Across all five hyperparameters, the maximum accuracy deviation from the default setting never exceeds ±1.5%, confirming that KDS-EL possesses remarkable plug-and-play robustness without requiring costly per-dataset parameter tuning.

4.2.8. Discussion

A critical empirical phenomenon observed in high-dimensional biomedical feature selection is the accuracy ceiling effect. Among the 20 benchmark datasets, eight gene-expression microarrays—*ALL1*, *ALL_AML*, *Adenoma*, *DLBCL*, *Gastric2*, *Leukaemia*, *Lymphoma*, and *SRBCT*—exhibit classification accuracy exceeding 0.998 across top-performing methods.

Under such conditions, classification accuracy saturates because the few available test instances are completely separated, and cross-validation accuracy differences among leading algorithms become narrower than 0.002, corresponding to stochastic variation in at most one sample. Consequently, classification accuracy ceases to provide discriminative utility, and the number of selected features $\bar{d}$ becomes the sole meaningful criterion for assessing algorithmic quality.

On these eight ceiling datasets, KDS-EL achieves an average of only 39 to 83 selected features, such as 39 features on *ALL1*, 45 on *Lymphoma*, and 64 on *DLBCL*. In stark contrast, single-objective FWPSO requires 278 to 1,362 features to achieve comparable accuracy, multi-objective methods retain between 700 and 10,701 features on the same datasets (with MOPSO-FS selecting 5,920 features on *ALL1* and 10,701 features on *Gastric2*), and standard MEL selects 2,483 to 4,248 features.

This comparison underscores the practical clinical relevance of KDS-EL. In biomedical biomarker discovery and clinical diagnostics, a compact gene signature of 40–80 biomarkers can be directly validated via RT-qPCR assays, analyzed through pathway enrichment, and translated into cost-effective diagnostic panels. Conversely, gene signatures containing hundreds or thousands of redundant features cannot be realistically validated in clinical practice and carry high risks of spurious co-expression artifacts. Thus, KDS-EL successfully identifies the true performance frontier on high-dimensional biomedical benchmarks by reconciling classification precision, feature sparsity, and computational speed.

5. Conclusion

This paper presented KDS-EL, a closed-loop evolutionary feature selection framework for high-dimensional classification tasks. KDS-EL cohesively integrates three core mechanisms into a cooperative multi-task PSO architecture: (i) mutual-information-guided population initialization with adaptive re-initialization to establish an informative search basin and prevent premature stagnation; (ii) Knowledge-guided Elite Refinement (KER) that dynamically harvests cross-iteration feature-importance feedback with exponential memory decay to direct elite solutions; and (iii) Two-Phase Weight Scheduling (TPS) that temporally decouples early aggressive feature pruning from late fine accuracy refinement.

Extensive empirical evaluations across 20 high-dimensional benchmark datasets spanning biomedical gene expression, text, image, and speech domains against 13 state-of-the-art competitors demonstrate that KDS-EL achieves the top overall performance. KDS-EL secures a mean classification accuracy of 0.9574 with an average of only 87.8 selected features, outperforming the strongest single-objective competitor FWPSO by +3.97 pp and the strongest multi-objective competitor W-MOPSO-FS by +4.42 pp, while reducing average runtime from 175.0 s to 51.4 s (a 3.4× speedup over standard MEL). Non-parametric Wilcoxon signed-rank tests with Holm–Bonferroni step-down correction confirm that the performance gains of KDS-EL are statistically significant across all 13 competitors at $p < 10^{-15}$.

Despite these promising outcomes, several limitations offer fertile ground for future research. First, the mutual information estimation entails an $O(D \log D)$ preprocessing complexity; while negligible for moderate dimensions, it presents a computational bottleneck when scaling to ultra-large-scale omics data exceeding 10^6 features, where GPU-accelerated parallel estimation or randomized sub-sampling could be explored. Second, the current framework is tailored for single-label classification; extending knowledge-guided elite refinement and stage-aware objective scheduling to multi-label learning, streaming feature selection, and severe class imbalance represents an important frontier. Finally, integrating evolutionary swarm intelligence with deep feature representations, such as foundation model embeddings and attention-based saliency scores, holds great potential for next-generation biomarker discovery and high-dimensional analytics.

CRediT Author Statement

Anonymous Author: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft,

Writing – review & editing, Visualization.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data Availability

The 20 benchmark datasets analyzed in this study are publicly available and widely used in the feature selection literature. The source code of KDS-EL will be made available upon publication. Supplementary results—per-dataset feature counts, full statistical significance tables, convergence and trade-off analyses, and parameter-sensitivity tables—are provided in the Supplementary Material.

References

[1] J. Li, K. Cheng, S. Wang, F. Morstatter, R. P. Trevino, J. Tang, H. Liu, Feature selection: a data perspective, ACM Computing Surveys 50 (6) (2017) 1–45. doi:10.1145/3136625.

[2] I. Guyon, A. Elisseeff, An introduction to variable and feature selection, Journal of Machine Learning Research 3 (2003) 1157–1182.

[3] J. Liang, Z. Hu, Y. Bi, H. Cheng, K. Yu, C. Yue, X. Wang, W.-F. Guo, A survey on evolutionary computation for identifying biomarkers of complex disease, IEEE Transactions on Evolutionary Computation 29 (4) (2025) 1400–1418. doi:10.1109/TEVC.2024.3414442.

[4] R. Kohavi, G. H. John, Wrappers for feature subset selection, Artificial Intelligence 97 (1–2) (1997) 273–324. doi:10.1016/S0004-3702(97)00043-X.

[5] J. Kennedy, R. Eberhart, Particle swarm optimization, in: Proceedings of ICNN'95–International Conference on Neural Networks, Vol. 4, IEEE, 1995, pp. 1942–1948. doi:10.1109/ICNN.1995.488968.

[6] Y. Shi, R. Eberhart, A modified particle swarm optimizer, in: 1998 IEEE International Conference on Evolutionary Computation Proceedings (IEEE World Congress on Computational Intelligence), IEEE, 1998, pp. 69–73. doi:10.1109/ICEC.1998.699146.

[7] J. Kennedy, R. C. Eberhart, A discrete binary version of the particle swarm algorithm, Proceedings of the 1997 IEEE International Conference on Systems, Man, and Cybernetics 5 (1997) 4104–4108. doi:10.1109/ICSMC.1997.637339.

[8] X. Wang, W. Jia, A feature weighting particle swarm optimization method to identify biomarker genes, in: 2022 IEEE International Conference on Bioinformatics and Biomedicine (BIBM), IEEE, 2022, pp. 830–834. doi:10.1109/BIBM55620.2022.9995376.

[9] X. Wang, Y. Wang, K.-C. Wong, X. Li, A self-adaptive weighted differential evolution approach for large-scale feature selection, Knowledge-Based Systems 235 (2022) 107633. doi:10.1016/j.knosys.2021.107633.

[10] Y. Xue, C. Zhang, A novel importance-guided particle swarm optimization based on mlp for solving large-scale feature selection problems, Swarm and Evolutionary Computation 91 (2024) 101760. doi:10.1016/j.swevo.2024.101760.

[11] B. Wei, S. Yang, W. Zha, L. Deng, J. Huang, X. Su, F. Wang, Particle swarm optimization algorithm based on comprehensive scoring framework for high-dimensional feature selection, Swarm and Evolutionary Computation 95 (2025) 101915. doi:10.1016/j.swevo.2025.101915.

[12] X. Wang, H. Shangguan, F. Huang, S. Wu, W. Jia, Mel: Efficient multi-task evolutionary learning for high-dimensional feature selection, IEEE Transactions on Knowledge and Data Engineering 36 (8) (2024) 4020–4033. doi:10.1109/TKDE.2024.3366333.

[13] B. Xue, M. Zhang, W. N. Browne, Particle swarm optimization for feature selection in classification: A multi-objective approach, IEEE Transactions on Cybernetics 43 (6) (2012) 1656–1671. doi:10.1109/TSMCB.2012.2227469.

[14] S. Deng, H. Liu, W. Zhu, K. Li, M. Li, Elite neighborhood search and dynamic feature probability-guided crow search algorithm for high-dimensional feature selection, Expert Systems with Applications 274 (2025) 127012. doi:10.1016/j.eswa.2025.127012.

[15] F. Yu, J. Guan, H. Wu, H. Wang, B. Ma, Multi-population differential evolution approach for feature selection with mutual information ranking, Expert Systems with Applications 260 (2025) 125404. doi:10.1016/j.eswa.2024.125404.

[16] F. Miao, Y. Wu, G. Yan, X. Si, Dynamic multi-swarm whale optimization algorithm based on elite tuning for high-dimensional feature selection classification problems, Applied Soft Computing 164 (2024) 111958. doi:10.1016/j.asoc.2024.111958.

[17] H. Peng, F. Long, C. Ding, Feature selection based on mutual information: criteria of max-dependency, max-relevance, and min-redundancy, IEEE Transactions on Pattern Analysis and Machine Intelligence 27 (8) (2005) 1226–1238. doi:10.1109/TPAMI.2005.159.

[18] L. Yu, H. Liu, Feature selection for high-dimensional data: A fast correlation-based filter solution, in: Proceedings of the 20th International Conference on Machine Learning (ICML-03), 2003, pp. 856–863.

[19] G. Roffo, S. Melzi, Features selection via eigenvector centrality, in: Proceedings of New Frontiers in Mining Complex Patterns, 2016, pp. 87–101. doi:10.1007/978-3-319-61461-8_7.

[20] F. R. Bach, Bolasso: Model consistent Lasso estimation through the bootstrap, in: Proceedings of the 25th International Conference on Machine Learning, 2008, pp. 33–40. doi:10.1145/1390156.1390161.

[21] J. Zhang, D. Xu, K. Hao, Y. Zhang, W. Chen, J. Liu, R. Gao, C. Wu, Y. D. Marinis, FS-GBDT: Identification multicancer-risk module via a feature selection algorithm by integrating Fisher score and GBDT, Briefings in Bioinformatics 22 (3) (2021) bbaa189. doi:10.1093/bib/bbaa189.

[22] C.-L. Huang, C.-J. Wang, A GA-based feature selection and parameters optimization for support vector machines, Expert Systems with Applications 31 (2) (2006) 231–240. doi:10.1016/j.eswa.2005.09.024.

[23] M. H. Aghdam, N. Ghasem-Aghaee, M. E. Basiri, Text feature selection using ant colony optimization, Expert Systems with Applications 36 (3) (2009) 6843–6853. doi:10.1016/j.eswa.2008.08.022.

[24] D. Karaboga, B. Basturk, A powerful and efficient algorithm for numerical function optimization: Artificial bee colony (ABC) algorithm, Journal of Global Optimization 39 (3) (2007) 459–471. doi:10.1007/s10898-007-9149-x.

[25] B. Xue, M. Zhang, W. N. Browne, X. Yao, A survey on evolutionary computation approaches to feature selection, IEEE Transactions on Evolutionary Computation 20 (4) (2016) 606–626. doi:10.1109/TEVC.2015.2504420.

[26] Y. Xue, B. Xue, M. Zhang, Self-adaptive particle swarm optimization for large-scale feature selection in classification, ACM Transactions on Knowledge Discovery from Data 13 (5) (2019) 1–27. doi:10.1145/3340848.

[27] Y. Prasad, K. Biswas, Gene selection in microarray datasets using progressively refined PSO scheme, in: Proceedings of the AAAI Conference on Artificial Intelligence, Vol. 29, 2015, pp. 4288–4289. doi:10.1609/aaai.v29i1.9283.

[28] B. Wei, J. Huang, L. Deng, S. Yang, J. Zheng, Y. Huang, Reinforcement learning-based particle swarm optimization with adaptive scoring mechanism for high-dimensional feature selection, Swarm and Evolutionary Computation 98 (2025) 102104. doi:10.1016/j.swevo.2025.102104.

[29] K. Chen, K. Liu, P. Wang, J. Liang, Z. Yang, K. Yu, An efficient evolutionary feature selection algorithm with divide-and-conquer strategy for classification, Expert Systems with Applications 302 (2025) 130416. doi:10.1016/j.eswa.2025.130416.

[30] T. M. Cover, J. A. Thomas, Elements of Information Theory, John Wiley & Sons, New York, NY, USA, 1991. doi:10.1002/0471200611.

[31] Y. Zhang, Q. Yang, A survey on multi-task learning, IEEE Transactions on Knowledge and Data Engineering 34 (12) (2021) 5586–5609. doi:10.1109/TKDE.2021.3070203.

[32] N. Zhang, A. Gupta, Z. Chen, Y.-S. Ong, Evolutionary machine learning with minions: A case study in feature selection, IEEE Transactions on Evolutionary Computation 26 (1) (2021) 130–144. doi:10.1109/TEVC.2021.3099289.

[33] K. Chen, B. Xue, M. Zhang, F. Zhou, An evolutionary multitasking-based feature selection method for high-dimensional classification, IEEE Transactions on Cybernetics 52 (7) (2020) 7172–7186. doi:10.1109/TCYB.2020.3042243.

[34] K. Chen, B. Xue, M. Zhang, F. Zhou, Evolutionary multitasking for feature selection in high-dimensional classification via particle swarm optimization, IEEE Transactions on Evolutionary Computation 26 (3) (2022) 446–460. doi:10.1109/TEVC.2021.3100056.

[35] C.-L. He, Y. Zhang, D. Gong, X. Song, X. Sun, A multitask bee colony band selection algorithm with variable-size clustering for hyperspectral images, IEEE Transactions on Evolutionary Computation 26 (6) (2022) 1566–1580. doi:10.1109/TEVC.2022.3159253.

[36] S. Liu, Q. Lin, L. Feng, K.-C. Wong, K. C. Tan, Evolutionary multitasking for large-scale multiobjective optimization, IEEE Transactions on Evolutionary Computation 27 (4) (2022) 863–877. doi:10.1109/TEVC.2022.3166482.

[37] K. Deb, A. Pratap, S. Agarwal, T. Meyarivan, A fast and elitist multiobjective genetic algorithm: NSGA-II, IEEE Transactions on Evolutionary Computation 6 (2) (2002) 182–197. doi:10.1109/4235.996017.

[38] Q. Zhang, H. Li, MOEA/D: A multiobjective evolutionary algorithm based on decomposition, IEEE Transactions on Evolutionary Computation 11 (6) (2007) 712–731. doi:10.1109/TEVC.2007.892759.

[39] C. A. C. Coello, G. T. Pulido, M. S. Lechuga, Handling multiple objectives with particle swarm optimization, IEEE Transactions on Evolutionary Computation 8 (3) (2004) 256–279. doi:10.1109/TEVC.2004.826067.

[40] Y. Hu, Y. Zhang, D. Gong, Multiobjective particle swarm optimization for feature selection with fuzzy cost, IEEE Transactions on Cybernetics 51 (2) (2020) 874–888. doi:10.1109/TCYB.2019.2946898.

[41] E. Hancer, B. Xue, M. Zhang, An evolutionary filter approach to feature selection in classification for both single- and multi-objective scenarios, Knowledge-Based Systems 280 (2023) 111008. doi:10.1016/j.knosys.2023.111008.

[42] Z. Ye, Y. Zeng, T. Cai, J. Shen, W. Zhou, Q. He, M. Mei, Y. Cen, KMHBO: A knowledge-guided multi-niche hybrid breeding optimization algorithm for high-dimensional multimodal feature selection, Expert Systems with Applications 298 (2026) 129911. doi:10.1016/j.eswa.2025.129911.

[43] S. Liu, H. Wang, W. Peng, W. Yao, Surrogate-assisted evolutionary algorithms for expensive combinatorial optimization: a survey, Complex & Intelligent Systems 10 (4) (2024) 5933–5949. doi:10.1007/s40747-024-01465-5.

[44] F. Han, F. Li, Q. Ling, H. Han, T. Lu, Z. Jiao, H. Zhang, A feature selection method based on feature-label correlation information and self-adaptive MOPSO, Neural Processing Letters 56 (2024) 110. doi:10.1007/s11063-024-11553-9.

[45] J. Demšar, Statistical comparisons of classifiers over multiple data sets, Journal of Machine Learning Research 7 (2006) 1–30.

[46] M. Friedman, A comparison of alternative tests of significance for the problem of m rankings, The Annals of Mathematical Statistics 11 (1) (1940) 86–92. doi:10.1214/aoms/1177731944.

[47] S. Holm, A simple sequentially rejective multiple test procedure, Scandinavian Journal of Statistics 6 (2) (1979) 65–70. doi:10.2307/4615733.